

Estimating Tax Elasticities along the Income Distribution

Clara von Bismarck-Osten

Jakob Wegmann

September 2026

Preliminary and incomplete. The latest version is available [here](#).

Abstract

The elasticity of earnings with respect to taxation governs the trade-off between redistribution and efficiency, but existing estimates rest on strategies confined to one part of the income distribution at a time and cannot recover how the elasticity varies with income. This paper estimates elasticities along the income distribution from variation that joint taxation supplies at every level of income: when the tax unit changes (a child leaves, a couple marries), an individual's tax liability function shifts, and whether the marginal tax rate changes at a given own income depends on partner income. In the universe of French tax returns, earnings elasticities decline with income, most steeply at the bottom, and are higher for women and at older ages. Among the lowest-earning women, the participation response comes from those with higher-earning partners, not from those whose partners also earn little. Accounting for this heterogeneity flattens optimal tax rates at low incomes and raises the assessed gains from reforms towards individual taxation.

We thank Uta Schönberg, Christian Dustmann, Richard Blundell, Costas Meghir, Kirill Borusyak, Pierre Boyer, Emmanuel Saez, Tim Obermeier, Andreas Peichl and Che-Yuan Liang.

Clara acknowledges the Institute for Fiscal Studies, the UCL Stone Centre, the Chartered Institute of Taxation and the Centre for Research and Analysis on Migration for supporting this project's infrastructure expenses, and the Economic and Social Research Council, the Institute for Fiscal Studies and the Thomas and Ulla Kolbeck Foundation for supporting her work. Jakob's work was funded by the German Research Foundation (DFG) through CRC TR 224, Project A05.

1 Introduction

How much revenue a tax raises, and how much it harms economic activity, depend on how elastic (responsive) earnings are to taxation. This elasticity may differ between income groups, but its estimates exploit local variation: bunching designs use a specific income threshold of the tax schedule; tax-reform-IV designs use an income bracket a reform affects, e.g., top incomes for a cut to top rates. An elasticity identified at one income threshold, or for one income bracket, need not be representative of all taxpayers and cannot recover how responses vary with income.

This paper proposes a strategy to estimate heterogeneous tax elasticities of earnings. Tax rates are a function of income, and the effort to earn income is a function of how much it is taxed, which makes an observed change in the tax rate an individual faces endogenous and calls for exogenous variation to instrument it. A tax reform arguably provides such variation, usually for a specific part of the income distribution (the change it predicts at pre-reform incomes serves as the instrument for the observed, endogenous one). The strategy proposed in this paper is no different conceptually, but uses a source that provides many such local changes, throughout the income distribution.

The strategy builds on the observation that, under a tax system that taxes spousal income jointly and gives tax advantages for children,¹ whether a life-cycle event (a child leaving the family tax unit, or marriage) changes an individual's marginal tax rate depends on partner income. Consider two women with the same earnings but different partner incomes. When a child reaches the age at which it can no longer be counted in the family tax unit (the event), they may face different rate changes: one remains in her bracket; the other, with the higher-earning partner, is pushed into the bracket above. The strategy formalises such comparisons within narrow cells of own income and bands of partner income, between individuals predicted (from pre-event incomes) to cross a threshold and those predicted to remain in their bracket. Both groups face the same event, and identification comes from the difference between them, on the assumption that their earnings would have continued to grow alike without the tax change. Since the groups are formed from pre-event incomes, the assumption rules out bunching in

¹Under family taxation such as in France, the couple's income is divided by a quotient that rises with the number of children (2 for a married couple, 2.5 with one child, 3 with two) before the tax schedule applies.

anticipation (holding income just below thresholds that will apply to them only after the event). It also has two testable implications: before the event, the two groups' earnings paths should move alike, and over the range of partner income where the event changes no rate, earnings growth should not vary with partner income. Because children leave the tax unit, and couples marry, at all levels of income, and the schedule has several brackets, there are many such tax changes an event can induce, each of which can be analysed on its own.²

For identifying causal tax elasticities, the variation has further advantages over reform-based designs. A tax reform typically changes the rates of a particular income bracket, and the elasticity is identified by comparing the income growth of affected and unaffected groups, under the assumption that, without the reform, the incomes of the two groups would have grown alike. Across income groups, the assumption is a strong one, for two main reasons: unusually high incomes tend to fall back towards the average (mean reversion), and changes in inequality mean that incomes grow at different rates along the distribution for reasons unrelated to taxation. Controlling for income does not easily address these two sources of bias, because which taxpayers a reform affects is itself a function of income: the controls then tend to absorb the variation that identifies the elasticity, a problem a large literature has proposed ways to handle.³ By contrast, our variation is *within-income*, in the sense that exposure to the tax change depends on partner income. Own income can then be held fixed, which alleviates both concerns. A second advantage is that the tax change is a shift of the bracket thresholds of an unreformed schedule. Monotonicity—the individual's tax rate moving in the direction predicted by the event—is then more likely to hold than under a reform,⁴ and the rate change comes without the other, potentially confounding provisions of a reform.⁵

²Analysing each on its own also avoids the negative weights that pooling changes of different strength produces when responses differ across individuals (Graber et al., 2026).

³Gruber and Saez (2002) and H. J. Kleven and Schultz (2014) use richer income controls, Weber (2014) longer lags of the income that defines exposure, and Jakobsen and Sogaard (2022) and Graber et al. (2026) differencing against placebo periods with graphical validation; grouping estimators use group-level rather than individual variation (Blundell et al., 1998; Burns & Ziliak, 2017), and Kumar and Liang (2020) allow responses to vary across individuals.

⁴Monotonicity requires that no individual's realised marginal rate moves against the change the schedule predicts. Under a reform that raises the rates of one income group, for example, an individual who stops working in response ends up at a zero rate where a rise was predicted. Family events move the bracket thresholds, not the rates, and partner income limits how far the household's rate can fall when own income falls, so a violation needs a much larger fall in income than under a reform (Appendix H gives the formal argument and a simulation).

⁵A reform is publicised and draws attention, and taxpayers respond less to changes they do not notice, so an elasticity from a well-advertised reform and one from a quiet shift of the schedule need not be the same (Saez et al., 2012). A reform also rarely moves one rate alone, and usually changes the tax base with it (Kopczuk,

We apply the strategy to the universe of French tax returns for 2003 to 2023, reshaped from a panel of tax units into a panel of individuals that follows each person across tax units. The code is open source, and the panel can serve other studies of French tax data. With around forty million individuals a year, it can be split finely by own and partner income and by the other dimensions along which we study heterogeneity.

We find the elasticity of own earnings (work income) with respect to the net-of-tax rate,⁶ the earnings elasticity, to decline with income, most steeply at the bottom, for both genders. Women’s earnings elasticity falls from 0.45 for own incomes below €10,000 to 0.31 at incomes approaching €70,000; men’s from 0.32⁷ to 0.19.⁸ Averaged over the French taxpayer distribution, the earnings elasticity is 0.23 (0.29 for women, 0.17 for men), and the participation elasticity, with respect to the same rate, is 0.25. Splitting the estimates by partner income, which the several thresholds of the schedule make possible, shows which households drive the large elasticities at low incomes. Among the lowest-earning women, the participation response comes from those with higher-earning partners rather than from those whose partners also earn little, with participation elasticities of 0.69 and 0.35 in the middle and top thirds of partner income, against 0.02 in the bottom third.

Both normative models, which ask how taxes should be set, and positive predictions of the effects of reforms are sensitive to the heterogeneity we document. Optimal-tax models trade off redistribution against the earnings that taxation discourages (Mirrlees, 1971), and Saez (2001) expresses the optimum as a formula in which the marginal rate at each level of income depends on three things: the elasticity there, the welfare weights of those above it and the shape of the income distribution. Feeding the estimated profile of elasticities over income into an illustrative calibration of that formula (Brewer et al., 2010; Saez, 2001), rather than assuming a single elasticity (as is common), flattens optimal marginal tax rates at the bottom of the income distribution, lowering them by about 17 percentage points on average at incomes below €10,000 in 2015 prices. Feeding the profiles over own and partner income into the evaluation of

2005); and since individuals optimise over the whole schedule rather than over the rate at their own income, an estimate that uses only that rate leaves out the other rates the reform moved (Kumar & Liang, 2020).

⁶The household net-of-tax rate: one minus the marginal tax rate the household faces at the individual’s income.

⁷The value for own incomes of €10,000 to €20,000; below €10,000, only one of the men’s comparisons meets the precision screens (0.45, standard error 0.49).

⁸The estimates cover own incomes up to €70,000 in 2015 prices; in nominal terms, at least 95.6 percent of men and 97.0 percent of women filing a return in 2023 earn less.

reforms towards more individual taxation (Bierbrauer et al., 2024), rather than one elasticity per spouse, raises the assessed welfare gains overall and relocates them, concentrating them where both partners earn little and lowering them over much of the remaining distribution.

A limitation of the variation is that, while it is not confined to one part of the income distribution, it is confined to the life-cycle event considered, a child’s exit from the tax unit. Parents are in their early fifties (51 to 53 on average) when their first-born leaves, and the estimates may reflect the responses of people approaching retirement rather than those of taxpayers in general. The variation is, however, not tied to this one event. It arises in the same way at other family events, and so at other ages, such as marriage, which we use as a second design (Appendix G), or the birth of a child. Age also varies within this event, with 85 percent of the women and 82 percent of the men aged 45 to 59, which allows (i) age fixed effects in every estimation and (ii) sample splits by age as a further heterogeneity analysis. The split shows that elasticities are lower below the minimum retirement age.⁹

To the literature on heterogeneity in tax responsiveness by income, the paper adds a profile of tax elasticities along the income distribution, estimated within a single setting, design and outcome and reaching down to low incomes. Figure 1 collects the quasi-experimental estimates of earnings (work-income) responses published in top field or general-interest journals. Each is either local to a narrow slice of the income distribution or a single average over a wide range, and most concern the top three deciles. The variance in the available estimates is large (Neisser, 2021), and it is difficult to establish how much reflects heterogeneity in responsiveness and how much differences in design and in the income measure studied.

To the literature on the taxation of couples, the paper adds evidence along partner income, the dimension its variation also spans. Normative work gives opposing prescriptions on optimal rates depending on the assumed difference between primary- and secondary-earner elasticities (Boskin & Sheshinski, 1983; Golosov & Krasikov, 2025).¹⁰ Positive work on actual switches between joint and individual taxation, or on comparable policy variation, finds that joint taxation lowers the employment and earnings of secondary earners but not of primary earners,

⁹Restricting the sample to parents below 59 at the event lowers women’s elasticities but keeps the shape of their profile, and cuts men’s earnings response from 0.17 to 0.06 (own incomes from €10,000 up).

¹⁰With a common elasticity, optimal marginal rates on secondary earners are on average *higher* than on primary earners (Golosov & Krasikov, 2025); with the secondary earner assumed more responsive, the inverse-elasticity rule, rates inversely proportional to elasticities, prescribes the opposite (Boskin & Sheshinski, 1983).




Figure 1: Published estimates of work-income responses

Notes: The figure shows quasi-experimental estimates of work-income tax elasticities from studies published in top field or general-interest journals. Colours denote the outcome (in-figure key); circles denote reform designs and squares bunching designs. Work income incl. contributions is measured gross of employer contributions (payroll contributions in Lehmann et al., employer-administered pension contributions in the Danish studies). Point markers sit at kinks or local-regression centres; horizontal dashed segments give the income range from which the identifying variation comes. Estimates without a locatable range appear in the right panel, where horizontal position is uninformative. Vertical lines are 95% confidence intervals where an elasticity standard error is reported. Percentile ranks refer to each study's ranking universe. Papers estimating taxable-income elasticities for populations broader than wage earners^h or Frisch elasticitiesⁱ are not included. Exact populations, horizons, estimands, treated ranges and mapping grades are reported in Appendix Table I.1.

^a One-child (left) and two-plus-children (right) EITC kinks, wage earners only. ^b EITC phase-in (left) and phase-out (right) regions. ^c Top-20 and nested top-10 percent samples. ^d All treated job switchers and, above, the top decile; long-run elasticities, whereas the other reform estimates are one-to-four-year responses. ^e Husbands (left) and wives (right), both uncompensated. ^f Confidence interval truncated below; the segment covers the bracket-creep kinks holding roughly nine-tenths of treated taxpayers (full support about P21–P98). ^g Approximate: the treated interval of 0.3 to 2.1 minimum wages, which covers half of French employees. ^h Feldstein (1995) and Gruber and Saez (2002); surveyed in Saez et al. (2012). ⁱ From temporary tax holidays: 0.41 in Iceland (Sigurdsson, 2025), 0.025 in Switzerland (Martínez et al., 2021). ^j Top-kink bunching in taxable income among wage earners; no elasticity standard error is reported. The right-panel entry is the same study's schedule-wide wage-earnings reform estimate. ^k First central-government kink, taxable labour income; the paper reports 0.00 (rounded), without a standard error. ^l Taxable labour income, uncompensated; the marker is the male estimate, 0.21 (0.07); the female estimate, 1.44 (0.83), exceeds the axis and its confidence interval does not exclude zero. ^m Taxable wage income; the reform estimate, 2.42 (0.19), exceeds the axis; the same study's bunching estimates on the same schedule are 0.09–0.41. ⁿ Observed (friction-attenuated) elasticity of taxable income at the wage-earner notches, derived as the paper instructs: 0.000–0.004 across the four notches, no standard error applicable. The friction-corrected structural bounds are 0.001–0.074. ^o The dashed line connects Jakobsen-Sogaard's local estimates at successive income thresholds (DKK 400,000–750,000); it orders estimates from the same study and does not denote interpolation or an estimated continuous profile.

who are typically men (Herold et al., 2025; Herold & Wallossek, 2024; Isaac, 2023; Kalíšková, 2014; LaLumia, 2008; Selin, 2014), which is consistent with the heterogeneity we map by own and partner income.

Other dimensions of heterogeneity to which we contribute evidence are the income measure, gender and the time horizon. Broad earnings elasticities, which add self-employment income to work income, closely track the earnings elasticities.¹¹ Women respond more than men in our estimates, but by less than in the literature,¹² which may reflect the life-cycle stage our estimates cover. With the time horizon recently highlighted as central to the gap between micro and macro estimates (H. Kleven et al., 2025), we provide estimates by year since the event, up to four years.

The rest of the paper is structured as follows. Section 2 describes the French tax system and the construction of the individual-level panel. Section 3 develops the empirical framework and Section 4 presents the results. Section 5 examines how accounting for heterogeneity affects optimal tax schedules and the evaluation of reforms towards more individual taxation. Section 6 concludes.

2 Background

2.1 Institutional setting

The variation explored in this paper stems from two features of the French tax system: (i) the dependence of the tax liability on household composition through the *family quotient*, and (ii) the step-wise shape of the marginal tax schedule.

What makes the French tax system one of “family taxation” is that the tax burden (equation (1)) is a function of the family quotient, which increases with the number of family members: the quotient is 1 for singles, 2 for childless married (or civil-union) couples, 2.5 with

¹¹From own incomes of €10,000 up. Our variation does not identify responses of taxable or capital income, the object of a large literature (e.g., Feldstein, 1995; Gruber & Saez, 2002; Neisser, 2021; Saez et al., 2012).

¹²The literature finds elasticities for married women and single mothers several times those of men (Blomquist & Selin, 2010; Eissa, 1996; Eissa & Liebman, 1996).

one dependent child and 3 with two; each further child raises it, and the tax advantage a child confers is capped.^{13,14}

The tax liability is defined (i) under individual filing and (ii) under joint filing as

$$(1) \quad (i) \ T_t(z_{it}), \quad (ii) \ T_t^H(z_{it} + z_{pt}; Q_{it}) = Q_{it} T_t\left(\frac{z_{it} + z_{pt}}{Q_{it}}\right),$$

where T_t denotes the year- t French income tax formula for a single individual, z_{it} and z_{pt} are own and partner taxable income, and Q_{it} is the family quotient.


Figure 2 shows the 2015 schedule. Panel (a) gives the marginal tax rate by household composition: the brackets of the single schedule scale with the quotient, so that a childless couple reaches each rate at twice the income of a single, and the child-related shifts are truncated by the cap. Panel (b) shows the implication for a married couple in own and partner income: given the quotient, only the sum of the two incomes enters the tax formula, so the marginal rate is constant along diagonals and, at given own income, rises with partner income. The resulting tax liabilities are shown in Appendix A.

The effective schedule differs from the statutory one through the *décote*, a tax reduction for households with a small liability, which is phased out as the liability rises. Over the phase-out range the effective marginal rate exceeds the statutory rate (in 2015, 24.5 instead of 14 percent); above it, the statutory rate applies. At the child's fiscal exit, the instrument and the realised rates of Section 3 are computed on the effective schedule, including the *décote* and the cap on the child advantage.

Under family taxation, unlike under individual taxation, family events change the tax function that applies to an individual's income—the variation studied subsequently. The fiscal exit of a dependent child lowers the family quotient and leaves the tax base unchanged; marriage merges two tax units into one, pooling the spouses' incomes and doubling the quotient. Table 1 summarises what each event changes.

¹³The tax advantage associated with children is capped, such that the benefit from an additional child does not exceed a predetermined threshold.

¹⁴Other factors also influence the family quotient. For example, it increases for single parents and parents of disabled children. We do not consider these cases in our analysis.



(a) By household composition

(b) Married couple, by own and partner income

Figure 2: Marginal tax rates under family taxation

Notes: Statutory schedule of 2015, before the *décote*, with the cap on child-related advantages applied; Q is the family quotient. Panel (a): by household composition and total taxable income of the tax unit. Panel (b): married couple without children (two fiscal parts), by own and partner taxable income; the 14, 30 and 41 percent brackets begin at total taxable income of €19,400, €53,582 and €143,652, twice the single thresholds. The corresponding tax liabilities are shown in Appendix Figure A.1. Source: Own calculations based on OpenFisca France.

Table 1: Family events and the tax liability function

Event	Tax base	Family quotient Q
Child's fiscal exit	unchanged: the couple's pooled income	2.5 $\rightarrow$ 2 with one dependent child; 3 $\rightarrow$ 2.5 with two
Marriage	own income $\rightarrow$ the spouses' pooled income	1 for each spouse $\rightarrow$ 2 for the joint unit

Notes: What each event changes in the tax liability function of equation (1). Households with more dependent children see the analogous quotient change, subject to the cap on the child advantage; the marriage row refers to two previously single, childless individuals.

2.2 Data

We use the French administrative tax panel known as POTE.¹⁵ The dataset covers all income tax returns filed in France between 2003 and 2023—all individuals required to file a French income declaration. POTE includes all variables submitted in the tax return (Form No. 2042), intermediate calculations for determining the final tax liability, and all information shown to the household on its official tax notice. It also includes detailed tax-relevant

¹⁵POTE—Panel On The Taxation of Households, INSEE-DGFiP. Data accessed via CASD (<https://www.casd.eu/>).

characteristics—the number of dependent children, marital status and filing status, gender, and the municipality—relevant to the subsequent analysis.

Tax returns are filed at the tax-unit level (the *foyer fiscal*), but the POTE database has the important advantage of including individual-level identifiers. It permits the (to our knowledge) first construction of a longitudinal individual-level panel, essential to our design: it lets us link individuals across tax units and follow their incomes before and after family events that change the tax liability function, such as marriage.

The result is a panel of approximately 50 million individuals per year before cleaning; the cleaned year files used in the analysis contain between 39.1 million (2003) and 42.9 million (2023) individuals. In 2023, 62 percent of men and 57 percent of women declare positive work income, and 50 and 44 percent, respectively, are married or in a civil union. From this panel we draw the two event samples: individuals experiencing the fiscal exit of their first child, and individuals marrying; Appendix F.1 describes their construction and characteristics.

3 Empirical framework

This section sets out the model and the identification problem (Section 3.1), introduces the reform-based and event-based designs (Section 3.2), defines the cells within which comparisons are made, states the identifying assumptions and establishes the causal interpretation of the estimands (Section 3.3), and describes aggregation, estimation and inference (Sections 3.4 and 3.5).

3.1 The model and the identification problem

We model a jointly taxed couple, extending the standard one-person model (Jakobsen & Sørensen, 2022; Saez et al., 2012) to a unitary household maximising a quasi-linear and iso-elastic utility function, with the costs of earning income additively separable across spouses.

The household solves

$$\max_{z_{it}, z_{pt}} c_t - \sum_{s \in \{i, p\}} \frac{n_{st}}{1 + \frac{1}{\epsilon_s}} \left(\frac{z_{st}}{n_{st}} \right)^{1 + \frac{1}{\epsilon_s}} \quad \text{s.t.} \quad c_t \leq z_{it} + z_{pt} - T_t^H(z_{it} + z_{pt}),$$

where z denotes income, with subscripts i for the individual and p for the individual's partner, c household consumption, and T_t^H the household tax function of Section 2.1.¹⁶ The cost of earning income is governed by two individual parameters, the elasticity ϵ_i and potential income n_{it} , the income i would earn at a zero marginal tax rate (Saez et al., 2012). Any change in income that is not a response to the rate runs through n_{it} . Quasi-linearity rules out income effects, so compensated and uncompensated elasticities coincide; income effects are added in Appendix D.

Because the couple is taxed jointly on the sum of spousal incomes, the marginal net-of-tax rate is common to the two spouses, $\tau_{it}^H \equiv 1 - T_t^H(z_{it} + z_{pt}) = \tau_{pt}^H$. With separable costs, the spouses' incomes are linked only through the household rate. The household problem then separates into two individual problems, and each spouse's first-order condition yields the income-supply function $z_{it} = n_{it} (\tau_{it}^H)^{\epsilon_i}$ (Appendix B relaxes separability).¹⁷ Taking logs and first differences gives the estimating equation

$$(2) \quad \Delta \ln z_{it} = \epsilon_i \Delta \ln \tau_{it}^H + \Delta \ln n_{it}.$$

The parameter of interest is ϵ_i , the earnings elasticity: the elasticity of own income with respect to the marginal net-of-tax rate the individual faces.¹⁸ The error term, $\Delta \ln n_{it}$, varies across individuals and years—transitory shocks, life-cycle and secular growth—and, at a family event, may respond to the event itself; the design does not assume it exogenous, as Section 3.3 will make clear. Elasticities are individual throughout; the objects of interest are averages of ϵ_i over subpopulations defined by own and partner income (characterised in Section 3.4). The iso-elastic form makes equation (2) exact and ϵ_i a constant. For any convex cost of earning income, log-linearising the first-order condition around the pre-event optimum gives the same equation, with ϵ_i the local elasticity at the pre-event income and rate (Appendix B). Our estimates are averages of local elasticities, for rate changes of the size the events induce;

¹⁶The model is written for a couple filing jointly, the tax unit of the child's fiscal exit design, where the event lowers the quotient Q_{it} of equation (1). Marriage instead merges two individually filing units into one joint unit. Before marriage, the household pays $T_t(z_{it}) + T_t(z_{pt})$, so the rate in i 's first-order condition, τ_{it}^H , is the own rate $1 - T_t'(z_{it})$.

¹⁷This is the formal counterpart of the standard practice of pairing the one-person model with the marginal rates a household tax system produces (Bick & Fuchs-Schündeln, 2018; Gruber & Saez, 2002).

¹⁸Appendix B relaxes separability: equation (2) is unchanged, with ϵ_i read as the elasticity of own income with respect to the household rate, $\epsilon_i^H \equiv \epsilon_{ii} + \epsilon_{ip}$ —the sum of its responses to the own and to the partner's rate, each a jointly re-optimised household response. The appendix shows how the two events identify the components separately.

the design does not distinguish whether they differ across the income distribution because preferences differ or because a common cost function is evaluated at different incomes and rates.

A naive estimation of equation (2) would not produce causal estimates of ϵ_i . Because the tax schedule depends on income, the marginal net-of-tax rate is a deterministic function of the dependent variable, which means that any change in potential income changes realised income, and with it $T_t^{H'}$ —the regressor of equation (2) is thus correlated with its error term by construction.

The literature handles this endogeneity in two ways: (i) bunching designs, which impose enough structure to infer the elasticity from excess mass at kinks of the schedule (H. J. Kleven, 2016; Saez, 2010); and (ii) IV designs that draw arguably exogenous variation from tax reforms—the approach we build on, and against which we compare our design in the next subsection.

3.2 Our variation relative to reform variation

The reform-based instrument of Gruber and Saez (2002) applies the post- and pre-reform schedules to the same pre-reform income, predicting each individual’s rate change from the legislated change alone,¹⁹

$$\Delta \ln \tau_{it}^{\text{reform}} \equiv \ln(1 - T_t^{\text{post}}(z_{i,t-k})) - \ln(1 - T_t^{\text{pre}}(z_{i,t-k})).$$

Validity requires

$$\text{cov}(\Delta \ln \tau_{it}^{\text{reform}}, \Delta \ln n_{it} \mid z_{i,t-k}) = 0:$$

conditional on lagged income, potential-income growth must be unrelated to the predicted rate change. The assumption is demanding because exposure to reforms is usually *between*-income. A reform changes rates in one bracket and not in others, so the instrument sorts individuals by lagged income, and treated and untreated individuals differ in the level of income. The level of income predicts income growth for reasons unrelated to taxes—a high base-year income is in part transitory and reverts, and incomes at different points of the distribution have grown

¹⁹See Saez et al. (2012) for an overview of the literature.

at different rates over time—so the instrument is correlated with potential-income growth. The literature mitigates these concerns with richer income controls (Gruber & Saez, 2002; H. J. Kleven & Schultz, 2014), longer instrument lags (Weber, 2014), and differencing against placebo periods or cohorts with graphical validation (Graber et al., 2026; Jakobsen & Sogaard, 2022).

A further concern, developed in Appendix H, is that the IV estimand weights individual responses by the realised change in their marginal rate, and reform variation can move realised rates in opposite directions. When a reform raises the rate, an individual who responds by leaving employment, or by cutting hours enough to fall below the reformed bracket, faces a lower rate where a rise was predicted. Monotonicity (Assumption 3 below) is then violated, and the estimand no longer aggregates behavioural responses with convex weights. The events we use shift the thresholds of the schedule rather than its gradient, so that a violation requires income to fall not merely back below the threshold the event pushed it across but below the lower threshold of the pre-event bracket, and partner income places a floor under the rate an individual can reach by reducing own income.

Instead of using reform variation, we exploit the fact that family events change the tax function that applies to an individual’s income (Table 1). Our main design uses the fiscal exit of a dependent child, which lowers the family quotient and leaves the tax unit otherwise unchanged (Section 2.1); marriage, the second event we use, instead merges two tax units. Where a reform changes the bracket formula, the event changes the income to which the formula is applied. We index the two schedules by $g \in \{0, 1\}$: the individual is taxed under the pre-event schedule ($g = 0$) or the post-event one ($g = 1$). From here on, “schedule” refers to the household tax function T^H under configuration g , and “formula” to the calendar-dated bracket formula T_t of equation (1). Dates are expressed in event time, $r \equiv t - m_i$, where m_i is the calendar year of the event—for the fiscal exit, the year in which fiscal attachment of the first-born ends, at 18 or, if extended, at 21 (25 if in education)—and the schedule indicator in event time is $g_{ir} \equiv \mathbb{1}[r \geq 0]$: every individual in our designs is taxed under the pre-event schedule for $r < 0$ and under the post-event schedule from $r = 0$ onwards.²⁰ The design does not require the

²⁰In its simplest form, the design is therefore a two-period comparison—the pre-event year $r = -1$ against a post-event year; the event-time profiles of Section 4 extend it horizon by horizon.

event to be exogenous. Comparisons are made within cells of individuals who all experience the event; Section 3.3 states what identification requires of it.

Per-part income under schedule g , the income to which the bracket formula applies, is total tax-unit income divided by the quotient, $b_{ir}(g)$ —for an individual in a couple filing jointly, $b_{ir}(g) = (z_{ir} + z_{pr})/Q_g$; for an unmarried individual, $b_{ir}(0) = z_{ir}$. Like the reform instrument, the event instrument is a simulation. It applies the pre-event year’s formula to pre-event ($r = -1$) per-part income twice, under $g = 0$ and under $g = 1$, and takes the difference in log net-of-tax rates. The two applications differ only in g —the quotient and, at marriage, the tax unit:

$$(3) \quad Z_{ir} \equiv g_{ir} \times \underbrace{\left[\ln(1 - T'_{-1}(b_{i,-1}(1))) - \ln(1 - T'_{-1}(b_{i,-1}(0))) \right]}_{\equiv Z_i, \text{ the dose}},$$

where the formula is dated in event time as well, T_{-1} being that of the individual’s pre-event calendar year.²¹ The dose Z_i is thus fixed at the pre-event year, and the indicator g_{ir} switches it on once the post-event schedule applies. Within a cell the dose takes few values, the log-rate differences between brackets, so that a dose δ names a bracket transition. It is zero for individuals whose per-part income is predicted to remain in its bracket.²² At the child’s fiscal exit the tax unit is unchanged and the quotient falls—from 2.5 to 2 with one dependent child, from 3 to 2.5 with two—so per-part income weakly rises at every level of household income, and for almost every individual the event predicts a weakly higher marginal rate, $Z_i \leq 0$.²³

This instrument definition leaves *within-income* variation: conditional on own income and the pre-event bracket, whether or not an individual is pushed over a bracket threshold is a function of partner income (Figure 2, Panel (b)): at fixed own income, the predicted rate change

²¹ Z_{ir} is the excluded instrument of Graber et al. (2026)—the interaction of a post-reform indicator with the simulated rate change—written in event time.

²²The dose and the realised net-of-tax rate τ_{ir}^H of Section 3.5 come from one and the same effective schedule: the hypothetical bracket under the post-event quotient and the realised bracket of every year are assigned from the same set of kinks and effective marginal rates, by year and number of parts (the statutory brackets and the features of the effective schedule described in Section 2.1, generated from OpenFisca’s *barème*), so that these features enter instrument and outcome identically.

²³For 5 percent of the women and 6 percent of the men, the effective schedule of Section 2.1 predicts a lower rate at the child’s exit; these rate cuts enter the estimation as doses of their own. At marriage, two tax units with quotient 1 merge into a single unit with quotient 2; predicted increases and decreases coexist, and we restrict the analysis to observations with predicted increases.

varies with partner income).²⁴ Own income can therefore be held fixed without absorbing the identifying variation.

3.3 Identification and an illustration

Given the two schedules, we let $Y_{ir}(g)$ denote potential log income at event time r under schedule $g \in \{0, 1\}$, and $\text{NTR}_{ir}(g)$ the corresponding potential log net-of-tax rate; the post-event schedule applies from $r = 0$ onwards (g_{ir} , Section 3.2). $Y_{ir}(0)$ is not a no-event counterfactual: the event occurs under both potential outcomes; the schedule alone is varied. The schedule change shifts the log net-of-tax rate by $\Phi_{ir} \equiv \text{NTR}_{ir}(1) - \text{NTR}_{ir}(0)$. The construction is stated for earnings, with log work income as the outcome; the participation indicator, $P_{ir} \equiv \mathbb{1}\{z_{ir} > 0\}$, is treated in the same way, with its model counterpart in Appendix D and its estimation in Section 3.5. The potential outcomes map to observed outcomes through

$$(4) \quad \begin{aligned} Y_{ir} &= Y_{ir}(0) + \mathbb{1}[Z_{ir} \neq 0] [Y_{ir}(1) - Y_{ir}(0)], \\ \text{NTR}_{ir} &= \text{NTR}_{ir}(0) + \mathbb{1}[Z_{ir} \neq 0] [\text{NTR}_{ir}(1) - \text{NTR}_{ir}(0)]. \end{aligned}$$

For zero-dose individuals the two schedules prescribe the same marginal rate at their pre-event income—the schedule change is a treatment only where $Z_i \neq 0$ —so the mapping assigns the predicted-control group its untreated outcomes at every horizon; the group’s realised rate changes, its first stage, provide a direct check (Appendix E).

In line with the model of Section 3.1, we assume potential log income is linear in the potential rate, with the slope invariant across the two schedules (as in Graber et al., 2026; Saez et al., 2012):

$$(5) \quad Y_{ir}(\text{NTR}, g) = \epsilon_i \times \text{NTR} + \ln n_{ir}(g),$$

where ϵ_i is the individual elasticity of Section 3.1 and $n_{ir}(g)$ the zero-rate income of Section 3.1 under schedule g , whose schedule argument collects every channel from the schedule change to income other than the marginal rate—the argument Assumption 2 below restricts.²⁵

²⁴Jakobsen and Sogaard (2022) develop the associated identifying assumption of constant trend differentials and its graphical validation.

²⁵With non-separable costs, the household model of Appendix B delivers the same form locally, with ϵ_i read as the elasticity of own income with respect to the household rate, $\epsilon_i^H = \epsilon_{ii} + \epsilon_{ip}$. The proof of Proposition 1

For both the identification argument and the estimation that follows, we compare at a disaggregated level, for two reasons—the credibility of the identifying assumptions and the heterogeneity of the elasticities. Comparisons are made within cells of individuals alike in the determinants of income growth, and each dose δ in the support of Z is compared separately to the zero-dose group of its cell. The identifying assumptions then need to hold only among near-identical individuals, where they are easiest to defend and to test; and with heterogeneous elasticities, pooling doses can place negative weights on individual responses (Callaway et al., 2026; Graber et al., 2026), whereas dose-by-dose contrasts aggregate them transparently.²⁶ Formally, cells fix own income within narrow bins, the event year and the event type exactly, and partner income within an interval. Let $X_i \equiv (z_{i,-1}, z_{p,-1}, m_i, e_i)$ collect these predetermined variables, with e_i the event type (the pre-event constellation: the number of dependent children and single or joint filing); let the cells $\mathcal{I}_c$ be the elements of the resulting partition of the support of X , and write $c(X_i)$ for the cell of individual i .²⁷ Throughout, expectations subscripted by c , $\mathbb{E}_c[\cdot]$, are taken over the individuals of cell $\mathcal{I}_c$.

Within cells, we impose three assumptions:

Assumption 1 (Conditional common trends). *For every cell, horizon r , and dose δ in the within-cell support of Z —writing $\Delta_r V \equiv V_r - V_{-1}$ for growth relative to the pre-event year—*

$$\begin{aligned}\mathbb{E}_c[\Delta_r Y(0) \mid Z_i = \delta] &= \mathbb{E}_c[\Delta_r Y(0) \mid Z_i = 0], \\ \mathbb{E}_c[\Delta_r \text{NTR}(0) \mid Z_i = \delta] &= \mathbb{E}_c[\Delta_r \text{NTR}(0) \mid Z_i = 0].\end{aligned}$$

Assumption 1 holds whenever untreated growth does not vary within cells at all, $\mathbb{E}[\Delta_r Y(0) \mid X_i = x] = f_y^r(c(x))$ for all x , and likewise for $\Delta_r \text{NTR}(0)$: since own income, the event year, and the event type are fixed by the cell, this requires that the exact position of partner income within the cell’s partner-income band does not predict untreated growth. By equation (5),

does not need linearity either. Equation (5) enters it only through $Y_{ir}(1) - Y_{ir}(0) = \epsilon_i \Phi_{ir}$; for any model in which the schedule affects income only through the rate and through $n_{ir}(g)$ —so that, under Assumption 2, an unchanged rate means unchanged income—the proposition holds with ϵ_i read as the ratio of the individual’s income change to her rate change, $[Y_{ir}(1) - Y_{ir}(0)]/\Phi_{ir}$. That ratio may then differ across horizons and with the size of the change, which the linear form rules out.

²⁶Conversely, were elasticities homogeneous within cells, the dose separation and the monotonicity condition below would be dispensable—the pooled within-cell contrast identifies the common ϵ . Even then the cells would remain necessary, because treated and controls differ unconditionally in own income, event year and event type, the dimensions along which untreated growth varies—Assumption 1 holds only within them.

²⁷Given X_i , both the pre-event tax bracket and the instrument of equation (3) are determined.

$\Delta_r Y(0) = \epsilon_i \Delta_r \text{NTR}(0) + \Delta_r \ln n(0)$, so the restriction concerns the untreated evolution of net-of-tax rates and of zero-rate income—the latter including the direct effects of the event, which occur under both schedules and difference out between the two groups. Although stated conditional on the constructed dose Z_i , the assumption restricts only predetermined variation: within a cell, Z_i is a deterministic function of the position of the pre-event incomes within the cell’s bands—predominantly of partner income, own income moving only within its narrow bin—so conditioning on $Z_i = \delta$ partitions the cell by predetermined income positions, and Assumption 1 is a restriction on untreated growth across these predetermined groups. Sorting of the event’s timing on the tax consequence—say, delaying a child’s fiscal exit by extending attachment when the predicted bracket change is costly—would violate this restriction within cells, which fix the event year; so would anticipatory bunching below thresholds that bind only after the event; the density tests of Appendix E target both.

Table 2: The example cell: characteristics by predicted treatment status

	Men		Women	
	Control	Treated	Control	Treated
Work income (€)	34,960	35,022	34,885	34,919
Partner’s work income (€)	22,098	28,979	22,592	30,524
Age	54.0	54.7	53.3	53.8
Partner’s age	52.1	53.1	54.3	54.8
Marginal rate at $r = 0$				
Predicted	14%	30%	14%	30%
Realised (mean)	16%	24%	16%	24%
N	4,295	2,586	1,687	1,735

Notes: The cell comprises individuals with pre-event own income in $[\text{€}34,000, \text{€}36,000]$, a pre-event marginal rate of 14% and 2.5 fiscal parts, whose first-born exits the tax unit at $r = 0$ in 2006–2020. The two groups are balanced on own income and age, and differ in partner income, which determines whether the fall in the family quotient pushes per-part income across the threshold of the post-event schedule. Incomes are measured in the pre-event year, ages at the event; the rate rows refer to $r = 0$, the first year under the post-event schedule.

How can Assumption 1 be assessed, and what does it mean economically? Table 2 and Figures 3 and 4 explain with an example: individuals with pre-event own income between €34,000 and €36,000, at a pre-event marginal rate of 14%, in 2.5-part tax units whose first-born exits at $r = 0$ (event years 2006–2020). Individuals whose partner earns enough are predicted into the 30% bracket once the quotient falls to 2 ($Z_i = \delta = \ln(0.70/0.86) \approx -0.21$), and the rest are predicted to remain in the same bracket ($Z_i = 0$). Figure 3 plots the income paths of the two groups; their alignment before the event is the assumption’s primary testable




Figure 3: Work income around the first-born's fiscal exit, by predicted treatment status

Notes: Mean work income in euro by year relative to the event; bands are 95% confidence intervals; the dashed vertical line marks the event. Raw means for the example, without matching or covariate adjustment, so the figure shows exactly the comparison the design uses. The series distinguish individuals predicted to change tax bracket at the event ($Z_i = \delta < 0$) from those predicted not to ($Z_i = 0$). Before the event the two paths lie within €150 of each other for women and €270 for men; two years after it, treated women earn €1,600 (5 percent of the pre-event level) less than control women and treated men €230 less than control men, and the shares in work are 0.89 against 0.94 for women and 0.87 against 0.89 for men. Table 2 describes the sample.

implication.²⁸ Figure 4 shows the split of the sample by partner income into a *treatment region*, where the exit is predicted to raise the marginal rate, and a *validation region*, where it is not. Economically, the assumption requires partner income to have no channel to income growth of its own within the cell's partner-income band. Such channels are conceivable. A higher-earning partner may allow the individual to reduce work once the child has left, or the individual may adjust own work to changes in the partner's (Appendices B and D take up cross-spouse responses and income effects). Either would appear as a gradient of income growth in partner income. The predicted rate change is a step function of partner income, whereas such channels would be smooth in it. The validation region, where partner income varies with no accompanying rate change, tests for the gradient directly. Absent such a channel, income growth, the vertical distance between the two series of Figure 4, is flat in partner income across the region. In a single cell such as the example, this flatness is a second check, next to the alignment before the event. Section 4 reports the same check aggregated over all cells of the design, in the pooled analogue of Figure 4 (Figure 5).²⁹

²⁸Every exhibit of the example applies the civil-status restriction of the estimation sample (Appendix F.1): individuals whose civil status changes at the event are excluded.

²⁹Narrower partner-income bands weaken the requirement, and the matching implementation of Section 3.5 additionally aligns pre-event income paths within cells.




Figure 4: Identifying variation at the first-born's fiscal exit

Notes: Sample: women of the example (pre-event own income in $[\text{€}34,000, \text{€}36,000]$, pre-event marginal rate 14%, 2.5 fiscal parts, first-born exiting the tax unit at $r = 0$ in 2006–2020). Left axis: mean work income at $r = -2$ (before) and, for each individual, the average of $r = +1$ and $r = +2$ (after), by 20 equal-count bins of partner income (about 170 women per bin; the lowest bin, with mean partner income below $\text{€}10,000$, lies outside the plotted range); bands are 95% confidence intervals. Right axis: the share of the bin predicted into the 30% bracket. Partner income is the household's taxable income at $r = -1$ net of the individual's own earnings; taxable wages carry the 10% deduction for professional expenses, so the measure is about nine tenths of the partner's wage less a tenth of own earnings, which is why the amounts look small. The validation region comprises the bins with mean partner income above $\text{€}11,000$ in which the predicted-treated share is below 5% (partner income up to about $\text{€}16,900$), the treatment region those in which it exceeds 95% (from about $\text{€}22,200$). The figure follows the graphical validation approach of Jakobsen and Sogaard (2022). In the validation region the event predicts no rate change. A gradient of income growth (the vertical distance between the two series) in partner income there would indicate a direct partner-income channel, violating Assumption 1. The corresponding figure for men is Figure E.1 in Appendix E; the participation counterpart is Figure E.2 there.

Assumption 2 (Exclusion restriction). $n_{ir}(0) = n_{ir}(1)$ with probability one.

Zero-rate income—including any direct effect of the event, through time use, home production, or transfers—is unaffected by the schedule change: the event occurs under both schedules and is left unrestricted, while the schedule change affects income only through the marginal net-of-tax rate. This is the potential-outcomes form of the no-income-effects property of the quasi-

linear model of Section 3.1.³⁰ Under quasi-linearity the assumption holds by construction.³¹ With income effects, the change in tax *liability* is a second channel from the schedule change to income: the fall in the quotient raises every household’s liability, and households predicted across a threshold pay more still, on top of the rate change that is the treatment. Assumption 2 then fails and the estimands below recover a total response to the schedule change; because the liability change works against the rate change when leisure is normal, the estimates are lower bounds on the compensated elasticity. Appendix D shows how to relax the assumption, and Figure D.1 there illustrates the approach.³²

Assumption 2, with the linear income function (5), also rules out responses to the position of the new threshold rather than to the rate: an individual predicted into the higher bracket who stays just below its threshold shows no change in the realised rate ($\Phi_{ir} = 0$) but a fall in income. Such responses enter the reduced form and not the first stage of the estimands below, and raise the cell-dose elasticity; the policy elasticity defined below (footnote 33), which divides by the predicted change, is unaffected.

Assumption 3 (Monotonicity). $Z_i [\text{NTR}_{ir}(1) - \text{NTR}_{ir}(0)] \geq 0$ with probability one, for all $r \geq 0$.

For every individual who is predicted a rate increase ($Z_i < 0$; Section 3.2), the schedule change may not lower the rate; for those who are predicted a cut, it may not raise it. Rates remain free to move either way for reasons unrelated to the event. This is the standard no-defiers condition ensuring that Wald estimands reflect positively weighted averages of individual parameters (Angrist & Imbens, 1995; Imbens & Angrist, 1994).

Section 3.2 explained why reform variation easily violates the condition and why our events make a violation less likely. Appendix H formalises the argument and illustrates it with a simulation.

³⁰Formally, decompose $\ln n_{ir}(g) = \mu_{ir} + \xi_{ir}(g)$ with $\xi_{ir}(0) \equiv 0$: μ_{ir} carries life-cycle and calendar-time growth and the event’s direct effects—changed time use, home production, transfers—which occur under both schedules and remain unrestricted; Assumption 2 sets $\xi_{ir}(1) = 0$. Reform settings impose the analogous restriction when they assume potential income unaffected by the reform.

³¹Quasi-linearity also implies that ϵ_i is the compensated elasticity, which coincides with the uncompensated one (Graber et al., 2026, Propositions 3–4).

³²The event-induced liability change scales with the amount of income shifted across the threshold and can be predicted from pre-event incomes in the same way as the rate change. Appendix D reports an attempt to add this simulated liability change to the cell regressions as a second regressor: within a cell it is largely collinear with the dose, so the design cannot separate the two channels.

For each cell, dose δ , and horizon, the reduced-form and first-stage contrasts compare realised income growth $\Delta y_{ir} \equiv \ln z_{ir} - \ln z_{i,-1}$ (the observed counterpart of $\Delta_r Y$) and net-of-tax rate growth $\Delta \text{NTR}_{ir} \equiv \ln \tau_{ir}^H - \ln \tau_{i,-1}^H$ across dose groups:

$$\begin{aligned} RF_c^r(\delta) &\equiv \mathbb{E}_c[\Delta y_{ir} \mid Z_i = \delta] - \mathbb{E}_c[\Delta y_{ir} \mid Z_i = 0], \\ FS_c^r(\delta) &\equiv \mathbb{E}_c[\Delta \text{NTR}_{ir} \mid Z_i = \delta] - \mathbb{E}_c[\Delta \text{NTR}_{ir} \mid Z_i = 0], \end{aligned}$$

and our headline estimand is the cell-dose elasticity³³

$$\beta_c^r(\delta) \equiv \frac{RF_c^r(\delta)}{FS_c^r(\delta)}.$$

Proposition 1 (Causal interpretation). *Under the mapping (4), the income function (5), and Assumptions 1–3, for every cell, dose $\delta \neq 0$, and horizon $r \geq 0$ with $FS_c^r(\delta) \neq 0$:*

$$\beta_c^r(\delta) = \mathbb{E}_c[\omega_{ir} \epsilon_i \mid Z_i = \delta], \quad \omega_{ir} \equiv \frac{\Phi_{ir}}{\mathbb{E}_c[\Phi_{ir} \mid Z_i = \delta]},$$

where the weights $\omega_{ir} \geq 0$ average to one: the cell-dose elasticity is a positively weighted average of individual elasticity parameters among the predicted treated.³⁴ (Proof: Appendix C.)

3.4 Aggregation

We aggregate in two steps.³⁵ Each estimate $\beta_c^r(\delta)$ rests on its own regression sample (Section 3.5): the dose’s predicted treated and the cell’s zero-dose group, observed at horizon r . We write $N_c^r(\delta)$ for its size and $\mathcal{D}_c$ for the within-cell support of non-zero doses. Within a cell, the dose-specific elasticities are combined with weights proportional to their regression samples,

$$\beta_c^r \equiv \sum_{\delta \in \mathcal{D}_c} w_c^r(\delta) \beta_c^r(\delta), \quad w_c^r(\delta) \equiv \frac{N_c^r(\delta)}{\sum_{\delta' \in \mathcal{D}_c} N_c^r(\delta')}.$$

³³Dividing the reduced form by the predicted change instead, $RF_c^r(\delta)/\delta$, scales the response by the predicted rather than the realised rate change—the policy elasticity of Hendren (2016). It needs only Assumption 1 and is unaffected where frictions or threshold-crossing responses attenuate the realised rate change.

³⁴Proposition 1 is the analogue, for event-based instruments, of the characterisation in Graber et al. (2026) of when reform-based ETI estimands are causal: their two conditions—a binary instrument contrast, and income controls rich enough to reproduce the conditional mean of the instrument—appear here as a feature of the design (each estimand contrasts $\mathbb{1}\{Z_i = \delta\}$ against the zero-dose group) and as Assumption 1.

³⁵The aggregation is stated for earnings and applies to participation in the same way (Section 3.5).

The zero-dose group belongs to the regression sample of every dose it is compared with, so it is counted once per dose. In the second step, we aggregate all cells contained in a group h of interest—own-income intervals for the income profiles, and own-by-partner-income rectangles for the maps of Section 4:

$$\epsilon_h^r \equiv \sum_{c \in h} v_h^r(c) \beta_c^r, \quad v_h^r(c) \equiv \frac{\sum_{\delta \in \mathcal{D}_c} N_c^r(\delta)}{\sum_{c' \in h} \sum_{\delta \in \mathcal{D}_{c'}} N_{c'}^r(\delta)},$$

so that $v_h^r(c)$ is cell c 's share of the group's estimation sample. Together, the two steps weight every estimate $\beta_c^r(\delta)$ of the group by its regression sample; we call these the frequency weights. Only estimates that pass the screens of Appendix F.2 enter the sums. We report ϵ_h^r as an event-time profile in Appendix F.5. The headline displays of Section 4 pool the post-event years $r = 0$ to 3 in the same way: every estimate enters with its own regression sample, so each year counts in proportion to the group's estimation sample in that year.

Because both steps use non-negative weights, the summary parameters and their pooled averages inherit the causal interpretation of Proposition 1. Substituting, ϵ_h^r is a weighted average of the individual parameters ϵ_i over the predicted treated of the group. An individual with dose δ in cell c has a weight proportional to $\omega_{ir} N_c^r(\delta)/n_c^r(\delta)$, where $n_c^r(\delta)$ is the number of the dose's predicted treated in its regression sample. The factor $N_c^r(\delta)/n_c^r(\delta)$ is one plus the ratio of the zero-dose group to the dose's predicted treated, so doses and cells with large comparison groups relative to their predicted treated count more than their share of the treated. Weighting by the predicted treated alone lowers the French-taxpayer aggregates of Section 4.2.1 to 0.18 (earnings) and 0.19 (participation),³⁶ with the difference concentrated in the cells of households predicted to cross from the 14 to the 30 percent bracket (Appendix F.2). The aggregation weights $v_h^r(c)$ and $w_c^r(\delta)$ are known sample frequencies rather than estimator-implied—in contrast to pooled regressions across treatment intensities, whose implicit weights can turn negative under heterogeneity (Callaway et al., 2026; Graber et al., 2026)—and al-

³⁶Weighted by the predicted treated alone, the French-taxpayer aggregates are 0.13 (men), 0.23 (women) and 0.18 (total) for earnings (0.17, 0.29 and 0.23 under the frequency weights; Table 4) and 0.15, 0.23 and 0.19 for participation (0.21, 0.28 and 0.25; Appendix Table F.4). From €10,000 of own work income up, they are 0.13, 0.09 and 0.11 for earnings (0.17, 0.23 and 0.20) and 0.10, 0.08 and 0.09 for participation (0.15, 0.18 and 0.16).

ternative policy-relevant weightings can be substituted without affecting the identification argument.³⁷

3.5 Estimation and inference

Estimation is by two-stage least squares within cells, separately for women and men, over an event window spanning $r = -3, -2$ and $r = 0, \dots, 3$, with $r = -1$ the base year. A regression is run for each cell, dose and period, a period being a year r relative to the event; its coefficient $\beta_c^r(\delta)$ is a cell–dose–period estimate. The regression sample of dose δ in period r pools the dose’s predicted treated with the cell’s zero-dose group, and we estimate the regression

$$(6) \quad \Delta y_{ir} = a_c^r(\delta) + \beta_c^r(\delta) \Delta \text{NTR}_{ir} + W_i' \gamma + u_{ir},$$

instrumenting the realised rate change ΔNTR_{ir} with Z_{ir} ; the controls W_i are fixed effects for own income at $r = -1, -2$ and -3 (in bins), age, marital status, the number of children at $r = -1$, and the event year.

For earnings, which we measure as work income without self-employment income, the outcome is log work-income growth. The coefficient is the earnings elasticity: the elasticity of own earnings with respect to the household net-of-tax rate, one minus the household’s marginal tax rate at the individual’s income, conditional on positive work income in both years, $r = -1$ and r .³⁸ The total (level) response of Section 4.2.1, estimated by Poisson pseudo-maximum likelihood (PPML) in levels including zeros and divided by the cell’s first stage, combines earnings and participation (Appendix Table F.5). For participation, $\Delta_r P_i \equiv P_{ir} - P_{i,-1}$ replaces Δy_{ir} in (6), estimated on the full event sample, which includes individuals without work income at $r = -1$. The coefficient is the participation semi-elasticity, the change in the participation probability per unit change in the log household net-of-marginal-tax rate. We

³⁷Inverse-variance weights are also non-negative and preserve the causal interpretation, but they weight estimates by their precision rather than by their sample frequency. The headline displays of Section 4 use the frequency weights throughout; Section 4.2.1 reports an inverse-variance-weighted aggregate for comparison.

³⁸This is the earnings elasticity of optimal-tax models, taken with respect to one minus the marginal rate and defined on the earnings of those who work rather than on hours (H. J. Kleven et al., 2009; Saez, 2001). The reform model of Section 5.2 uses it in the same sense (Bierbrauer et al., 2024).

divide it, cell by cell, by participation two years before the event to obtain the participation elasticity with respect to the household net-of-marginal-tax rate.³⁹

The estimation cells fix the own-income bin, the partner-income tercile within the bin and the pre-event bracket. The event year and the number of children, which the cells of Section 3.3 fix exactly, are pooled in estimation and enter through the fixed effects; filing status is fixed by the sample.⁴⁰

A dose is estimated if at least 30 predicted-treated individuals carry it in the cell and its regression sample, predicted treated and zero-dose group together, holds at least 300 outcomes observed at $r = -3$. In the lowest own-income bin, €0–10,000, the cells are in addition formed separately for individuals with at least €400 of work income at $r = -1$, the earner sub-sample, and for those below, the entrant sub-sample, each group with its own partner-income terciles, so that a cell compares individuals of the same pre-event earner status; the bin’s estimates pool the cells of both groups with the frequency weights and screens of every other bin (Appendix F.2). Estimating the reduced form and first stage on Z_i at pre-event horizons $r < 0$ —where $Z_{ir} = 0$, so (6) itself has no first stage—produces the placebo contrasts implied by Assumption 1, which we report alongside the post-event estimates.

The matching implementation replaces the regression comparison within each cell with nearest neighbours (with replacement) on the Mahalanobis distance of pre-event work income, measured at $r = -3$, $r = -2$ and $r = -1$, and age, with exact matching on the event year and a caliper of 0.8. The reduced form is then the matched average difference in income growth over all individuals of the cell whom the caliper keeps, predicted treated and zero-dose alike, rather than over the predicted treated to whom Proposition 1 refers; the first stage is the difference in realised rate changes between the two groups on the same sample, without matching; and the elasticity is their ratio (footnote 29).

Standard errors are clustered by person. Within a comparison group—a cell’s zero-dose group and the doses compared with it—the estimates for different doses and periods rest on the

³⁹The price of participation is the participation tax rate rather than the marginal rate (Appendix D.2); rescaled cell by cell by the change in the log net-of-participation-tax rate, the French-taxpayer aggregate is about 3.4 times larger, with a factor of about two in the bottom own-income bin and seven to fifteen above it (Appendix F.7).

⁴⁰The reduced form $RF_c^r(\delta)$ and first stage $FS_c^r(\delta)$ — Δy_{ir} and ΔNTR_{ir} regressed on Z_i with the same controls—are estimated alongside; the two-stage least squares coefficient equals their ratio, and without the controls the coefficients equal the group-mean contrasts of Section 3.3 divided by the dose δ —a factor common to both that cancels in the ratio.

same individuals, and their person-clustered covariance propagates to the aggregates β_c^r and ϵ_h^r , whose weights are sample frequencies with negligible estimation error; comparison groups contain different individuals and are treated as independent. At the median bin of the income profiles the clustered standard error is 1.39 times the conventional one, which treats every cell–dose–period estimate as independent, and at the median tile of the maps 1.45 times, mostly because an aggregate pools four post-event years of the same people.⁴¹ The pre-event placebo contrasts are pooled with the frequency weights of the estimate they accompany and tested with the same covariance. The estimates on the balanced sample (individuals observed in every year of the event window, $r = -3$ to 3), the matching estimates and the marriage estimates treat cells as independent in their standard errors, as do the placebo tests of the marriage exhibits; the exhibits say so. Standard errors of totals over both genders and of gender differences treat the women’s and the men’s estimates as independent, although the two samples contain the two spouses of the same tax units.

4 Results

This section proceeds from validation to estimates. We first describe the comparison cells that carry the estimation and assess the identifying assumptions in the aggregate, across all cells. The main results then summarise the estimates in three ways: as population aggregates (Section 4.2.1), along the own-income distribution, with their event-time profiles (Section 4.2.2), and over the joint distribution of own and partner income (Section 4.2.3). Section 4.3 examines robustness to specification choices, and Section 4.4 reports heterogeneity across household characteristics.

⁴¹Dose-specific comparisons within a cell share the zero-dose control group; the clustered covariance includes the induced correlation across dose-specific estimates, which is small. Across the bins of the income profiles, the covariance across doses accounts for a median 0.2 percent of the variance of the earnings estimate (0.3 percent for participation), the covariance across the periods of the same dose for 36 percent (52), and the estimates’ own variances for 64 percent (47).

4.1 Comparison cells and pooled validation

The cells cover parents in couples,⁴² aged 51 to 53 on average at the first-born’s fiscal exit, with own work income below €70,000; the estimates are local to them, and within them to those whom the fall in the quotient moves across a threshold. Appendix F.1 describes the sample, Appendix G the marriage design. A cell fixes an own-income bin (€10,000 wide; the example of Section 3.3 is a €2,000 slice of one), the partner-income tercile within the bin and the pre-event bracket; within it, each dose is compared with the cell’s zero-dose group (Section 3.5). In the headline specification—work income, all event years—297 cell–dose comparisons for men and 248 for women reach the estimation floors of Section 3.5, over ten pre-event brackets and 25 distinct doses for each gender;⁴³ 93 and 102 of them are in the €0–10,000 bin, where they are formed separately for the earner and entrant sub-samples (Appendix F.2). The comparisons differ widely in size: the density of taxpayers varies along the income distribution, and several bracket crossings exist only in the subperiods in which the schedule created them. Table 3 reports the size distribution at the level of the 21 grid cells of own-income bin $\times$ partner tercile, the finest level disclosed. Every estimate below is an average over cell–dose comparisons, each weighted by its regression sample (Section 3.4) and screened as in Appendix F.2.

How do the tests of Assumption 1 perform in the aggregate? Because every cell splits into a predicted-treated and a predicted-control group, the checks of the example cell pool across cells. The two groups differ where assignment requires it. The predicted treated have partners who earn more, by 0.24 (women) and 0.30 (men) of a standard deviation, and a household taxable income higher by about one standard deviation, which is what assigns them; the two groups also differ in family size, which sets the quotient. Own work income and age differ by 0.2 to 0.3 of a standard deviation, and pre-event income growth and participation by less than 0.05 (Appendix Table F.2). Figure 5 pools the validation-region check of Figure 4 across cells, measuring household income from each household’s hypothetical post-event threshold. Before the event, income growth shows no systematic gradient in household income across the

⁴²Couples whose civil status changes at the event are excluded, so that the fall in the quotient is the only change to the household’s schedule (Appendix F.1).

⁴³The brackets are the distinct marginal rates of the effective schedule (Section 2.1) at which the comparisons start, across the schedules in force from 2003 to 2023, whose rates and thresholds changed over the period. A dose need not lead to the next rate up: the fall in the quotient can carry a household past more than one kink of the schedule.

Table 3: The comparison grid

	Women	Men
<i>Event sample below the income cap</i>		
Individuals	1,884,448	1,758,056
Predicted treated (share)	0.30	0.32
<i>Comparisons of the headline estimation</i>		
Cell-dose comparisons	248	297
of which rate cuts	64	73
Pre-event brackets	10	10
Distinct doses	25	25
Grid cells with a comparison (of 21)	18	21
Distinct doses per grid cell (median)	13	13
Individuals in the comparisons	1,401,229	1,388,055
predicted treated (share)	0.39	0.38
Individuals per grid cell: median	78,295	60,858
smallest	391	821
largest	186,175	158,205
<i>Behind the displayed estimates</i>		
Cell-dose comparisons	128	152
Grid cells	18	21
Individuals	1,089,165	1,213,681
predicted treated (share)	0.37	0.35

Notes: The upper panel counts the child’s fiscal exit sample below the income cap (Appendix F.1) and its predicted treated. The lower panels count the cell-dose comparisons of the headline estimation and the individuals in them at $r = -1$: the predicted treated of every comparison and the zero-dose comparison group of every cell, each counted once. A grid cell is an own-income bin (€10,000 wide) $\times$ within-bin partner tercile; in the €0–10,000 bin the comparisons are formed separately for the earner and entrant sub-samples (Appendix F.2), and a grid cell pools the two sub-samples’ terciles of the same rank. The comparison cells of Section 3.5 additionally fix the pre-event bracket; that finer split is not disclosed at cell level. Behind the displayed estimates: comparisons with at least one estimate that passes the screens of Appendix F.2 for earnings or participation.

validation region, nor on the other side of the threshold, for either gender (in every €2,000 band of household income, women’s work income grows by 1.7 to 2.0 log points from $r = -2$ to $r = -1$, men’s by 0.1 to 0.4). This is what the figure checks: a partner- or household-income channel of the kind Assumption 1 excludes would already show such a gradient. After the event the two genders move in opposite directions. Men, about two years older than women on average, see their work income fall, by 2.1 to 2.4 log points from $r = -1$ to $r = +1$ in every band, a little less in the validation region than above the threshold. Women’s work income still grows, and more in the validation region (0.9 to 1.5 log points by band) than above the threshold (0.5 or less).⁴⁴ Single bands are noisy, since the figure pools different

⁴⁴Just below the threshold women’s growth is lower than a few thousand euros further below (0.9 against 1.5 log points). A possible reading is that predicted controls within a few thousand euros of the post-event threshold are partly treated, by income drift across it after the event; the control arm’s realised rate change is not zero (Table F.2).




Figure 5: The validation-region check, pooled across cells

Notes: Log work-income growth relative to $r = -1$ (100 times the change in log work income, in log points) at $r = -2$ (before the event) and $r = +1$ (after), against household taxable income at $r = -1$ minus the household's statutory threshold at post-event fiscal parts, pooled across the comparison cells of the design. An individual enters once per comparison it belongs to (zero-dose controls once per dose comparison of their cell) and once per statutory threshold within €10,000 of its household taxable income, on whichever side its income lies. Points are means of 40 quantile bins of these rows, cut over the whole window and joined within each side of the threshold; the first bin right of the threshold (grey) can therefore also hold rows just below it. Growth is computed for individuals with work income above €1,000 in both years and winsorised at the 5th and 95th percentiles of the gender sample at each horizon. Bands are 95% confidence intervals clustered by person. The figure is the pooled analogue of the validation-region check of Figure 4: a pre-event profile that is flat across the threshold indicates that household income carries no income-growth gradient of its own before the event (Jakobsen & Sogaard, 2022). Sample: the child's fiscal exit sample of Appendix F.1, in the comparison cells of the design (Section 3.5); 874,000 (women) and 934,000 (men) rows at $r = -2$ and 864,000 and 908,000 at $r = +1$, 19,900 to 24,500 per bin.

comparison cells along its axis. Because the outcome is winsorised, which caps the large falls that carry the response, the figure checks the validation region and does not measure the response; Section 4.2 estimates it within cells. Sorting of the event's timing, and bunching below the post-event threshold, would show as a break in the density of household taxable income there; pooled across cells, the estimated log-density discontinuities are -0.002 (s.e. 0.010) for women and -0.001 (s.e. 0.009) for men (Appendix Figure F.1). The pre-event placebo contrasts of Section 3.5 accompany every estimate below, flagged where they reject (Appendix F.2).

4.2 Main results

The main results proceed from the most aggregated summary to progressively finer ones: population-level averages, the profile along the own-income distribution, and the map over the joint distribution of own and partner income.

4.2.1 Aggregated elasticities

What do our estimates imply for the aggregate elasticity, under the most common aggregation schemes applied in the literature? Table 4 reports four, by gender and combined; the participation counterpart is Table F.4 in Appendix F.6.

Table 4: Aggregated earnings elasticities

	Men	Women	Total
Population-weighted: French taxpayers	0.169 (0.035)	0.289 (0.051)	0.229 (0.031)
Population-weighted: event population	0.174 (0.033)	0.322 (0.048)	0.246 (0.029)
Inverse-variance weighted	0.125 (0.023)	0.102 (0.022)	0.114 (0.016)
Income-weighted	0.137 (0.031)	0.217 (0.037)	0.171 (0.024)

Notes: Aggregation of tile-level estimates: cell-dose elasticities pooled within own-by-partner-income rectangles under the profile-level display conventions of Appendix F.2. One bin-level screen applies: a bin contributes only if the cells passing the cell screens cover at least 10 percent of the bin’s estimation sample and at least 1,000 individuals underlie them; this removes the men’s €0–10,000 bin for earnings (one cell, 5.7 percent of the bin’s estimation sample) and, by the floor, the women’s €60,000–70,000 bin for earnings and participation (865 and 885 individuals). The €0–10,000 bin pools the earner and entrant sub-samples (Appendix F.2). *French taxpayers:* each own-income bin weighted by its share of prime-age taxpayers, from the INSEE distribution of annual salary income by gender (2021), with the non-employed in the bottom bin; the support of the estimates, own income up to €70,000 in 2015 prices, covers at least 95.6 percent of the men and 97.0 percent of the women filing a return in 2023. The averaged estimates remain those of parents aged around 50 at the event, older than the taxpayers who supply the weights. *Event population:* each own-income bin $\times$ partner-tercile group is weighted by its regression samples in the event year ($r = 0$), summed over its doses, with the zero-dose group counted once per dose as in Section 3.4. *Income-weighted:* each bin weighted by its share of aggregate earnings from the same INSEE distribution. In the French-taxpayer and income-weighted rows, a bin’s weight is split equally over the partner terciles whose cells pass the screens. The Total column combines genders with equal population shares (French-taxpayer and inverse-variance rows), the genders’ shares of the event-population counts (event row), and gender aggregate-income shares (income row). Standard errors in parentheses are clustered by person, and the inverse-variance weights use the same clustered variances; the standard errors of the Total column treat the two genders as independent.

Weighted to the population of French taxpayers, the average earnings elasticity, the elasticity of own earnings with respect to the household net-of-tax rate, is 0.229 (0.289 for women, 0.169 for men).⁴⁵ Weighted to the event population itself—parents at the fiscal exit of their first-born, each bin and partner tercile weighted by its regression samples in the event year—the average is 0.246 (0.322 for women, 0.174 for men). The two aggregates are close: reweighting from the event population to the general taxpayer distribution moves the total by less than 0.02, so the summary is not an artefact of the event population’s income composition. The difference between women and men under the French-taxpayer weights, 0.12, is 1.9 person-clustered standard errors (95% confidence interval from about zero to 0.24, treating the two genders as independent).

⁴⁵The French-taxpayer weights are constructed from the INSEE distribution of annual salary income by gender (2021), interpolated to a piecewise-linear CDF with a Pareto tail, with the non-employed entering the bottom income bin; see the table notes.

The literature has also aggregated estimates with inverse-variance (efficiency) weights. That these produce the smallest average (0.114) is itself informative: the cells with high elasticities are the less precisely estimated ones, so precision weighting shifts mass towards mid-income cells where responses are modest.

The convention in the taxable-income literature, however, is to weight each individual by their contribution to the tax base. This yields an earnings elasticity of 0.171 (0.024) overall—0.137 for men and 0.217 for women—below the population-weighted averages. The shortfall follows from the concentration of responses at low incomes for both genders, documented in the next subsection: weighting by income moves mass away from the most responsive part of the distribution.

We treat the French-taxpayer weighting as the headline aggregate; across the four schemes, the average earnings elasticity lies between 0.11 and 0.25. The participation elasticities are of comparable size: weighted to French taxpayers, 0.247 overall—0.283 for women, 0.212 for men (Appendix Table F.4). For both earnings and participation, the women’s aggregate lies above the men’s under every scheme of Table 4 except inverse-variance weighting (with treated weights, from €10,000 of own income up, the ordering reverses; footnote 36).

Earnings and participation combine in the response of mean work income, which we report as the total (level) response: the reduced form of each cell estimated in levels, including zeros, by Poisson pseudo-maximum likelihood (PPML) and divided by the cell’s first stage, which puts it on the scale of the elasticities above. The earnings elasticity conditions on positive work income in the year before the event and in the outcome year; the total response does not. Weighted to French taxpayers from €10,000 of own income up, it is 0.21 (0.03) for men and 0.20 (0.03) for women, against 0.32 and 0.41 for the sum of the earnings and participation elasticities: a third (men) to a half (women) smaller.⁴⁶ For women the total response is at or below the earnings elasticity alone in every bin from €10,000 up that rests on more than one cell, so in euros their participation response adds little to the change in mean income. In the bottom bin the two measures describe different objects (for women, 0.08 against a sum of 0.78): the earnings elasticity averages log changes of small incomes that carry few euros.

⁴⁶The two elasticities overlap: a fall to a few hundred euros counts both as an exit and as a large fall in earnings (Appendix Table F.13, *Floor 400*), and exits from small earnings carry few euros.

Appendix Table F.5 reports the earnings and participation elasticities, their sum and the total response by own-income bin.

Our income-weighted earnings elasticity of 0.171, with respect to the household net-of-marginal-tax rate, is within the range of quasi-experimental estimates of the elasticity of taxable income surveyed by Neisser (2021).

4.2.2 Profiles along own income

The population averages conceal systematic variation in responsiveness, first and foremost along own income.

Figure 6 reports elasticities aggregated within own-income bins, pooled over the post-event years $r = 0$ to 3. For women, shown in Panel (b), elasticities are largest at the bottom of the distribution, at 0.45 in the lowest bin, and lower everywhere above it: 0.18 and 0.02 in the €10,000–20,000 and €30,000–40,000 bins and 0.19–0.31 in the other bins up to €60,000; the top bin rests on a single cell–dose–period estimate (−0.28) and is drawn in grey. The €10,000–20,000 bin fails its pre-event placebo and is displayed hollow; its depression relative to the neighbouring bins partly reflects the flagged differential trend rather than behaviour. For men, shown in Panel (a), the profile is highest in the €10,000–20,000 bin, at 0.32, among the bins resting on more than one estimate, and lower in every bin above it: 0.17 and 0.13 through the middle of the distribution and 0.02–0.19 in the three top bins; the lowest bin rests on a single cell–dose–period estimate, drawn in grey (it enters no aggregate; its placebo and later-year values are in the figure note).

Event-time profiles. The estimates by year relative to the event serve as a validity check, since the same construction yields placebo estimates for the pre-event years; they are collected in Appendix F.5. Earnings and participation differ sharply in their dynamics. The earnings elasticity is an impulse: largest in the event year, at 0.37 for men and 0.36 for women, it falls to 0.14 and 0.30 at $r = 1$ and is statistically indistinguishable from zero from $r = 2$. The participation elasticity is more persistent: it appears immediately and remains between 0.15 and 0.23 throughout the window for women, while for men it declines from 0.20 in the event year to 0.08 at $r = 3$. The headline earnings estimates pool $r = 0$ to 3 and so lie between the




Figure 6: Earnings elasticities along the own-income distribution

Notes: Elasticities are largest at low own incomes for both genders. In the raw data, falls in work income of more than one log unit after the event are 0.2 to 1.7 percentage points more frequent among the predicted treated than in their comparison groups in every bin from €20,000 up with more than one estimate, so part of the response there may reflect a small share of large falls. Sample: the child's fiscal exit sample of Appendix F.1, all event years. Outcome: work-income growth relative to $r = -1$, pooled over $r = 0$ to 3. Construction, screens and weights follow Appendix F.2, with own-income bins as the aggregation groups. The €0–10,000 bin pools the earner and entrant sub-samples (Appendix F.2): 18 cell-dose-period estimates on 159,108 individuals for women and a single estimate on 4,421 individuals for men, all from the earner sub-sample. The men's single estimate has a robust first-stage F of 615 and a classical standard error of 0.49, the value the cap of 0.5 is applied to; its placebo estimates at $r = -3$ and -2 are 0.00 and -0.03 , and the same comparison gives 0.22, 0.27 and 0.33 at $r = 1$ to 3. A hollow marker flags a bin that fails its pooled pre-event placebo ($|t| \geq 1.96$; here only the women's €10,000–20,000 bin); flagged bins are kept (Appendix F.2). Points are at bin midpoints of own work income at $r = -1$; shaded areas are 95% confidence intervals from person-clustered standard errors. The vertical axis is set by the bins with at least two cell-dose-period estimates and is shared by the two genders. Bins resting on a single estimate are drawn in grey with their interval cut at the axis: the men's €0–10,000 bin, 0.45 with interval -0.50 to 1.39 , and the women's €60,000–70,000 bin, -0.28 with interval -0.90 to 0.35 ; neither enters an aggregate (Appendix F.2). The participation counterpart is Figure F.4.

event-year response and the smaller later-year responses; for participation, pooling matters less, and its pre-event placebo estimates are indistinguishable from zero for both genders. For earnings, the women's placebo at $r = -3$ is negative (-0.10 , $t = -2.65$), the same differential drift flagged in the €10,000–20,000 bin of Figure 6, and so is the men's at $r = -2$ (-0.05 , $t = -2.00$).⁴⁷ Pooled over all pre-event cells with the weights of the estimates, the placebo reduced form does not reject at the 5 percent level for men, for earnings or participation, or for women's participation; for women's earnings the test is borderline ($t = -1.9588$, $p = 0.050$; $t = -2.30$ if the pre-event cells are weighted by inverse variance instead). The event-year estimate combines the immediate behavioural response with any income realisation coinciding with the timing of the fiscal exit itself, so the impulse shape of the earnings response warrants caution in reading its event-year peak as a steady-state elasticity.

⁴⁷Pre-event points are scaled up relative to post-event ones: they divide the placebo reduced form by the pre-event contrast in realised rates, 0.10 to 0.13 against 0.43 to 0.45 in the event year; on the event-year scale the two placebos are -0.03 , 8 percent of the women's event-year response, and -0.02 .

Relation to the literature. The larger responses at low own incomes mirror the concentration of participation and hours responses at low earnings documented in the labour-supply literature.

4.2.3 Maps over own and partner income




Figure 7: Earnings and participation elasticities by own and partner income

Notes: The double-low corner, the lowest own-income bin with a low-earning partner, carries the largest women's earnings tile with a passing pre-event placebo (0.63); the one larger coloured tile (0.72) fails it. Sample: the child's fiscal exit sample of Appendix F.1, all event years. Outcome: work-income growth relative to $r = -1$ in Panels (a) and (b), an indicator for positive work income in Panels (c) and (d), pooled over $r = 0$ to 3; the estimates are elasticities with respect to the household net-of-marginal-tax rate, those for participation divided by participation two years before the event. Each tile pools the cell-dose elasticities of its own-by-partner-income rectangle, with the pooled standard error beneath and, for a dashed cell, the number of underlying individuals above. Standard errors are clustered by person, except those of the bracketed (screened) values, which are classical. Colour classes, the four display states and the flag rule are those of Appendix F.2: a coloured tile passes the screens; a dashed cell with a bracketed value has no estimate within the cap of 0.4, or its pooled estimate fails the cap or the floor of 1,000 individuals; n.a. marks a rectangle none of whose cells passes the screens other than the cap; a blank rectangle has no cell; corner ticks and a grey value mark a failed pre-event placebo; the estimate is kept. A grey fill marks an estimate below -0.075 . The €0–10,000 column pools the earner and entrant sub-samples (Appendix F.2). In the men's low and middle participation tiles the two sub-samples carry opposite signs (0.47 and 0.31 for earners against -0.37 and -0.26 for entrants), and the pooled 0.02 and -0.15 are within two standard errors of zero; neither tile is read. Both incomes are measured at $r = -1$; partner terciles are formed within each own-income bin, so a row is not at constant partner income. The grid ends at €60,000, the own-income range supported for both genders. Tercile boundaries, estimates, standard errors and cell counts for every tile are reported in Appendix Tables F.7 and F.8.

The own-income profile uses one of the design’s two dimensions. Within an own-income bin, the identifying variation comes from partner income (Section 3.2), and it is available along the whole partner-income distribution. Because the comparisons do not need wide partner-income intervals, we can divide the partner-income distribution and study heterogeneity along the source of identification itself. This permits asking which households drive the high elasticities documented at low own incomes: those insured by a high-earning partner, or those in which both partners earn little.⁴⁸ Figure 7 disaggregates the own-income profile accordingly, reporting one estimate per own-income bin and within-bin partner-income tercile; the vertical axis orders partner income *within* each own-income bin, so the comparison across terciles holds the own-income composition fixed.

For women, the answer differs between earnings and participation. For earnings (Panel b), the large bottom-bin response concentrates among households in which the partner also earns little: the low-partner tercile of the bottom bin, the double-low corner of the map, carries an elasticity of 0.63, against 0.29 in the middle and -0.02 in the top tercile, and no other tile of the women’s map whose placebo passes is as large (the €10,000–20,000 middle-tercile tile, 0.72, fails its placebo). For participation (Panel d) the pattern reverses: the bottom-bin elasticity is close to zero where the partner also earns little (0.02) and large where the partner earns more (0.69 in the middle and 0.35 in the top tercile). The participation response of women with a higher-earning partner is not confined to the event year: in the lowest bin, the middle- and top-tercile tiles are 0.80 and 0.44 over the three following years, against 0.48 and 0.11 in the event year.⁴⁹ Women’s and men’s bottom-bin participation elasticities are of similar size (0.33 and 0.29; Appendix Figure F.4) but come from different couples: men respond in couples that paid no income tax before the event (0.95 against 0.10 in couples already taxed; 7 and 20 cell–dose–period estimates), women in couples already taxed (0.73 against -0.01 ; 39 and 28). Men are on average two years older at the event (Appendix Table F.1), and the retirement margin matters more for them (Section 4.3). Women with a higher-earning partner adjust participation; women in households where both earn little adjust earnings while largely remaining in work. The double-low corner is, however, where the design has least support.

⁴⁸In the lowest own-income bin, the low partner-income tercile ends at €24,700 of partner income for women and at €13,800 for men in the earner sub-sample; Appendix F.2 gives the boundaries of both sub-samples.

⁴⁹Each value pools the tile’s cell–dose–period estimates at $r = 0$, or at $r = 1$ to 3, with the screens and weights of Figure 7; the person-clustered standard errors are 0.10 and 0.07 in the event year and 0.14 and 0.09 over the three following years.

For men, the earnings map (Panel a) shows no clear partner-income gradient within own-income bins; the largest tile is 0.32, at €10,000–20,000 with a low-earning partner, and the €0–10,000 column is screened.⁵⁰ On the participation map (Panel c), responses are sizeable across most of the map. The largest (0.42) is in the top partner tercile of the bottom bin, whose two other tiles pool responses of opposite sign and are not read (see the figure notes). In the bottom bin the two sub-samples measure different responses: the cells of the entrant sub-sample can record only entry into work, those of the earner sub-sample mostly exit from it. Over the bin as a whole, predicted-treated entrants enter work less often than their comparison groups, and predicted-treated earner men leave work more often; the entrant cells carry 75 percent of the weight of the men’s bottom bin (Appendix F.7).

Appendix F.9 shows the equivalent income profiles for an alternative outcome definition: broad earnings, which add self-employment income to work income. From €10,000 of own work income up, the broad earnings elasticities closely track the earnings elasticities.

Is the map robust to the sample period and to the event definition itself? Section 4.3 shows that the map survives dropping the 2020–2023 event years, which the headline specification includes (Appendix F.10). For the jump from the 14 to the 30 percent bracket, the marriage design (Appendix G) gives earnings elasticities similar to those at the child’s fiscal exit for men and somewhat lower for women. In the €10,000–20,000 bin it gives 0.32 for men and 0.14 for women, against 0.32 and 0.18 at the child’s fiscal exit (Figure 6, all bracket changes); in the €20,000–30,000 bin it gives 0.16 and 0.18, against 0.17 and 0.30. Above €30,000 no bracket jump of the marriage design gives a number that is stable across weightings.

Four readings are consistent with the maps. (i) Partner income as insurance: a lower earner whose partner earns considerably more can afford to leave work, and where both earn little, earnings adjust while employment is kept. (ii) Hours may be more adjustable at low pay, where part-time work is common, so that earnings can be adjusted without leaving work more easily at the bottom of the distribution than above it. (iii) The event arrives late in working life, and much of the response from €10,000 up comes from those approaching retirement ages: below age 59 at the event, the women’s earnings profile shifts down with its shape intact and the men’s lies below 0.10 from €10,000 to €60,000 (Appendix Figure F.8), so the timing of the

⁵⁰No men’s earnings cell in that column is within the maps’ standard-error cap of 0.4, so its three tiles show bracketed values. The single cell of the men’s lowest bin in Figure 6 passes the profiles’ cap of 0.5.

exit from work may be the margin that moves. (iv) Part of the gender gap may be the partner’s adjustment transmitted within the household rather than a difference in individual parameters (Appendix B); since both spouses’ responses are estimated within the same cells, the pair of estimates is informative about such reallocation. Under all four, responsiveness depends on the household’s circumstances—the partner’s income, the jobs available, the distance to retirement—rather than on preferences alone, in line with Attanasio et al. (2018), who find that women’s labour-supply elasticities vary with age, wealth, hours and wages.

In sum, responsiveness is largest at low own incomes for both genders. The participation responses of the lowest-earning women are driven mostly by those with higher-earning partners rather than by those whose partners also earn little; their earnings responses, in contrast, are largest where both partners earn little. The heterogeneity around the aggregates is large; Section 5 examines its implications for tax design.

4.3 Robustness

The earnings map of Section 4.2.3 is robust to the estimation variants of Table 5; the exception is nearest-neighbour matching, which does not reproduce the men’s map. The table re-estimates the map under each variant, on the same sample and grid, and reports the correlation of tile values with the headline map, the mean absolute tile difference and the income gradient. The balanced sample and the sample without the 2020–2023 event years reproduce the map: correlations above 0.8, gradients within 0.03 of the headline’s. Matching, which replaces the regression comparison within cells (Section 3.5), reproduces the women’s map but not the men’s; the matched gradient has the headline’s negative sign for both genders. PPML estimates measure a different object, the response of the level of work income with zeros included, and are reported with the aggregates as the total response (Section 4.2.1).

From €10,000 up, the responses come mostly from parents close to leaving work. Because the fiscal exit of a first-born typically comes as the parents approach retirement, Appendix Table F.13 repeats the estimates in restricted samples. Among individuals younger than 59 at the event, the aggregate earnings elasticity falls from 0.17 to 0.06 for men and from 0.23 to 0.10 for women (Appendix Figure F.8). Excluding instead the few percent who already

Table 5: Specification agreement of the heterogeneity map

	Men				Women			
	Corr.	Rank corr.	mAD	Gradient	Corr.	Rank corr.	mAD	Gradient
Headline	1.00	1.00	0.00	-0.07 (0.03)	1.00	1.00	0.00	-0.07 (0.03)
Balanced sample	0.84	0.91	0.04	-0.06 (0.02) [†]	0.98	0.94	0.05	-0.10 (0.03) [†]
Nearest-neighbour matching	0.27	0.24	0.12	-0.03 (0.02) [†]	0.70	0.82	0.16	-0.09 (0.03) [†]
Without 2020–2023 event years	0.82	0.92	0.04	-0.06 (0.03)	0.98	0.97	0.05	-0.09 (0.04)

Notes: Each row re-estimates the earnings map of Figure 7 under one estimation variant, with the map’s grid and display conventions (Appendix F.2; tiles flagged by the pre-event placebo included). Corr. and Rank corr. are the Pearson and Spearman correlations of tile values with the headline map, and mAD is the mean absolute tile difference, all on the tiles both maps keep (own-income bin $\times$ partner tercile). Gradient is the slope of tile values in own income per €10,000, each tile weighted by its estimation sample, with its standard error in parentheses, clustered by person; standard errors marked [†] treat cells as independent. From €10,000 up the women’s correlations are 0.99 (balanced sample), 0.72 (matching) and 0.99 (without 2020–2023). The women’s matching gradient of -0.09 comes from the bottom-bin tiles; from €10,000 up it is -0.02 , against -0.04 for the headline. The first-stage screen uses the robust t -ratio, except in the balanced sample (classical t -ratio) and under matching (no first-stage screen).

draw some pension income in the year before the event lowers it to 0.01 for men and 0.06 for women, and reduces the combined participation aggregate.

Weak instruments. The displayed estimates and their confidence intervals can be read without a weak-instrument correction: a display screen keeps the weakest cells out, and two checks show that the cells it admits need none. A small predicted-treated group can leave a cell’s just-identified first stage weak, so a cell enters an exhibit only if its heteroskedasticity-robust first-stage F , with one instrument the effective F of Montiel Olea and Pflueger (2013), is at least 9 (Appendix F.2). We use the robust statistic because the first-stage residual variance differs between the predicted treated and the zero-dose group; at the frequency-weighted median displayed cell it is about 0.8 times the classical F . Weak cells carry almost no weight: those with F below 23.1, the Montiel Olea and Pflueger (2013) critical value for a worst-case bias of 10 percent, hold 0.05 percent of the earnings and 1.4 percent of the participation weight (Appendix Table F.9). Weak-instrument-robust inference agrees with the cells’ robust Wald intervals: robust Anderson–Rubin sets (Anderson & Rubin, 1949) are bounded in every displayed cell with a robust first stage, at the median 0.4 and 1.0 percent longer, and reject zero wherever those do (Appendix Table F.12); the tF correction of Lee et al. (2022), applied to each cell and carried to the bin standard errors as an approximation, costs no income-profile bin its significance at 5 percent (Appendix Table F.11).

A stricter screen tests whether the aggregates lean on the weakest displayed cells. Requiring $F \geq 104.7$, above which the tF critical value is 1.96 and the two-stage least squares t -ratio

needs no correction, moves every income-profile bin by less than one standard error except women’s participation at €10,000–30,000, whose two bins fall by 0.09 and 0.08, about two standard errors (Appendix Table F.10). Those falls are a matter of weighting, not of weak-instrument bias. They are driven by seven cells of households at the 14 percent rate predicted to cross into the 30 percent bracket. These cells are not weak by the conventional criteria: their robust F of 28 to 97 clears 23.1, and their Anderson–Rubin sets exclude zero in five of the seven; with 77 to 403 predicted treated, they fall short only of 104.7. What the stricter screen changes is the weighting: under the frequency weights of Section 3.4 the seven cells hold 7 and 11 percent of the two bins’ weight but 0.2 percent of their predicted treated, and with predicted-treated weights the same screen moves the bins by -0.01 and 0.00 .

4.4 Heterogeneity




Figure 8: Elasticities by household subgroup

Notes: Each row pools the cells of one subgroup with own income from €10,000 up, both genders together, weighting each estimate by its regression sample (Section 3.4); the screens are those of Appendix F.2. Estimates are elasticities with respect to the household net-of-marginal-tax rate; those for participation are divided by participation two years before the event. Whiskers are 95 percent confidence intervals, with standard errors clustered by person within gender and the two genders treated as independent. A hollow marker means that the row fails its pre-event placebo at the 5 percent level, tested with the estimate’s weights and covariance. A filled row can still hide one gender that fails its own placebo: women in the All and Married rows for earnings and in the Younger siblings remain row for participation (Appendix Figure F.10 tests each gender on its own). The All row (0.19 for earnings, 0.15 for participation) is not weighted to French taxpayers; the French-taxpayer aggregates from €10,000 up are 0.20 and 0.16. Subgroups are defined at $r = -1$, except the child’s age, measured at the event; “First and only child” means that no other child is ever observed in the tax unit. The civil-union (PACS) row rests on 20 cells for earnings and 57 for participation (358 and 681 for All), and both its intervals include zero.

The design permits heterogeneity cuts beyond own income, partner income and gender. Figure 8 reports pooled elasticities by household subgroup, both genders' cells pooled and own income from €10,000 up: the age gap between the parents, marriage versus civil union (PACS), whether the exiting first-born is the household's only child, and the age of the exiting child. Responses are larger when the age gap between the parents is four or more years (0.19 against 0.11 for earnings, 0.17 against 0.12 for participation), and larger when the exiting child is the first and only one (0.18 against 0.11, and 0.23 against 0.10 for participation)—in line with the first fiscal exit also being the last for these households. Both contrasts come with a caveat: the four-or-more-years row fails its pooled pre-event placebo for earnings and participation, and the first-and-only-child row for earnings (hollow markers). The age-gap contrast is the men's: within men the earnings elasticity is 0.26 with a gap of four or more years against 0.09 below it, within women 0.12 against 0.16, and the cut moves the father's age at the event (53.4 against 51.2 years), not the mother's. The children contrast has the same sign for both genders and is largest for women's participation (0.30 against 0.07). By the age of the exiting child, households whose child is 20 or younger respond least (0.10 for earnings and 0.11 for participation, against 0.15–0.19 and 0.14–0.18 in the two older groups); above that age the ordering differs by gender. Appendix Figure F.10 reports the by-gender breakdown.

5 On the relevance of accounting for heterogeneity

Section 4 documented substantial heterogeneity in estimated elasticities, both across the income distribution and by the composition of income within the household. In this section, we examine whether this heterogeneity matters for tax design. The two exercises below use the estimated earnings and participation elasticities, both with respect to the household net-of-marginal-tax rate, pooled over the event year and the three years after it ($r = 0$ to 3).⁵¹ The income distributions they run on are calibrated from aggregate statistics, so the analyses remain illustrative. The elasticities are those of parents in couples at the fiscal exit of a child, applied here to the whole distribution.

⁵¹The earnings response is short-lived: it is largest in the event year and statistically indistinguishable from zero from $r = 2$. The pooled earnings elasticities of this section average large early responses and near-zero later ones. The participation response persists for women and declines for men (Section 4.2.2).

We assess this question against two methodologically distinct benchmarks. The first is the normative optimum: the Mirrlees–Saez tradition solves for the tax schedule that maximises a social welfare function subject to resource and incentive constraints (Mirrlees, 1971; Saez, 2001), and its counterpart for couples solves the corresponding multidimensional screening problem (Golosov & Krasikov, 2025; H. J. Kleven et al., 2009). The second is the evaluation of reforms: the perturbation approach of Bierbrauer et al. (2024) takes the observed tax system as given and derives sufficient statistics under which a move towards more individualised taxation of couples is Pareto- or welfare-improving. Section 5.1 conducts the first exercise using our estimates of tax elasticities by own income; Section 5.2 conducts the second using our results on heterogeneity in tax responsiveness by the division of income within couples.

5.1 In optimal tax models

The optimality conditions of the Mirrlees–Saez framework are local: the optimal marginal rate at earnings level z is governed by the elasticity *at* z (Mirrlees, 1971; Saez, 2001). In the framework, a government chooses a tax schedule to maximise a social welfare function subject to its resource constraint and to incentive compatibility; because ability is unobserved, the schedule can condition on earnings but not directly on ability. Saez (2001) expresses the resulting optimality conditions in terms of empirically meaningful sufficient statistics: behavioural elasticities, the earnings distribution, and social marginal welfare weights. The optimal marginal rate at z balances the mechanical revenue and welfare effects of a small rate change against the behavioural response of taxpayers at z ; nothing in the formula requires the elasticity to be common across the distribution, and the constant-elasticity calibrations used in practice are a convenience rather than a feature of the theory.⁵²

We adapt the Mirrlees–Saez calculation of Brewer et al. (2010) to a skill distribution inferred from the 2023 administrative work-income aggregates and compare a constant elasticity with the estimated profile $e(z)$.⁵³ Figure 9 provides an illustrative comparison. The grey dashed

⁵²Following Brewer et al. (2010), the calibration infers a skill distribution from observed earnings under a flat reference tax rate and holds that distribution fixed while solving for earnings and taxes under each elasticity schedule. We evaluate the estimated cross-sectional $e(z)$ pointwise as earnings change; this is an illustrative approximation rather than an estimate of a complete structural labour-supply model. The profile is estimated at the rates of the existing schedule and could itself shift with rates if local responses depend on them.

⁵³The figure interpolates the estimated income-group elasticities and pools genders equally at shared income-bin midpoints. Pre-trend-flagged bins remain in the calibration to avoid selecting on pre-period noise. Where one gender’s estimate in an edge bin enters no aggregate, the other gender’s estimate represents that bin.




Figure 9: Optimal marginal tax rates under constant and heterogeneous elasticities: 2023 administrative work-income distribution

Notes: The figure plots an illustrative Mirrlees–Saez calculation following the skill-inversion approach of Brewer et al. (2010). The earnings distribution is reconstructed from 2023 administrative gender-specific counts, participation shares, and mean work incomes in nominal €10,000 bins. Multiplying each all-filer count by its positive-earner share and dividing the all-filer count by that share yields the positive-earner count and conditional mean. A monotone-cubic quantile curve matches those bin counts and means. Above nominal €70,000, the administrative total fixes the tail mass and an INSEE-shaped Pareto tail supplies its within-tail distribution; the tail mean is not observed in this export. Nominal 2023 incomes are divided by 1.1747, so the x-axis and all monetary amounts are in 2015-price euros. The original nominal-income bins are retained in the reconstruction; their exact allocation across 2015-price €10,000 bins is not observed. Men and women receive equal distribution weights. The grey dashed line uses a constant elasticity of 0.181, the positive-earner-weighted headline intensive-margin benchmark from the administrative population aggregation. The solid navy line uses the gender-pooled estimates by own-income-bin midpoint, retaining pre-trend-flagged bins. Where one gender’s estimate in an edge bin enters no aggregate (men below €10,000, women at €60,000–€70,000), the other gender’s estimate represents that bin. A flat 41% reference marginal rate maps the fitted earnings distribution into a skill distribution, which is held fixed while earnings respond to each counterfactual tax schedule. The calculation uses logarithmic social welfare ($\gamma = 1$) and a revenue requirement of €3,405 per earner in 2015 prices; welfare weights and the balancing transfer are recomputed for each schedule. The seven displayed points are probability-weighted mean tax rates in the deflated nominal-income bins; shape-preserving interpolation between them is a visual guide and does not enter the calculation. The figure stops near €60,000, the 2015-price equivalent of the original nominal €70,000 bin limit; the estimated elasticity profile extends to €70,000 in 2015 prices.

line assumes the positive-earner-weighted headline benchmark elasticity of 0.181; the navy line uses the profile at the corresponding 2015-price income. The inferred skill distribution is held fixed, while earnings, welfare weights, and the revenue-balancing transfer are recomputed for each schedule. High estimated elasticity at low earnings lowers the estimated-profile rate by about 17.4 percentage points on average below €10,000 in 2015 prices; the transfer falls from €10,321 to €9,670 in the same units. The estimated-profile rate exceeds the constant-elasticity rate most in the displayed bin centred near €38,300, where the pooled elasticity is low. The two schedules nearly meet again near €55,300. The administrative export reports nominal-income bins, so the within-bin shape and the division of their mass across 2015-

price bins are fitted rather than observed. The revenue requirement applies to the full fitted distribution, including its Pareto tail. The last estimated elasticity is carried forward above its €65,000 midpoint, including beyond the estimated €70,000 range, and should not be read as a new empirical finding. The comparison illustrates how the income profile of responses changes the rates prescribed by the calculation.

This exercise exploits only the variation of responses with own earnings. By construction, a one-dimensional formula has no notion of a partner and therefore cannot absorb the variation of elasticities with the composition of household income. Using that dimension requires a framework for couples, to which we now turn.

5.2 In the assessment of jointness reforms

For couples, solving for the optimal two-dimensional tax schedule is a multidimensional screening problem, which is analytically demanding even under strong simplifying assumptions (Goloso & Krasikov, 2025; H. J. Kleven et al., 2009)—and, under assumptions common in the perturbation literature, not recoverable by perturbing the status quo (Bierbrauer et al., 2026). In particular, the analytical characterisation in Goloso and Krasikov (2025) relies on a constant elasticity that is common to all individuals and therefore symmetric across spouses; its results would not survive elasticities that vary with the division of earnings within the couple. We therefore follow Bierbrauer et al. (2024), who evaluate reforms of the *existing* system rather than solving for the optimum; Bierbrauer et al. (2026) generalise the approach to multi-dimensional screening problems, the taxation of couples in particular. Starting from status-quo couple taxation, they derive sufficient statistics that determine whether reforms towards more individualised taxation—reforms that allow the marginal rates of the two spouses to differ, lowering them on secondary relative to primary earnings in a revenue-neutral way—are Pareto-improving and, where they are not, whether they raise welfare under given welfare weights.

Their baseline calibration sets the elasticity of primary earnings to 0.25 and that of secondary earnings to 0.75. To isolate the role of our estimated heterogeneity, our constant-response benchmark instead uses the headline French-taxpayer-weighted earnings elasticity, $e = 0.293$, for each spouse and the corresponding participation elasticity with respect to the household

net-of-marginal-tax rate, $\eta = 0.386$, for the household-level participation term. The estimated-response specification replaces these constants with the estimated surfaces over own income and the within-bin partner-income tercile, pooling the men’s and women’s estimates equally at each observed tile.

Figure 10 displays the welfare effects of local reforms towards a more individualised tax system under equal welfare weights. Panels (a) and (b) use the constant-response calibration and our estimated earnings and participation surfaces, respectively; Panel (c) reports their difference. Appendix Figure F.11 repeats the comparison with near-Rawlsian welfare weights.




Figure 10: Welfare effects of local reforms towards more individualised taxation

Notes: Colours report the calibrated welfare effect of a one-percentage-point jointness-reducing perturbation integrated over a €500-by-€500 local income cell following Bierbrauer et al., 2024; the scales are in model welfare units under this normalization, not euro equivalents. Blue denotes gains; red in Panel (c) denotes that the estimated-response calibration produces a smaller gain than the constant-response calibration. Panel (a) holds the earnings elasticity at $e = 0.293$ for both spouses and the participation elasticity at $\eta = 0.386$; Panel (b) uses the own-income-by-partner-tercile tiles for earnings and participation; Panel (c) is Panel (b) minus Panel (a). Marginal rates apply the statutory 2019 French schedule (Section 2.1) to a two-part couple after the standard 10% wage-income deduction, before tax credits and child-related adjustments. Men and women are pooled equally within each tile; pre-trend-flagged aggregates are retained to avoid selection on pre-period noise. The men’s €20,000–€30,000 middle-tercile earnings tile fails a disaggregated pre-trend test and is omitted before interpolation; the women’s estimate is used at that position. Primary earners are the higher-earning spouse; secondary-earner deciles are computed conditionally within each primary-earner decile. A decile cell is displayed only when the full simulated income range of both spouses lies within the estimates’ support, which omits the tenth primary-earner decile. The household participation term averages the two spouses’ predicted participation elasticities. All panels use equal welfare weights.

Heterogeneity changes both the level and the location of the effects. With equal welfare weights, the unweighted mean across the 90 supported local reforms rises from 470 to 510 model welfare units, and the largest gain rises from 965 to 1,654, in the corner where both partners earn little. Under near-Rawlsian weights (Appendix Figure F.11), the mean becomes less negative, from -483 to -442 . The large estimated responses where both partners earn little strengthen the local gain in that corner; elsewhere the effects move in both directions

rather than shifting uniformly—with equal weights, the assessed effect falls in 50 of the 90 supported cells.

The two exercises show that accounting for heterogeneity changes the shape of the optimal tax schedule and the assessment of reforms towards more individual taxation of couples.^{54,55,56}

6 Conclusion

This paper proposes a source of identifying variation for estimating tax elasticities, suited to heterogeneity analyses across gender, income, and the division of income within the household. Exploiting the institutional features of French family taxation, we use the fiscal exit of a couple’s first-born child—and, as a second design, marriage—as an event that changes the mapping from incomes to marginal tax rates, generating variation in marginal incentives through partner income at given own income, under testable assumptions.

Using the universe of French tax declarations from 2003 to 2023 and a panel structure that follows individuals across household-unit changes, we show that headline averages—weighted to the French taxpayer distribution, earnings elasticities of 0.17 for men and 0.29 for women, and participation elasticities of 0.21 and 0.28, both with respect to the household net-of-marginal-tax rate—mask substantial heterogeneity. Low own incomes raise three of these four averages: from €10,000 of own income up, women’s earnings elasticity is 0.23 and the participation elasticities are 0.15 (men) and 0.18 (women). Men’s earnings elasticity is 0.17 either way (their lowest income bin enters no aggregate). The participation responses of the lowest-earning women are driven mostly by those with higher-earning partners rather than by those whose partners also earn little, whereas their earnings responses are largest where both

⁵⁴The direction of the optimal primary–secondary differential is itself unsettled. With a common elasticity for all individuals, optimal marginal rates on secondary earners are on average *higher* than on primary earners, because low secondary earnings identify couples to whom a redistributive planner wants to target transfers (Golosov & Krasikov, 2025); the classic argument of Boskin and Sheshinski (1983) instead prescribes *lower* rates on the more responsive spouse. Which force dominates depends on the size and location of the elasticity gap between spouses—the object we estimate. Section 4.2.3 discusses possible sources of that gap.

⁵⁵Constant-elasticity calibrations are not an innocuous simplification in one dimension either: the welfare assessment of the post-war US tax cuts is sensitive to the assumed constant elasticity of taxable income (Bierbrauer et al., 2021); estimated elasticity profiles remove this ambiguity.

⁵⁶A further margin, which we do not pursue, is political feasibility: with heterogeneous elasticities, assessments of the political feasibility of moves towards individualisation may change (Bierbrauer et al., 2025).

partners earn little. The earnings response is largest in the event year and fades within two years; the participation response is more persistent (Section 4.2.2).

Replacing constant elasticities with profiles of the estimated shape changes the conclusions that standard frameworks for tax design reach: in the illustrative calibrations of Section 5, feeding the income profile into the sufficient-statistics formula flattens optimal marginal tax rates at the bottom of the income distribution, relative to the constant-elasticity benchmark; and replacing the constant per-spouse responses in the evaluation of reforms towards more individual taxation of couples concentrates the largest assessed gains where both partners earn little and reduces the assessed effect over much of the remaining joint distribution.

References

- Anderson, T. W., & Rubin, H. (1949). Estimation of the parameters of a single equation in a complete system of stochastic equations. *Annals of Mathematical Statistics*, 20(1), 46–63.
- Angrist, J. D., & Imbens, G. W. (1995). Two-stage least squares estimation of average causal effects in models with variable treatment intensity. *Journal of the American Statistical Association*, 90(430), 431–442.
- Attanasio, O., Levell, P., Low, H., & Sánchez-Marcos, V. (2018). Aggregating Elasticities: Intensive and Extensive Margins of Women’s Labor Supply. *Econometrica*, 86(6), 2049–2082.
- Bick, A., & Fuchs-Schündeln, N. (2018). Taxation and labour supply of married couples across countries: A macroeconomic analysis. *The Review of Economic Studies*, 85(3), 1543–1576.
- Bierbrauer, F. J., Boyer, P. C., & Peichl, A. (2021). Politically feasible reforms of nonlinear tax systems. *American Economic Review*, 111(1), 153–191.
- Bierbrauer, F. J., Boyer, P. C., Peichl, A., & Weishaar, D. (2024). *Pareto-improvements, welfare trade-offs and the taxation of couples* (CESifo Working Paper No. 11539). CESifo.
- Bierbrauer, F. J., Boyer, P. C., Peichl, A., & Weishaar, D. (2025). *The political economy of joint taxation* (CESifo Working Paper No. 11728). CESifo.
- Bierbrauer, F. J., Boyer, P. C., Peichl, A., & Weishaar, D. (2026). *Tax reforms and multi-dimensional screening* (Working paper, April 2026 version). Working paper.
- Blandhol, C., Bonney, J., Mogstad, M., & Torgovitsky, A. (2026). When is TSLS actually LATE? [Advance access]. *The Review of Economic Studies*.
- Blomquist, S., & Selin, H. (2010). Hourly wage rate and taxable labor income responsiveness to changes in marginal tax rates. *Journal of Public Economics*, 94(11-12), 878–889.
- Blundell, R., Duncan, A., & Meghir, C. (1998). Estimating labor supply responses using tax reforms. *Econometrica*, 66(4), 827–861.
- Blundell, R., Pistaferri, L., & Saporta-Eksten, I. (2016). Consumption inequality and family labor supply. *American Economic Review*, 106(2), 387–435.
- Borusyak, K., Hull, P., & Jaravel, X. (2022). Quasi-experimental shift-share research designs. *The Review of Economic Studies*, 89(1), 181–213.
- Boskin, M. J., & Sheshinski, E. (1983). Optimal tax treatment of the family: Married couples. *Journal of Public Economics*, 20(3), 281–297.
- Brewer, M., Saez, E., & Shephard, A. (2010). Means-testing and tax rates on earnings. In J. A. Mirrlees, S. Adam, T. Besley, R. Blundell, S. Bond, R. Chote, M. Gammie, P. Johnson, G. Myles, & J. M. Poterba (Eds.), *Dimensions of tax design: The Mirrlees review* (pp. 90–173). Oxford, Oxford University Press.

- Burns, S. K., & Ziliak, J. P. (2017). Identifying the elasticity of taxable income. *Economic Journal*, 127(600), 297–329.
- Callaway, B., Goodman-Bacon, A., & Sant’Anna, P. H. C. (2026). Difference-in-differences with a continuous treatment [Forthcoming]. *American Economic Review*.
- Chiappori, P.-A. (1992). Collective labor supply and welfare. *Journal of Political Economy*, 100(3), 437–467.
- Eissa, N. (1996). Tax reforms and labor supply. In J. M. Poterba (Ed.), *Tax policy and the economy, volume 10* (pp. 119–151). MIT Press.
- Eissa, N., & Liebman, J. B. (1996). Labor supply response to the Earned Income Tax Credit. *The Quarterly Journal of Economics*, 111(2), 605–637.
- Feldstein, M. (1995). The effect of marginal tax rates on taxable income: A panel study of the 1986 Tax Reform Act. *Journal of Political Economy*, 103(3), 551–72.
- Gelber, A. M. (2014). Taxation and the earnings of husbands and wives: Evidence from Sweden. *The Review of Economics and Statistics*, 96(2), 287–305.
- Golosov, M., & Krasikov, I. (2025). The optimal taxation of couples. *The Quarterly Journal of Economics*, 140(3), 2163–2211.
- Goux, D., Maurin, E., & Petrongolo, B. (2014). Worktime regulations and spousal labor supply. *American Economic Review*, 104(1), 252–276.
- Graber, M., Håvarstein, M., Mogstad, M., Torsvik, G., & Vestad, O. L. (2026). *Substitution and income effects of labor income taxation* (NBER Working Paper No. 34987). National Bureau of Economic Research.
- Gruber, J., & Saez, E. (2002). The elasticity of taxable income: Evidence and implications. *Journal of Public Economics*, 84(1), 1–32.
- Hendren, N. (2016). The policy elasticity. *Tax Policy and the Economy*, 30(1), 51–89.
- Herold, E., Koch, L., & Neisser, C. (2025, November). *Joint taxation and intra-household inequality: Evidence from same-sex couples* [Working paper]. Working paper.
- Herold, E., & Wallossek, L. (2024). The marriage earnings gap. *Working Paper*.
- Imbens, G. W., & Angrist, J. D. (1994). Identification and estimation of local average treatment effects. *Econometrica*, 62(2), 467–475.
- Immervoll, H., Kleven, H. J., Kreiner, C. T., & Saez, E. (2007). Welfare reform in European countries: A microsimulation analysis. *The Economic Journal*, 117(516), 1–44.
- Isaac, E. (2023). Suddenly married: Joint taxation and the labor supply of same-sex married couples after *U.S. v. Windsor*. *Journal of Human Resources*, 1220–11396R2.
- Jakobsen, K. M., & Søgaaard, J. E. (2022). Identifying behavioral responses to tax reforms: New insights and a new approach. *Journal of Public Economics*, 212, 104691.
- Kališková, K. (2014). Labor supply consequences of family taxation: Evidence from the czech republic. *Labour Economics*, 30, 234–244.

- Kleven, H., Kreiner, C., Larsen, K., & Søgaaard, J. (2025). Micro versus macro labor supply elasticities: The role of dynamic returns to effort. *American Economic Review*, 115(9), 2849–2890.
- Kleven, H. J. (2016). Bunching. *Annual Review of Economics*, 8(1), 435–464.
- Kleven, H. J., Kreiner, C. T., & Saez, E. (2009). The optimal income taxation of couples. *Econometrica*, 77(2), 537–560.
- Kleven, H. J., & Schultz, E. A. (2014). Estimating taxable income responses using Danish tax reforms. *American Economic Journal: Economic Policy*, 6(4), 271–301.
- Kopczuk, W. (2005). Tax bases, tax rates and the elasticity of reported income. *Journal of Public Economics*, 89(11), 2093–2119.
- Kumar, A., & Liang, C.-Y. (2020). Estimating taxable income responses with elasticity heterogeneity. *Journal of Public Economics*, 188, 104209.
- LaLumia, S. (2008). The effects of joint taxation of married couples on labor supply and non-wage income. *Journal of Public Economics*, 92(7), 1698–1719.
- Lee, D. S., McCrary, J., Moreira, M. J., & Porter, J. (2022). Valid t-ratio inference for IV. *American Economic Review*, 112(10), 3260–3290.
- Martínez, I. Z., Saez, E., & Siegenthaler, M. (2021). Intertemporal labor supply substitution? evidence from the Swiss income tax holidays. *American Economic Review*, 111(2), 506–546.
- Mirrlees, J. A. (1971). An exploration in the theory of optimum income taxation. *The Review of Economic Studies*, 38(2), 175–208.
- Montiel Olea, J. L., & Pflueger, C. (2013). A robust test for weak instruments. *Journal of Business & Economic Statistics*, 31(3), 358–369.
- Neisser, C. (2021). The elasticity of taxable income: A meta-regression analysis. *The Economic Journal*, 131(640), 3365–3391.
- Saez, E. (2001). Using elasticities to derive optimal income tax rates. *The Review of Economic Studies*, 68(1), 205–229.
- Saez, E. (2010). Do taxpayers bunch at Kink Points? *American Economic Journal: Economic Policy*, 2(3), 180–212.
- Saez, E., Slemrod, J., & Giertz, S. H. (2012). The elasticity of taxable income with respect to marginal tax rates: A critical review. *Journal of Economic Literature*, 50(1), 3–50.
- Selin, H. (2014). The rise in female employment and the role of tax incentives: An empirical analysis of the swedish individual tax reform of 1971. *International Tax and Public Finance*, 21(5), 894–922.
- Sigurdsson, J. (2025). Labor supply responses and adjustment frictions: A tax-free year in Iceland. *American Economic Journal: Economic Policy*, 17(4), 30–71.
- Stock, J. H., & Yogo, M. (2005). Testing for weak instruments in linear IV regression. In D. W. K. Andrews & J. H. Stock (Eds.), *Identification and inference for econometric*

models: Essays in honor of thomas rothenberg (pp. 80–108). Cambridge, Cambridge University Press.

Weber, C. E. (2014). Toward obtaining a consistent estimate of the elasticity of taxable income using difference-in-differences. *Journal of Public Economics*, 117, 90–103.

A Tax liability by household composition

Figure A.1 plots the income tax liability under the 2015 schedule for the household compositions of Figure 2, Panel (a), and for a married couple with one child. The child-related advantage is capped, so beyond the income at which the cap binds the one- and two-child schedules run parallel to the childless couple’s, at a constant distance equal to the cap.




Figure A.1: Income tax liability by household composition

Notes: Statutory income tax liability under the 2015 schedule of Figure 2, by household composition and total taxable income of the tax unit, with the cap on child-related advantages applied. Source: Own calculations based on OpenFisca France.

B The household model and the estimand

This appendix relaxes the separability assumption of Section 3.1: it derives the household estimating equation under a general convex joint cost of earning, shows that joint filing collapses it back to equation (2) with the coefficient read as the elasticity of own income with respect to the household rate, and makes precise which elasticities the two events identify—including for the marriage design, which switches the filing regime.

We model a household choosing both spouses’ incomes against a joint budget constraint,

$$\max_{z_{it}, z_{pt}} c_t - \Psi(z_{it}, z_{pt}; n_{it}, n_{pt}) \quad \text{s.t.} \quad c_t \leq z_{it} + z_{pt} - \mathcal{T}_t(z_{it}, z_{pt}),$$

where z_{pt} is the partner's income, $\mathcal{T}_t$ is the couple's tax function—whose form depends on the filing regime (Section 2.1), a dependence our design uses and reform settings usually abstract from—and Ψ is a convex joint cost of earning income. The joint cost may be additively separable across spouses or non-separable, reflecting leisure complementarities, home production, or coordinated retirement (Gelber, 2014; Goux et al., 2014).⁵⁷

The potential incomes n_{it} and n_{pt} index each spouse's cost of earning income; we normalise them so that n_{it} is the income i would earn if both spouses faced marginal tax rates of zero—as in the separable case of Section 3.1.

Writing $\tau_{it} \equiv 1 - \partial\mathcal{T}_t/\partial z_{it}$ and $\tau_{pt} \equiv 1 - \partial\mathcal{T}_t/\partial z_{pt}$ for the spouses' marginal net-of-tax rates, the first-order conditions are $\partial\Psi/\partial z_{it} = \tau_{it}$ and $\partial\Psi/\partial z_{pt} = \tau_{pt}$. Totally differentiating the pair around the pre-event optimum and solving the resulting two-equation system yields the elasticities

$$(B.1) \quad \epsilon_{ii} = \frac{\tau_{it}}{z_{it}} \frac{\Psi_{pp}}{D}, \quad \epsilon_{ip} = -\frac{\tau_{pt}}{z_{it}} \frac{\Psi_{ip}}{D}, \quad D \equiv \Psi_{ii}\Psi_{pp} - \Psi_{ip}^2 > 0,$$

where subscripts on Ψ denote second derivatives with respect to the spouses' incomes, evaluated at the optimum, and $D > 0$ by strict convexity—and the household estimating equation in changes,⁵⁸

$$(B.2) \quad \Delta \ln z_{it} = \epsilon_{ii} \Delta \ln \tau_{it} + \epsilon_{ip} \Delta \ln \tau_{pt} + \Delta \ln n_{it}.$$

Every coefficient in (B.2) is an elasticity of i 's *own* income; the subscripts index whose rate varies: ϵ_{ii} is the response of i 's income to i 's own rate, ϵ_{ip} its response to the *partner's* rate.

⁵⁷Quasi-linearity in household consumption shuts down income effects at the household level: neither the partner's net income nor the change in the household's tax *liability* affects income choices, and it makes the intra-household allocation of consumption irrelevant for those choices: with transferable utility, any sharing rule supported by lump-sum transfers leaves the first-order conditions unchanged. For the same reason the model's predictions for incomes coincide with those of an efficient collective household à la Chiappori (1992) with quasi-linear private utilities—we need not take a stance on the household decision process. Gelber (2014) rejects the income-pooling restriction of the unitary model; the rejection operates through income effects, which quasi-linearity shuts down, so the income predictions used here do not rest on the rejected restriction. See Appendix D for the extension with income effects.

⁵⁸For additively separable, iso-elastic Ψ the relation is exact and $\epsilon_{ip} = 0$; in general, the elasticities are local to the pre-event optimum. The error term is composite: with non-separable costs, changes in the *partner's* potential income also shift i 's income—they enter through the same cross-derivative Ψ_{ip} —and $\Delta \ln n_{it}$ stands for the combined contribution of both spouses' potential-income changes to i 's income. That this component does not vary with the dose within comparison cells is what Assumption 1 restricts; the model does not deliver it.

Both are derivatives of the household’s optimal income-supply function $z_i^*(\tau_{it}, \tau_{pt})$ —jointly re-optimised responses, in which the partner’s income also adjusts—and the cross-derivative Ψ_{ip} governs both: it is the numerator of ϵ_{ip} , and it enters ϵ_{ii} through D , so that even when i ’s rate alone changes, the partner’s re-optimisation raises ϵ_{ii} above the one-person elasticity $\tau_{it}/(z_{it}\Psi_{ii})$, by the factor $1/(1 - \rho^2)$ with $\rho \equiv \Psi_{ip}/\sqrt{\Psi_{ii}\Psi_{pp}}$. The split between ϵ_{ii} and ϵ_{ip} is thus by price, not by mechanism. Under quasi-linearity the cross-rate response is transmitted entirely through the partner’s behavioural adjustment, and it is zero when effort costs are separable, $\Psi_{ip} = 0$.⁵⁹

Under joint filing, the marginal rate is common to both spouses, $\tau_{it} = \tau_{pt} = \tau_{it}^H$ (Section 2.1), so any change in it is common as well, and equation (B.2) collapses to the estimating equation (2) of the main text,

$$\Delta \ln z_{it} = \underbrace{(\epsilon_{ii} + \epsilon_{ip})}_{\equiv \epsilon_i^H} \Delta \ln \tau_{it}^H + \Delta \ln n_{it},$$

with the coefficient read as the *elasticity of own income with respect to the household rate*, $\epsilon_i^H \equiv \epsilon_{ii} + \epsilon_{ip}$: the response of own income to the common rate, including the component transmitted through the partner’s adjustment. Explicitly, from (B.1), $\epsilon_i^H = (\tau_{it}^H/z_{it})(\Psi_{pp} - \Psi_{ip})/D$: it nests the one-person elasticity at $\Psi_{ip} = 0$ and exceeds it under complementarity ($\Psi_{ip} < 0$) through both channels, the cross term and the feedback.

Under non-separable costs, all three elasticities are couple-level parameters: both spouses’ cost curvatures enter each of them. They are local to the couple—the same individual would exhibit a different elasticity if single—and estimates from couples are not directly comparable to estimates on samples of singles. A gradient of responses in partner income at fixed own income can accordingly arise from four sources: sorting of individual elasticities across partner types, income effects (Appendix D), the non-separability studied here, or a common cost function whose local elasticity varies with the household rate, which differs across partner terciles at fixed own income; the design identifies the gradient, not its decomposition.

⁵⁹Separability is the assumption underlying most quantitative work on couples taxation (Bick & Fuchs-Schündeln, 2018).

At the child’s fiscal exit, the household files jointly before and after the event, so the event-induced rate change is common to both spouses, and the design recovers the elasticity of own income with respect to the household rate, ϵ_i^H .⁶⁰

At marriage, by contrast, spouses are individually taxed before the event, so the two spouses’ predicted changes generally differ—often only one spouse is predicted to see any change. Among couples for whom the event predicts a change in i ’s rate only, the event-induced variation moves $\Delta \ln \tau_{it}$ alone, and the reduced form identifies the own-rate elasticity ϵ_{ii} —itself inclusive of the partner’s re-optimisation (the factor $1/(1 - \rho^2)$ above). Among the same couples, the *partner’s* income response to i ’s rate change identifies the cross-rate elasticity ϵ_{pi} directly—the design of Gelber (2014) reproduced within our setting. The current implementation does not condition on the partner’s predicted change: among treated individuals the partner—typically the higher earner—may simultaneously be predicted to cross a threshold downwards, in which case the reduced form mixes ϵ_{ii} with a term in ϵ_{ip} . The contamination is bounded by the cross-rate elasticity, which is small by the estimates cited below.

The gap between ϵ_i^H and ϵ_{ii} is likely small: exploiting the Swedish tax reform of 1990–1991, in which spouses’ marginal rates moved independently, Gelber (2014) estimates compensated cross-rate elasticities of 0.05 for husbands and 0.07 for wives, an order of magnitude below his own-rate estimates of 0.41 and 0.47—and positive, so that ϵ_i^H weakly overstates ϵ_{ii} . Where the distinction matters, the marriage design can in principle provide internal estimates, and external ones can be imputed instead: $\epsilon_{ii} = \epsilon_i^H - \hat{\epsilon}_{ip}$.⁶¹

⁶⁰The income-weighted average of ϵ_i^H across spouses gives a different object: the response of the couple’s *combined* income to the household rate, the one relevant for revenue and Laffer calculations—whichever way the response splits across own- and cross-rate channels.

⁶¹In more detail, two routes lead from the identified total to the structural own-rate elasticity. **Route 1** (*spousal separability*): with $\Psi_{ip} = 0$, $\epsilon_{ip} = 0$ and $\epsilon_i^H = \epsilon_{ii}$; specialising to the iso-elastic form, the household problem separates into two copies of the one-person model, each facing the household rate—the formal justification for the standard practice of pairing the one-person model with household rate variation. **Route 2** (*external adjustment*): impute $\hat{\epsilon}_{ip}$ from Gelber (2014) and report $\epsilon_{ii} = \epsilon_i^H - \hat{\epsilon}_{ip}$ as an adjustment or as bounds; relative to our estimates, which are largest at low own incomes, the implied adjustment is modest but not uniformly negligible; the adjustment can be applied where the distinction matters, and none is applied here. The imputation matches the model’s object: because the model is quasi-linear, ϵ_{ip} is a compensated cross-rate elasticity, and Gelber’s estimates are compensated as well. The uncompensated added-worker channel operates through income effects—absent here, treated in Appendix D—and is opposite-signed and potentially large after major income shocks (Blundell et al., 2016); for bracket-switch variation, the associated liability changes are small relative to lifetime resources (for the predicted treated, 0.9 percent of household taxable income, about €400–€600 a year), which confines it to second order. Gelber’s estimates remain gender-specific and from a different population and a different tax change, so we treat them as orders of magnitude rather than point corrections, and would apply the adjustment to the earnings elasticities only. The distinction matters most for the heterogeneity results: gender differences in ϵ_i^H are individual parameters only under Route 1, or after the

C Proof of Proposition 1

Proof. By (4), (5) and Assumption 2, $\Delta y_{ir} = \Delta_r Y_i(0) + \mathbb{1}[Z_{ir} \neq 0] \epsilon_i \Phi_{ir}$ and $\Delta \text{NTR}_{ir} = \Delta_r \text{NTR}_i(0) + \mathbb{1}[Z_{ir} \neq 0] \Phi_{ir}$. By Assumption 1, the conditional means of the untreated components are equal for the $Z_i = \delta$ and $Z_i = 0$ groups of the same cell, so they cancel in the contrasts: $RF_c^r(\delta) = \mathbb{E}_c[\epsilon_i \Phi_{ir} \mid Z_i = \delta]$ and $FS_c^r(\delta) = \mathbb{E}_c[\Phi_{ir} \mid Z_i = \delta]$. Taking the ratio gives the weighted-average representation, with weights averaging to one by construction; under Assumption 3 all Φ_{ir} share the sign of δ , hence the sign of their conditional mean, and the weights are non-negative. $\square$

D Model extensions

D.1 Income effects

Relaxing quasi-linearity, we replace linear consumption with a concave utility $v(c_t)$: the household problem becomes

$$\max_{z_{it}, z_{pt}} v(c_t) - \Psi(z_{it}, z_{pt}; n_{it}, n_{pt}) \quad \text{s.t.} \quad c_t \leq z_{it} + z_{pt} - T_t^H(z_{it} + z_{pt}; Q_{it}) - \kappa_{it},$$

where κ_{it} collects lump-sum household expenses, which may change at the event (expenses for the child). The first-order conditions become

$$\frac{\partial \Psi}{\partial z_{jt}} = v'(c_t) \tau_{jt}^H, \quad j \in \{i, p\}:$$

income choices now depend on the marginal utility of household consumption, and therefore on partner income, on the level of the tax *liability*, and on needs κ_{it} . Log-linearising adds a virtual-income term to equation (2), in the spirit of Gruber and Saez (2002). Two remarks follow for our design. First, households predicted to cross a bracket at the event also experience a larger

Route 2 adjustment; without either, part of a measured gender gap can reflect the other spouse's adjustment transmitted within the household—mechanically so, since symmetry of Ψ restricts the cross responses in levels, not in elasticities: $\epsilon_{ip}/\epsilon_{pi} = \tau_{pt}z_{pt}/(\tau_{it}z_{it})$, the income ratio under a common rate, so the lower earner's cross-rate elasticity is the larger one (consistent with Gelber's larger estimate for wives), and the adjustment removes exactly this component. Because we estimate both spouses' responses to the same household shock within the same cells, the pair of estimates is itself informative about such reallocation, and we return to this when interpreting the heterogeneity results (Section 4.2.3).

increase in tax liability than their within-cell controls, so, absent a virtual-income instrument, our estimates should be read as total responses to the event-induced tax change. The two channels oppose each other—the higher marginal rate reduces income, while the larger liability raises it when leisure is normal—so under income effects our estimates are lower bounds on the compensated elasticity; the end of this subsection reports an attempt to separate the two responses. Second, the needs change is part of the event, not of the schedule change: the child leaves, and the expenses change, under both schedules (Section 3.3), so the income consequence of the needs change belongs to untreated growth and is restricted by Assumption 1 rather than Assumption 2. The classification has teeth: with concave v , even a needs change common to every household of a cell moves income through $v'(c_t)$, and consumption varies with partner income within the cell’s band—income effects can thus convert a common event shock into a partner-income *gradient* in growth. This is exactly the class of violation the validation region detects.⁶² Figure D.1 plots income growth against partner income on either side of the threshold: without income effects (Panel (a)), with an income effect that grows with the distance above the threshold (Panels (b) and (c), for a linear and a progressive schedule), and, in Panel (d), with a partner-income gradient in income growth that the event itself does not change, in place of the flat growth of Panels (a) and (b).

We tried to separate the two responses within the design by adding the mechanical change in the household’s tax liability at the event to the cell regressions as a second regressor. For the predicted treated this change is 0.9 percent of household taxable income on average, about €400–€600 a year. Within a cell it moves largely with the instrument: the predicted rate change is the same for every predicted-treated household of the cell, and the liability change differs from it only through the household’s distance above the threshold, which at given own income is mostly partner income. The liability term absorbs about half of the instrument’s variation within cells, and in the work-income regressions the unweighted median robust first-stage F over all estimated cells falls from 27 to 6.5 for women and from 23 to 5.0 for men. The coefficient on the liability term is negative: pooled over $r = 0$ to 3, -0.93 (0.20) for men and -0.43 (0.18) for women, for whom it is close to zero at $r = 0$ and negative afterwards (classical standard errors, which treat the event years as independent). Before the event it is close to

⁶²Appendix H takes income effects seriously for a different purpose: they are one reason why first-stage monotonicity can fail in reform-based simulated-IV designs, though not in ours.




Figure D.1: Income effects

Notes: Stylised income growth against partner income in $r = -1$ for one cell, before and after the event. Above the vertical dashed line the event moves the household into a higher bracket; below it the marginal rate is unchanged. Panel (a) has no income effects. Panels (b) and (c) add an income effect that grows with the distance above the threshold (dashed curve), under a linear and a progressive schedule. In Panel (d) income growth falls with partner income, with a gradient that the event itself does not change, in place of the flat growth of Panels (a) and (b). The dashed curve of Panel (d) is the before-event curve shifted by the event's direct effect alone, and the arrows split the fall above the threshold into that effect and the effect of the tax increase, which includes an income effect.

zero, and on participation it is not distinguishable from zero under the frequency weights of Section 3.4. A negative coefficient is the opposite of an income effect with normal leisure: among the predicted treated, households further above the threshold, whose partners earn more, reduce their income more. We read it as a partner-income gradient of the substitution response rather than as an income effect. The design does not separate the two channels, and the estimates remain total responses to the event-induced tax change.

D.2 The participation decision

The model of Section 3.1 cannot generate participation responses. Returning to the quasi-linear case, and following H. J. Kleven et al. (2009), consider a spouse deciding whether to work

at all, facing a fixed cost of participation q_{it} (commuting, childcare, forgone home production). Conditional on working, income is chosen optimally, so the value of participation is

$$V_{it}^w = \max_z \left\{ z - [T_t^H(z + z_{pt}; Q_{it}) - T_t^H(z_{pt}; Q_{it})] - \psi(z; n_{it}) \right\},$$

where the bracketed term is the *participation tax* (Immervoll et al., 2007): the increase in the household's tax liability caused by the spouse's income, given partner income and the household's tax schedule. The spouse works if and only if

$$V_{it}^w \geq q_{it}.$$

The relevant price of participation is therefore not the marginal rate at current income but the participation tax wedge—and its first-euro marginal rate, $1 - \tau^H$ evaluated at $z = 0$, is exactly what partner income and the quotient determine: because partner income alone places the household above the tax-free threshold, the second earner's first euro is taxed at the household's current bracket. A household pushed into a higher bracket by the child's fiscal exit thus sees the return to participation fall, and with heterogeneous fixed costs q_{it} the participation probability falls with the participation tax—the sign of the participation responses we document. The estimates themselves, however, are not expressed in this price: the implementation instruments the same object for earnings and participation, the event-induced change in the household *marginal* net-of-tax rate evaluated at current own income, whereas the wedge above is governed by the schedule between z_{pt} and $z_{it} + z_{pt}$, its first-euro rate by partner income and the quotient alone. Within cells both rise for the predicted treated, but the participation tax far less relative to the controls, as the ratio below shows; we therefore read the participation estimates, divided by pre-event participation, as responses to the marginal-rate change (Section 3.5).

To express them in the participation-tax price we compute, for every household of the estimation sample, the participation tax rate implied by the effective schedule at pre-event incomes, $[T_t^H(z_{pt} + z_{it}) - T_t^H(z_{pt})]/z_{it}$, under the pre-event and the post-event quotient, and multiply each cell estimate by the ratio of the cell's change in the log net-of-marginal-tax rate to its change in the log net-of-participation-tax rate, both as contrasts between the treated and the control households. The contrast in the participation tax is small because the control house-

holds' participation tax also rises: losing the half part raises the inframarginal tax on the spouse's earnings even where the marginal rate at household income is unchanged (for women earning €10,000–€20,000, the control households' participation tax rate rises from 11.2 to 13.9 percent when the half part is lost, the treated households' from 6.8 to 11.3 percent). The ratio averages 8.0 over the estimation cells (6.6 for women, 9.3 for men) and rises with own income, from about 2 in the bottom own-income bin to between 9 and 15 above €30,000, so the participation elasticities with respect to the net-of-participation-tax rate are about 3.4 times those of the main text in the French-taxpayer aggregate, with a different pattern across income: the bottom bin no longer stands out, and the rescaled profiles are largest, and least precise, in the middle and upper bins; Appendix F.7 reports them. The rescaling takes the realised pass-through of the mechanical change to be the same for both prices; a participation tax rate evaluated after the event at the pre-event own wage, with the household's realised taxable base and quotient, gives a ratio of 3.9 on the cell-periods for which it can be formed.

E Further case-study exhibits

This appendix collects the remaining exhibits of the example cell of Section 3.3: the men's validation-region figure (Figure E.1), the participation counterpart (Figure E.2), the first stage, the second-stage estimates by year relative to the event (Figure E.3) and the density test (Figure E.4).

The predicted-control group's realised marginal rate is a direct check of the mapping of Section 3.3: between $r = -1$ and $r = 0$ its household marginal rate rises by 1.9 percentage points on average for women and 1.7 for men, with a median change of zero, against 10.5 and 10.0 points for the predicted treated, whose predicted change is 16 points. Per unit of the predicted change in the log net-of-tax rate, the realised change is 0.53 (s.e. 0.01) for women and 0.52 (0.01) for men at $r = 0$, and declines to 0.38 and 0.40 by $r = 3$ as incomes move across brackets.

Translating the earnings paths of Figure 3 into elasticities with the estimation approach of Section 3 yields the second-stage estimates shown in Figure E.3.




Figure E.1: Identifying variation at the first-born's fiscal exit: men

Notes: Male counterpart of Figure 4: men of the example, same construction and horizons, with partner income in 20 equal-count bins of about 345 men (the two lowest, with mean partner income below €10,000, lie outside the plotted range). The validation region reaches up to about €17,200, the treatment region starts at about €22,000.






Figure E.2: Identifying variation at the first-born's fiscal exit: participation


Notes: Participation counterpart of Figures 4 and E.1: the outcome is an indicator for positive work income. Left axis: the share of the bin in work at $r = -2$ (before; close to one by construction, since the cell requires positive work income at $r = -1$) and the average of $r = +1$ and $r = +2$ (after), in 30 equal-count bins of partner income (about 115 women and 230 men per bin; bins with mean partner income below €11,000 are not shown); bands are 95% confidence intervals, capped at one. Right axis and sample as in Figure 4.




(a) Men, earnings elasticity



(b) Women, earnings elasticity



(c) Men, participation semi-elasticity



(d) Women, participation semi-elasticity

Figure E.3: Second-stage estimates for the case-study cell

Notes: Two-stage least squares estimates for the example cell of Figure 3, with the estimator of the headline analysis (Section 3), by year r relative to the first-born's fiscal exit ($r = -1$ is the reference); shaded areas are 95% confidence intervals from classical standard errors. The upper panels report the elasticity of own earnings with respect to the household net-of-tax rate, the lower panels the participation semi-elasticity: the change in the share in work per unit change in the log household net-of-marginal-tax rate. At $r = +1$ and $r = +2$ the earnings elasticities are 0.38 (s.e. 0.21) and 0.43 (0.38) for women and 0.21 (0.19) and 0.16 (0.26) for men; the participation semi-elasticities are 0.18 (0.07) and 0.57 (0.13) for women and 0.20 (0.06) and 0.25 (0.10) for men, and 0.65 (0.19) and 0.41 (0.13) at $r = +3$. The pre-event earnings estimates at $r = -2$ are 0.13 (0.11) for women and 0.15 (0.17) for men; those at $r = -3$, -0.37 (0.16) and -0.42 (0.20), rest on a pre-event first stage of 0.12 and 0.09 (those at $r = -2$ on 0.08 and 0.06), against 0.38 to 0.53 after the event.




Figure E.4: McCrary density test at the hypothetical joint-income threshold

Notes: Number of households in €200 bins of the distance of pre-event joint income to the threshold at which the marginal tax rate rises from 14% to 30% once the first-born's exit reduces the family quotient, within €4,000 of the threshold, for the example cell of Figure 3. Households above the threshold are predicted treated, households below it predicted control. If households anticipated the exit, by bunching below the future threshold or by timing the detachment, the pre-event density would show excess mass around zero. The McCrary test does not reject continuity: the estimated log-density discontinuity is -0.12 (s.e. 0.13 , $p = 0.35$) for men and 0.04 (s.e. 0.17 , $p = 0.79$) for women, on 3,388 and 1,735 households.

F Further results

F.1 The event samples

Cleaning of the raw panel drops households in the overseas departments and French residents abroad (1.29 million records in 2023), returns without any individual identifier (107,000) and individuals aged 75 and over (7.4 million), and resolves the individual-year duplicates that multiple declarations create (1.39 million in 2023). Table F.1 describes the two event samples; the figure notes of Section 4 refer to it. Panel A traces their construction. The child's fiscal exit sample comprises individuals in married or civil-union tax units whose first-born leaves the tax unit between 2004 and 2023 (1,932,140 women and 1,959,441 men); whose civil status is unchanged at the event, so that the fall in the quotient is the only change to the household's schedule (the restriction removes 0.9 percent of the women and 1.2 percent of the men); and whose own work income at $r = -1$ lies below €70,000 in 2015 prices, the cap removing a further 1.6 percent of the women and 9.2 percent of the men. In nominal terms, 95.6 percent of men and 97.0 percent of women filing a return in 2023 earn less than €70,000, so the cap covers at least that share of taxpayers. Work income at $r = -3$, on which the second estimation floor below (300 outcomes at $r = -3$) is counted, is observed for 87 percent of the sample. The

Table F.1: The two event samples

	Child's fiscal exit		Marriage	
	Women	Men	Women	Men
<i>Panel A: construction (individuals)</i>				
Event observed (2004–2023)	1,932,140	1,959,441	2,295,222	2,297,699
Marriage: predicted rate increase	–	–	575,517	203,268
Civil status unchanged; marriage: design screens ^a	1,914,489	1,936,085	309,067	121,554
Own work income at $r = -1$ below €70,000	1,884,452	1,758,070	308,483	121,154
Positive work income at $r = -1$	1,497,580	1,490,879	269,731	97,744
<i>Panel B: characteristics of the event sample at $r = -1$</i>				
Annual event cohort size (mean)	95,724	96,804	89,249 ^b	79,209 ^b
Age at $r = 0$ (mean)	51.09	53.28	33.61 ^c	35.49 ^c
Partner's age at $r = 0$ (mean)	53.39	50.94	35.65 ^c	33.52 ^c
Married at $r = 0$ (share)	0.98	0.98	0.53 ^c	0.55 ^c
Dependent children at $r = -1$ (mean)	1.92	1.90	0.27 ^c	0.35 ^c
Own work income (€, mean)	18,126	27,021	19,160	22,500
Partner's income (€, mean)	42,481	21,632	27,239	22,363
Positive work income (share)	0.79	0.85	0.93	0.93
Predicted into a rate change (share)	0.30	0.32	0.25	0.11
Mean $ \text{predicted } \Delta \ln(1 - \text{MTR}) $ if treated	0.12	0.12	0.12 ^c	0.12 ^c

Notes: Panel A counts individuals at each step of the sample construction described in the text; the fourth row is the event sample used throughout and the fifth the subsample on which the earnings outcome is defined. The marriage columns follow the spouses for whom the wedding predicts a rate increase; the design compares them with the zero-dose spouses of the same cells (941,102 women and 1,004,916 men below the income cap), so that the marriage estimation sample holds 1,249,585 women and 1,126,070 men. The child's fiscal exit counts of the fourth row exceed the grid totals of Table 3 by the 4 women and 14 men whose own work income at $r = -1$ is exactly €70,000, which the estimation's cap includes and the grid's bins exclude. Panel B reports descriptive statistics for the event samples below the income cap, measured in the pre-event year ($r = -1$) unless otherwise indicated; the marriage columns describe the estimation sample, predicted-increase and zero-dose spouses together. Partner's income sums work, self-employment, and pension income. Counts respect the data provider's disclosure rules. ^a Both spouses childless and separately taxed in the year before the wedding and filing jointly the year after. ^b Mean over the 2005 to 2017 cohorts; later cohorts are much smaller. ^c All first marriages below the income cap (2,282,283 women and 2,250,589 men), the marriage design's sample before its screens.

marriage sample comprises first marriages of spouses who were childless and separately taxed in the year before the wedding and file jointly the year after, restricted to spouses for whom the merger predicts a rate increase (Section 3.2) and compared within cells with the zero-dose spouses. Each sample serves two outcomes: the earnings outcome, log work-income growth, requires positive work income at $r = -1$ and in the outcome year; the participation outcome, an indicator for work income above €1,000 (2015 prices), does not (the participation sample includes individuals without work income at $r = -1$: two years before the event, 33 percent of the women and 22 percent of the men in the €0–10,000 comparisons are in work, against 97 percent or more in every bin from €10,000 up). In the €0–10,000 bin the comparison cells are formed separately for individuals with at least €400 of work income at $r = -1$, the

earner sub-sample, and for those below, the entrant sub-sample (Appendix F.2): at $r = -1$, 61 percent of the predicted-treated women and 85 percent of the predicted-treated men in the bin's comparisons have less than €400 of work income, and 36 and 14 percent are in work. Comparison cells enter estimation if at least 30 predicted-treated individuals carry the dose within the cell, each counted once at $r = -1$, and the dose's estimation sample, predicted treated and zero-dose comparison group together, holds at least 300 earnings outcomes at $r = -3$; the same floors gate the participation regressions of the dose; displayed estimates are further screened under the display conventions of Appendix F.2.

Panel B characterises the samples in the pre-event year. Parents at the first-born's fiscal exit are on average 51 to 53 years old and almost all married; 85 percent of the women and 82 percent of the men are aged 45 to 59 at the event. The exiting first-born is aged 19 to 26 in 86 percent of households (modal age 23; 3 percent are 18 or younger and 10 percent 27 or older; the age is known for 92 percent of households and is an upper bound where an older child left the tax unit before 2003). Partner incomes differ by gender, women's partners earning on average €42,481 and men's €21,632, against own work income of €18,126 and €27,021. Between 79 and 85 percent have positive work income at $r = -1$. About 30 percent are predicted into a rate change, with a mean predicted change in the log net-of-tax rate of 0.12 among them. The marriage sample is younger, 34 to 36 years old at the event, with fewer dependent children and higher participation; a quarter of the women and a tenth of the men are predicted a rate increase at the wedding.

F.2 Display conventions for the cell-level exhibits

The headline exhibits of Section 4 share a two-layer display convention, applied to the two-stage least squares cell-dose estimates of Section 3.5 (equation (6)), pooled over the post-event years $r = 0$ to 3 with the frequency weights of Section 3.4, which weight each cell-dose-period estimate by its regression sample, predicted treated and zero-dose comparison group together. Under these weights a cell whose comparison group is large relative to its predicted treated counts for more than its share of the treated; this matters most in the cells of households at the 14 percent rate predicted to cross into the 30 percent bracket, in which a few hundred predicted treated share regression samples of up to about 51,000 (Section 3.4 and

Appendix Table F.10). Standard errors and pre-event placebo tests follow Section 3.5: person-clustered within a comparison group, comparison groups and genders treated as independent; the balanced-sample, matching and marriage exhibits and the bracketed map values treat cells as independent.

The lowest own-income bin. In the €0–10,000 bin the comparison cells are formed separately for individuals with at least €400 of work income at $r = -1$, the earner sub-sample, and for those below, the entrant sub-sample, and each sub-sample cuts its own partner-income terciles: at €24,700 and €38,300 of partner income for women and at €13,800 and €26,200 for men in the earner sub-sample, and at €27,900 and €51,700 for women and at €11,600 and €26,400 for men in the entrant sub-sample. The bin’s aggregates pool the cell–dose–period estimates of both sub-samples with the frequency weights and screens of every other bin, and a map tile of the bin pools the two sub-samples’ cells of the same tercile rank.

Cell screens. The first layer decides which cell–dose–period estimates enter a displayed aggregate. An estimate enters if (i) its standard error is positive and at most 1 on the outcome scale and, on the elasticity scale, at most 0.5 in the income profiles, the event-time and subgroup exhibits and the aggregation table, and at most 0.4 in the maps, where a tile shows no confidence interval;⁶³ and (ii) its heteroskedasticity-robust first-stage F is at least 9, a t -ratio of 3 (Appendix F.11).⁶⁴ A failed pre-event placebo removes nothing; the aggregate is flagged and kept (*Pre-trend annotation* below).

Maps. The second layer decides how an aggregate is shown. Every rectangle of the own-income-by-partner-tercile grid is drawn, over the income range supported for both genders (own income up to €60,000), in one of four states. A rectangle is filled if at least one of its cells passes the cell screens, the standard error pooled over these cells is at most 0.4 and at least 1,000 individuals underlie them, predicted treated and comparison group together.⁶⁵ A

⁶³The outcome scale is that of the coefficient. For earnings the coefficient is the elasticity; for participation the elasticity divides the coefficient by pre-event participation, a share of at most one, and the elasticity-scale caps imply the outcome-scale bound of 1. Estimates above 5 in absolute value are dropped as outliers (the elasticity screen).

⁶⁴The classical first-stage t -ratio is used for the balanced-sample and marriage estimates, the level column of Appendix Table F.5 and one headline participation cell and dose, displayed at $r = 0$ and, in Appendix Figure F.3, as a pre-event point at $r = -2$; the matching estimates are not screened on the first stage.

⁶⁵Individuals are counted as each dose’s largest regression sample over $r = 0$ to 3, summed over doses, so a comparison group shared by several doses is counted once per dose; the coverage of a bin, below, uses the same count. The cap is evaluated on the pooled standard error that treats the rectangle’s cells as independent; the person-clustered standard errors printed on the maps play no part in it, and the largest of them among the filled tiles is 0.26.

rectangle is n.a. when none of its cells passes the elasticity, outcome-scale standard-error and first-stage screens; a rectangle whose cells pass those screens but none of which is within the elasticity-scale cap is drawn dashed, with the bracketed value pooled over them. A rectangle with cells within the cap is drawn dashed in the same way when their pooled standard error exceeds 0.4 or fewer than 1,000 individuals underlie them. A dashed cell prints the standard error below the bracketed value and the number of individuals behind it above; an n.a. cell prints the number behind all its cells. A rectangle with no cell is blank. Of the 72 rectangles of the four headline maps, 63 are filled; the men’s €0–10,000 earnings column and one women’s earnings rectangle have no cell within the cap, two women’s participation rectangles fail the floor, one women’s earnings rectangle is n.a. and two are blank (Appendix Tables F.7 and F.8).

Profiles, event-time and subgroup exhibits. These exhibits apply the cell screens, the event-time exhibits to their pre-event points as well, and no aggregate-level screen; an imprecise bin appears with a wide confidence interval rather than being suppressed. In the income profiles the vertical axis of a figure is set by the bins with at least two cell–dose–period estimates, and a bin resting on a single estimate is drawn in grey with its interval cut at the axis; the figure note gives the full interval.

The aggregation table. Table 4 adds one bin-level screen: a bin contributes only if the cells passing the cell screens cover at least 10 percent of the bin’s estimation sample (predicted treated and comparison group together) and at least 1,000 individuals underlie them. The first condition is one of representativeness and the second one of absolute size; neither is a precision rule, since the cells have passed the standard-error screens. Coverage is a rule of the table because the table multiplies a bin’s estimate by the population share of the whole bin: the men’s €0–10,000 bin carries 0.34 of the men’s French-taxpayer weight, while its one earnings cell that passes the screens holds 4,421 of the bin’s 77,460 individuals (5.7 percent). The screen removes that bin for earnings and, by the floor, the women’s €60,000–70,000 bin for earnings and participation (865 and 885 individuals, with full coverage). It applies to the table alone; the maps have the floor and the pooled cap in its place, and a bin shown in the income profiles may enter no aggregate.

Colour. The fill of a tile is a function of its printed estimate alone: classes with boundaries at ± 0.075 , 0.15, 0.225, 0.3, 0.4, 0.55, 0.8, 1.2 and 1.7, the class around zero symmetric, so that

estimates between -0.075 and 0.075 share the lightest tint whatever their sign, and estimates below -0.075 in grey. Precision enters no colour; it is the standard error printed beneath every value. Every map uses these classes, including those of the marriage design (Appendix G).

Pre-trend annotation. For every displayed aggregate, the pre-event placebo reduced forms of its cells at $r = -3$ and $r = -2$ (Section 3.5) are pooled with the aggregate’s own frequency weights—each placebo estimate weighted by its regression sample, as in the aggregate it annotates—and the pooled contrast is tested with the person-clustered covariance of the reduced forms. Aggregates with $|t| \geq 1.96$ are marked with four corner ticks and a grey value (maps; the tile keeps its colour), a dagger (tables) or a hollow marker (profiles and subgroup exhibits). The flag is an annotation: no estimate is dropped, altered or withheld from the colour scale on account of a failed pre-event placebo. The marriage exhibits use the same weights, with cells treated as independent. The event-time profiles carry no flag: their pre-event points are themselves placebo estimates.

F.3 Balance across predicted groups

The balance of predetermined characteristics behind the pooled validation of Section 4 is summarised in Table F.2: means for the predicted-control and predicted-treated groups and the standardised mean difference averaged across comparisons. The two groups are closely comparable on the outcome’s own pre-event dynamics—income growth and participation differ by less than 0.05 of a standard deviation. Differences of 0.2 to 0.3 appear in own and partner income levels and in age: within an own-income bin, the predicted treated are those whose partner earns more, so that their household income sits higher in the bin; own income, partner income and age are correlates of household taxable income, which determines assignment. The largest differences are in household taxable income itself and in family size, which shifts the number of fiscal parts and hence the predicted rate change. Behind the displayed estimates, the two arms differ in pension receipt at low incomes: between €10,000 and €30,000, in the cells behind the earnings estimates, 7 to 16 percent of the predicted treated drew a pension at $r = -1$ against 1 to 4 percent of their comparison groups (Table F.3).

Table F.2: Balance across predicted-treated and predicted-control comparisons

	Women			Men		
	Control	Treated	SMD	Control	Treated	SMD
<i>Pre-event outcomes</i>						
Work income, $r = -1$ (€)	16,558	17,167	0.23	27,116	27,710	0.23
Log work-income growth, $r = -3$ rel. $r = -1$ (%)	-2.35	0.68	0.04	11.24	14.28	0.02
Positive work income, $r = -1$ (%)	79.4	81.0	0.03	89.4	89.0	-0.01
<i>Composition</i>						
Partner work income, $r = -1$ (€)	25,783	28,382	0.24	16,085	17,679	0.30
Partner has positive work income, $r = -1$ (%)	84.6	87.2	0.04	77.3	82.3	0.10
Age	49.5	50.6	0.25	51.8	52.8	0.22
Partner age	52.1	52.9	0.17	49.4	50.6	0.28
Number of dependent children	2.24	1.97	-0.47	2.29	1.95	-0.56
Fiscal parts, $r = -1$	3.37	3.15	-0.45	3.41	3.13	-0.55
<i>Assignment and first stage (differ by design)</i>						
Household taxable income, $r = -1$ (€)	39,584	45,112	1.06	39,596	44,416	1.04
Predicted change in log net-of-tax rate ($\times 100$)	0.00	-11.42	-	0.00	-11.30	-
Realised net-of-tax-rate change, $r = +1$ (pp)	-1.36	-5.56	-	-1.49	-5.59	-
<i>Sample</i>						
Unique individuals	842,950	534,851		836,555	512,991	
Comparison-year cells		964			1,215	

Notes: Each treated tax-change dose is compared with the zero-change controls in the same own-income bin, partner-income tercile, pre-event bracket, and event year. Means are computed within each comparison-year arm and averaged across comparison-year cells with the cell's predicted-treated count as weight; SMD is the mean of the signed within-cell standardised differences with the same weights. Household taxable income determines assignment and differs between arms by construction (standardised difference 1.06 for women, 1.04 for men). The predicted rate change is zero in the control arm by construction; the realised change is not (controls whose income drifts across a threshold realise a change), and no standardised difference is reported for either row. The log-income-growth row requires positive work income at $r = -1$ and at $r = -3$. Controls may enter several dose comparisons; the counts report unique individuals. Levels are in euros of the event year, not deflated. Sample: the child's fiscal exit sample of Appendix F.1, in the comparison cells of the design. The complete table—all balance variables, the 90th percentile of absolute standardised differences, and per-row comparison counts—is available on request.

Table F.3: Composition of the two arms by own-income bin

Own income	Earnings cells				Participation cells			
	Draws a pension		Age		Draws a pension		In work	
	Treated	Control	Treated	Control	Treated	Control	Treated	Control
<i>Women</i>								
0–10k	0.15	0.06	51.3	49.3	0.23	0.11	0.78	0.76
10–20k	0.08	0.02	50.0	48.7	0.11	0.02	1.00	1.00
20–30k	0.07	0.01	51.3	49.8	0.08	0.01	1.00	1.00
30–40k	0.02	0.01	52.3	50.9	0.02	0.01	1.00	1.00
40–50k	0.02	0.01	52.3	51.5	0.02	0.01	1.00	1.00
50–60k	0.02	0.01	52.0	52.4	0.02	0.01	1.00	1.00
60–70k ^a	0.02	–	52.5	51.9	0.02	–	1.00	1.00
<i>Men</i>								
0–10k ^a	0.62	0.49	58.6	56.7	0.63	0.61	0.22	0.23
10–20k	0.16	0.04	52.7	51.4	0.17	0.05	1.00	1.00
20–30k	0.13	0.02	52.2	50.9	0.13	0.02	1.00	1.00
30–40k	0.12	0.02	53.2	51.5	0.10	0.02	1.00	1.00
40–50k	0.06	0.01	53.3	51.9	0.05	0.01	1.00	1.00
50–60k	0.03	0.01	53.3	52.2	0.04	0.01	1.00	1.00
60–70k	0.02	0.01	53.6	52.3	0.02	0.01	1.00	1.00

Notes: Predicted-treated and comparison arms at $r = -1$ behind the displayed profile bins of Figures 6 and F.4, each cell-dose-period estimate weighted by its regression sample as in the bins; ‘in work’ is work income above €1,000, which holds by construction from €10,000 up, and – marks a share resting on fewer than 11 individuals. ^a The earnings entries rest on a single cell-dose-period estimate (the grey bin of Figure 6).

F.4 Pooled density tests at the threshold




Figure F.1: Density tests at the hypothetical threshold, pooled across cells

Notes: Distribution of household taxable income at $r = -1$, measured relative to the hypothetical threshold at which the household would enter the 30 percent bracket after the child’s exit (at its post-exit number of fiscal parts) and pooled across the comparison cells of the design; unlike the observation counts of Figure E.4, the vertical axis is the share of households in each €500 bin, in percent. The estimated log-density discontinuities are -0.002 (s.e. 0.010, $p = 0.80$, $N = 782,811$) for women and -0.001 (s.e. 0.009, $p = 0.92$, $N = 776,649$) for men: no evidence at the threshold of sorting of the event’s timing or of bunching below it.

Figure F.1 shows the density of household taxable income around the hypothetical threshold, pooled across the comparison cells; its note gives the estimated discontinuities.

F.5 Event-time profiles

The event-time profiles of the estimates double as validity checks: the construction yields placebo estimates for the pre-event years (Figures F.2 and F.3). The earnings response is an impulse concentrated in the event year, and for women also the year after; the participation response is persistent throughout the window for women and declines for men.




Figure F.2: Event-time earnings elasticities around the child's fiscal exit

Notes: Sample: the child's fiscal exit sample of Appendix F.1 (all event years). The outcome is work-income growth relative to $r = -1$. Each point pools the cell-dose elasticities of one year relative to the event, each weighted by its regression sample across all own-income bins rather than by the population shares of the bins, so the points do not average to the aggregates of Table 4. The €0–10,000 bin holds 13 percent of the women's weight at $r = 0$ (3 to 4 percent at $r = 2$ and 3) and, for men, a single cell at $r = 0$. Shaded areas are 95% confidence intervals from person-clustered standard errors; a point pools a single year, so its standard error involves no covariance across years. The year before the event, $r = -1$, is the reference and is shown at zero (hollow marker: a reference point, not a failed placebo); estimates for $r = -3$ and $r = -2$ are placebo checks. They are scaled up relative to the post-event estimates: each divides the placebo reduced form by that year's pre-event contrast in realised rates, 0.10 to 0.13 against 0.43 to 0.45 in the event year; on the event-year scale the women's placebo at $r = -3$ is -0.03 , 8 percent of the event-year response, and the men's at $r = -2$ is -0.02 . Pooled over all pre-event cells with the weights of the estimates and tested with the person-clustered covariance, the placebo reduced form does not reject for men ($t = 0.53$); for women the test is borderline ($t = -1.9588$, $p = 0.050$; $t = -2.30$ if the pre-event cells are weighted by inverse variance instead). Cells enter under the cell screens of the income profiles (Appendix F.2); no aggregate-level screens apply. The participation counterpart is Figure F.3.




Figure F.3: Event-time participation elasticities around the child's fiscal exit

Notes: As Figure F.2, for participation: the outcome is an indicator for positive work income and the displayed estimates are participation elasticities, divided by participation two years before the event. The €0–10,000 bin holds 21 to 25 percent of the women's weight and 5 to 7 percent of the men's in each post-event year. The pooled pre-event placebo reduced form does not reject for either gender ($t = -1.08$ for women, 1.32 for men).

F.6 Aggregated elasticities

Weighted to the population of French taxpayers, the average participation elasticity is 0.247 (0.283 for women, 0.212 for men); across the four weighting schemes the average lies between 0.11 and 0.25. As with earnings, women's participation elasticities exceed men's under every scheme of Table F.4 except inverse-variance weighting. From €10,000 of own income up (Panel B of Table F.4), the French-taxpayer aggregate is 0.162 (0.178 for women, 0.147 for men); there it is an exit elasticity of established workers, whose participation two years before the event is 0.99.

F.7 Profiles along own income

This subsection reports the participation counterparts of the earnings profiles of Section 4.2.2. The outcome is an indicator for work income above €1,000 (2015 prices); bin estimates are participation elasticities (the change in the participation probability per unit change in the log household net-of-marginal-tax rate, divided by participation two years before the event), constructed as described in Section 3 with the samples, pooling window ($r = 0$ to 3) and screens of the earnings profiles. Because the earnings outcome requires only positive work income, a fall in work income to between €1 and €1,000 counts as an exit in the participation outcome

Table F.4: Aggregated participation elasticities

<i>A. All own-income bins</i>			
	Men	Women	Total
Population-weighted: French taxpayers	0.212 (0.039)	0.283 (0.029)	0.247 (0.024)
Population-weighted: event population	0.205 (0.033)	0.230 (0.033)	0.218 (0.023)
Inverse-variance weighted	0.127 (0.012)	0.096 (0.011)	0.111 (0.008)
Income-weighted	0.145 (0.016)	0.177 (0.022)	0.158 (0.013)

<i>B. Own work income at $r = -1$ of €10,000 or more</i>			
	Men	Women	Total
Population-weighted: French taxpayers	0.147 (0.017)	0.178 (0.024)	0.162 (0.015)
Population-weighted: event population	0.147 (0.018)	0.174 (0.027)	0.158 (0.015)
Inverse-variance weighted	0.124 (0.012)	0.086 (0.011)	0.105 (0.008)
Income-weighted	0.135 (0.015)	0.159 (0.023)	0.145 (0.013)

Notes: As Table 4, for participation: the outcome is an indicator for positive work income, and the entries are participation elasticities with respect to the household net-of-marginal-tax rate, the cell coefficients divided by participation two years before the event (Section 3.5). With respect to the net-of-participation-tax rate (Appendix D.2) the French-taxpayer aggregate is about 3.4 times larger (0.84 against 0.25; Table F.6). Panel B repeats the aggregation over the own-income bins from €10,000 up.

and as a large negative log change in earnings. Figure F.5 and Table F.6 show the same bins rescaled by the participation tax rate (Appendix D.2).

One difference from the earnings profiles stands out: the men's €0–10,000 bin, in which a single cell passes the screens for earnings, carries their largest participation elasticity (0.29, against 0.33 for women). The bin pools the earner and entrant sub-samples (Appendix F.2); the cells of the entrant sub-sample can record only entry into work, those of the earner sub-sample mostly exits. In the raw data, before the regression's controls and in each year from $r = 0$ to 3, predicted-treated entrants enter work at 2.4 to 4.7 percent a year (men) and 3.0 to 4.0 percent (women), against 3.2 to 4.6 and 4.4 to 7.6 percent in their comparison groups, and 41 to 61 percent of predicted-treated earner men leave work, against 33 to 58 percent (gaps of 3 to 12 percentage points). The women's bottom-bin value, 0.33 (67 cell-dose-period estimates, 360,546 individuals), pools 60 estimates of the earner sub-sample (0.17) and seven of the entrant sub-sample (1.11), all seven at one dose, the move from the 14 to the 30 percent bracket; without them the bin is 0.17, and leaving out one cell-dose at a time gives 0.20 to 0.40. The men's, 0.29 (27 estimates, 130,117 individuals), pools 16 estimates of the earner sub-sample (0.38) and 11 of the entrant sub-sample (0.26), which carry 75 percent of the bin's weight; leaving out one cell-dose at a time gives 0.19 to 0.50. Within the earner sub-sample

Table F.5: Earnings, participation and total responses by own income

Own work income at $r = -1$	Earnings	Participation	Sum	Level (PPML)
<i>Women</i>				
0–10k	0.45 (0.11)	0.33 (0.06)	0.78 (0.12)	0.08 (0.10)
10–20k	0.18 [†] (0.07)	0.14 (0.04)	0.33 (0.08)	0.13 (0.04)
20–30k	0.30 (0.07)	0.20 (0.04)	0.51 (0.08)	0.29 [†] (0.05)
30–40k	0.02 (0.05)	0.06 [†] (0.03)	0.08 (0.06)	0.02 (0.04)
40–50k	0.19 (0.09)	0.07 (0.04)	0.26 (0.10)	0.16 (0.07)
50–60k	0.31 (0.21)	0.19 (0.10)	0.50 (0.24)	0.27 (0.14)
60–70k ^a	-0.28 (0.32)	-0.11 (0.23)	-0.38 (0.39)	0.08 (0.24)
Population-weighted, 10k and up	0.23 (0.05)	0.18 (0.02)	0.41 (0.05)	0.20 (0.03)
<i>Men</i>				
0–10k ^a	0.45 (0.48)	0.29 (0.11)	0.73 (0.50)	2.40 [†] (0.39)
10–20k	0.32 (0.08)	0.17 (0.04)	0.49 (0.09)	0.31 (0.07)
20–30k	0.17 (0.05)	0.15 (0.03)	0.33 (0.06)	0.25 (0.04)
30–40k	0.13 (0.06)	0.15 (0.03)	0.28 (0.06)	0.18 (0.04)
40–50k	0.02 (0.07)	0.08 (0.03)	0.10 (0.08)	0.10 (0.05)
50–60k	0.05 (0.09)	0.06 (0.04)	0.11 (0.10)	0.07 (0.06)
60–70k	0.19 (0.13)	0.14 [†] (0.07)	0.32 (0.15)	0.18 (0.11)
Population-weighted, 10k and up	0.17 (0.03)	0.15 (0.02)	0.32 (0.04)	0.21 (0.03)

Notes: By gender and own-income bin at $r = -1$, pooled over $r = 0$ to 3: the earnings elasticity of Figure 6, conditional on positive work income in the year before the event and in the outcome year; the participation elasticity of Appendix Figure F.4; their sum; and the total (level) response of Section 4.2.1, the reduced form of each cell estimated in levels, including zeros, by Poisson pseudo-maximum likelihood (PPML) and divided by the cell's first stage. The population-weighted rows apply the French-taxpayer weights of Table 4 to the bins from €10,000 up. [†] The entry's pooled pre-event placebo fails ($|t| \geq 1.96$). ^a The earnings entry rests on a single cell–dose–period estimate (for the men's €0–10,000 bin also the Level entry); neither bin enters the population-weighted rows. Standard errors in parentheses are clustered by person; in the Level column each cell's first-stage sampling variance is added as if the cells' first stages were independent, and the standard error of the Sum treats the earnings and participation estimates as independent, although they are estimated on overlapping samples of the same people. Cells enter the Level column if their classical first-stage t -statistic is at least 3, the other columns if their robust one is. Three of the 278 women's Level cells have a negative first stage, all at doses that predict a rate cut (Section 3.2); dropping them changes the women's Level response from €10,000 up from 0.203 to 0.204 (standard error 0.026 in both cases).

the split by the couple's tax status before the event (Section 4.2.3) is the same as in the whole bin: men respond in couples that paid no income tax (0.74 against 0.04 in couples already taxed; 6 and 10 estimates), women in couples already taxed (0.50 against -0.01).




Figure F.4: Participation elasticities along the own-income distribution

Notes: As Figure 6, for participation: the outcome is an indicator for work income above €1,000, and the estimates are participation elasticities with respect to the household net-of-marginal-tax rate. The women's bottom bin includes seven estimates of the entrant sub-sample at one dose (1.11); without them it is 0.17. Hollow markers: the women's €30,000–40,000 bin and the men's €60,000–70,000 bin ($t = 1.96$) fail their pre-event placebo. The women's €60,000–70,000 bin rests on three cell-dose-period estimates and 885 individuals and enters no aggregate. No bin rests on a single estimate; the vertical axis is shared by the two genders.




Figure F.5: Participation elasticities with respect to the marginal and the participation tax rate

Notes: The bins of Figure F.4 (light blue: per unit change in the log household net-of-marginal-tax rate, the paper's series) and the same bins rescaled to the log net-of-participation-tax rate (navy). Each cell-dose-period estimate is multiplied by the cell's ratio of the change in the log net-of-marginal-tax rate to the change in the log net-of-participation-tax rate (Appendix D.2); cells whose participation-tax contrast is below 0.005 log points or of the wrong sign are dropped from the rescaled series (Table F.6 counts the cells behind each bin), and the bins are then aggregated with the frequency weights of Section 3.4. Hollow markers identify bins whose pooled pre-event placebo fails; shaded bands are 95% confidence intervals from person-clustered standard errors, which the rescaling widens roughly in proportion to the factor. The women's €60,000–70,000 bin rests on 885 individuals and enters no aggregate and is drawn in grey; its rescaled estimate, -1.06 , lies just below the end of the vertical axis at -1 , and its rescaled interval (-5.50 to 3.39) is cut at the axis. Sample, pooling window ($r = 0$ to 3) and screens as in Figure F.4.

Table F.6: Participation elasticities rescaled by the participation tax rate

Own income at $r = -1$	Marginal rate	Factor	Participation tax	Cells	
	elasticity		elasticity	paper	rescaled
<i>Women</i>					
0–10k	0.33 (0.06)	2.4	0.73 (0.17)	67	67
10–20k	0.14 (0.04)	6.7	-0.06 (0.26)	102	82
20–30k	0.20 (0.04)	8.6	0.44 (0.21)	102	84
30–40k [†]	0.06 (0.03)	12.1	1.11 (0.34)	60	51
40–50k	0.07 (0.04)	9.9	0.70 (0.42)	17	17
50–60k	0.19 (0.10)	9.3	1.81 (0.90)	5	5
60–70k	-0.11 (0.23)	10.0	-1.06 (2.27)	3	3
<i>Men</i>					
0–10k	0.29 (0.11)	1.6	0.52 (0.17)	27	27
10–20k	0.17 (0.04)	7.8	1.46 (0.40)	87	78
20–30k	0.15 (0.03)	8.9	1.04 (0.45)	122	93
30–40k	0.15 (0.03)	11.8	2.40 (0.36)	91	73
40–50k	0.08 (0.03)	11.9	0.99 (0.42)	49	48
50–60k	0.06 (0.04)	12.9	0.83 (0.55)	25	24
60–70k [†]	0.14 (0.07)	14.7	2.01 (1.12)	18	18

Notes: By gender and own-income bin at $r = -1$: the participation elasticity with respect to the household net-of-marginal-tax rate (the series of Figure F.4), the bin’s rescaling factor (the frequency-weighted mean over its cells of the ratio of the treated-minus-control change in the log net-of-marginal-tax rate to that in the log net-of-participation-tax rate), and the elasticity with respect to the net-of-participation-tax rate, each with its person-clustered standard error in parentheses; the last two columns count the cell–dose–period estimates behind the paper’s bin and behind the rescaled bin, the difference being the cells whose participation-tax contrast is below 0.005 log points or of the wrong sign. Because factor and estimate vary together within a bin, the rescaled bin is not the paper’s bin times the factor. [†] Bin whose pooled pre-event placebo fails. Weighted to the population of French taxpayers, the rescaled participation elasticity is 0.84, about 3.4 times the 0.25 of Table F.4. In the €0–10,000 bin, which pools the earner and entrant sub-samples, each cell takes the factor of the cell of the whole bin with the same partner-tercile rank, pre-event bracket and dose, on which the participation tax rates were computed; ignoring the tercile, or giving the cells of the entrant sub-sample the factor at a reference entry income of €5,000, moves the bin by at most 0.06.

F.8 Maps over own and partner income

Tables F.7 and F.8 report the estimate, standard error and number of underlying cell–dose–period estimates for every rectangle of the earnings and participation maps, in the display states of the maps. Five rectangles, four of them women’s, fail their pre-event placebo; their pooled pre-event $|t|$ lies between 2.09 and 2.63.

Table F.7: Tile-level estimates: women

Own income	Partner tercile	Partner income	Earnings		Participation	
			Estimate	Cells	Estimate	Cells
0–10k	low	≤ 24.7 / ≤ 27.9	0.63 (0.19)	4	0.02 (0.08)	24
	mid	24.7–38.3 / 27.9–51.7	0.29 (0.18)	4	0.69 (0.11)	21
	high	> 38.3 / > 51.7	-0.02 (0.14)	4	0.35 (0.08)	9
10–20k	low	≤ 23.1	0.01 (0.05)	23	-0.07 (0.04)	41
	mid	23.1–34.1	0.72 (0.25) [†]	13	0.64 (0.14)	30
	high	> 34.1	0.08 (0.05)	10	0.12 (0.02)	22
20–30k	low	≤ 25.3	0.44 (0.12) [†]	26	0.33 (0.09)	41
	mid	25.3–36.9	0.25 (0.06)	11	0.15 (0.03)	31
	high	> 36.9	0.05 (0.05)	11	0.07 (0.02)	22
30–40k	low	≤ 29.0	0.09 (0.09)	11	0.10 (0.04) [†]	28
	mid	29.0–42.3	-0.04 (0.07)	8	0.00 (0.04)	13
	high	> 42.3	-0.01 (0.13)	5	0.02 (0.07)	7
40–50k	low	≤ 32.8	0.24 (0.10) [†]	6	0.09 (0.04)	8
	mid	32.8–50.8	-0.05 (0.19)	3	0.03 (0.12)	6
	high	> 50.8	[1.22] (0.74)	1	[0.27] (0.14)	1
50–60k	low	≤ 36.3	0.31 (0.21)	3	0.20 (0.10)	4
	mid	36.3–60.3	n.a.	–	[0.01] (0.38)	1

Notes: One row per own-income bin (thousand euro) $\times$ within-bin partner-income tercile of Figure 7, for earnings and participation. Partner income gives the tercile's range of partner income at $r = -1$ in thousand euro; the lowest tercile includes negative values and the top tercile is open-ended. In the €0–10,000 bin the terciles are formed separately for the earner and entrant sub-samples (Appendix F.2) and pooled by rank; the first range is the earner sub-sample's, the second the entrant sub-sample's. Estimates follow the map-level display conventions of Appendix F.2 and match the printed tiles; standard errors in parentheses are clustered by person, and for bracketed entries treat the cell estimates as independent. Daggers mark the corner-ticked tiles of the maps, whose pooled pre-event placebo fails ($|t| \geq 1.96$); bracketed entries are estimates screened out of that map (dashed cells), n.a. marks rectangles with no usable cell, and rectangles with no cell in either map are omitted. Cells columns give the number of underlying cell-dose-period estimates; with the pooling window $r = 0$ to 3, one cell-dose can contribute up to four.

Table F.8: Tile-level estimates: men

Own income	Partner tercile	Partner income	Earnings		Participation	
			Estimate	Cells	Estimate	Cells
0–10k	low	≤ 13.8 / ≤ 11.6	[2.46] (0.57)	2	0.02 (0.26)	2
	mid	13.8–26.2 / 11.6–26.4	[0.70] (0.39)	5	-0.15 (0.22)	9
	high	> 26.2 / > 26.4	[0.32] (0.30)	4	0.42 (0.16)	8
10–20k	low	≤ 13.6	0.32 (0.13)	8	0.24 (0.07)	28
	mid	13.6–22.0	0.12 (0.10)	12	0.01 (0.06)	26
	high	> 22.0	0.29 (0.17)	5	0.22 (0.06)	19
20–30k	low	≤ 15.4	0.22 (0.05)	29	0.12 (0.04)	47
	mid	15.4–23.6	0.24 (0.17)	22	0.33 (0.08)	35
	high	> 23.6	0.11 (0.05)	12	0.16 (0.03)	29
30–40k	low	≤ 16.5	0.09 (0.11)	20	0.11 (0.06)	37
	mid	16.5–27.1	0.19 (0.08)	9	0.28 (0.04)	22
	high	> 27.1	0.10 (0.06)	8	0.07 (0.03) [†]	16
40–50k	low	≤ 16.6	0.05 (0.13)	9	0.14 (0.06)	19
	mid	16.6–28.9	0.06 (0.08)	5	0.06 (0.04)	13
	high	> 28.9	-0.01 (0.10)	4	-0.02 (0.07)	9
50–60k	low	≤ 15.8	0.07 (0.12)	5	0.09 (0.05)	12
	mid	15.8–29.9	-0.06 (0.12)	4	0.00 (0.07)	6
	high	> 29.9	0.13 (0.18)	2	0.16 (0.12)	5

Notes: As Table F.7, for men. No men’s earnings cell in the €0–10,000 bin is within the standard-error cap; the three bracketed entries pool the cells that pass the other screens. The €20,000–30,000 mid-tercile earnings tile (0.24) is dominated by a single comparison whose dose group fails a disaggregated pre-event placebo test, and is about zero without it; the rectangle-level placebo does not flag it, so it carries no dagger.

F.9 Alternative outcomes

Claimed professional-expense deductions (*frais réels*) were estimated as an outcome and are not reported: because the predicted bracket is computed from taxable income, which is net of the deductions themselves, households predicted to cross a threshold are selected on low deductions in the base year within their cell, and the pre-event placebo contrasts reject in the income bins that hold over 90 percent of the cells (four of six bins for men, three of five for women; reduced forms of -0.03 to -0.05 at $r = -3$ and -2 for both genders). Self-employment income on its own is estimated but not reported: few individuals in the event sample have it (7 percent of men and 3 percent of women before the event), its log change is about four to six times as dispersed as that of work income (standard deviation 0.64 against 0.11 to 0.18), so two to three percent of the cell–dose–period estimates pass the display screens against 21 percent for work income, and the pre-event estimates reject zero in 30 to 57 percent of the participation cells. Broad earnings, which add self-employment income to work income, are reported in Figure F.6.





Figure F.6: Broad earnings elasticities along the own-income distribution


Notes: As Figure 6, with broad earnings (work plus self-employment income) as the outcome, on cells defined by own work income as in the headline. Self-employment income sums the declared industrial, commercial, non-commercial and agricultural profit boxes (turnover before the abatement under the micro regimes, net profit under the real regimes); turnover declared under the micro-entrepreneur flat-rate scheme (*versement libératoire*) is not included. From €10,000 of own work income up, the broad earnings profile lies within 0.05 of the work-income profile for women and 0.02 to 0.08 above it for men. In both panels the bottom bin and the €30,000–40,000 bin fail their pre-event placebo (hollow markers); the bottom-bin estimates, 1.13 (women) and 1.05 (men), are not read. The women’s €60,000–70,000 bin rests on a single cell–dose–period estimate and is drawn in grey with its interval cut at the axis: -0.32 , interval -0.94 to 0.30 . The broad earnings outcome has no participation counterpart.

F.10 Excluding the 2020–2023 event years


The headline specification includes all event years. Figure F.7 re-estimates the maps without the 2020–2023 event years, which removes about a sixth of the predicted-treated individuals (8 to 22 percent in the bins below €40,000). The maps change little: from €10,000 of own income up, the tiles correlate with those of Figure 7 at 0.99 (women) and 0.82 (men) for earnings and at 0.99 and 0.96 for participation, with mean absolute differences of 0.02 to 0.05 (on all earnings tiles 0.98 and 0.82; Table 5). Every tile shown in the headline maps is shown here, and the women’s three earnings tiles whose pre-event placebo fails remain flagged, with larger values (0.82, 0.68 and 0.25 against 0.72, 0.44 and 0.24); for participation one tile’s flag differs, the women’s €30,000–40,000 tile in the bottom partner tercile, which fails its placebo in the headline map only.




(a) Men, earnings elasticity



(b) Women, earnings elasticity



(c) Men, participation elasticity



(d) Women, participation elasticity

Figure F.7: Earnings and participation elasticities by own and partner income: excluding event years 2020–2023

Notes: As Figure 7, excluding the 2020–2023 event years, with the same screens and display conventions (Appendix F.2).

F.11 First-stage strength

Appendix Table F.9 reports the robust first-stage F for all post-event cell estimates and for the cells behind each exhibit. Appendix Tables F.10, F.11 and F.12 show the income profiles under the stricter requirement $F \geq 104.7$, the tF -adjusted bin standard errors and the Anderson–Rubin sets. Section 4.3 summarises what they show.

Table F.9: First-stage strength of the cell estimates behind each exhibit

Cells	Women						Men					
	Number	Median F	Ratio	$F < 16.4$	$F < 23.1$	$F < 104.7$	Number	Median F	Ratio	$F < 16.4$	$F < 23.1$	$F < 104.7$
<i>Panel A: earnings</i>												
All estimates	835	275	0.89	14.7	16.2	33.8	948	353	0.82	15.2	17.9	29.3
Passing the elasticity and s.e. screens	174	1,402	0.87	0.0	0.1	3.3	203	1,016	0.79	0.0	0.0	1.0
Displayed: income and event-time profiles	174	1,402	0.87	0.0	0.1	3.3	203	1,016	0.79	0.0	0.0	1.0
of which in the maps	142	1,825	0.88	0.0	0.1	0.8	154	1,197	0.79	0.0	0.0	0.4
of which in the aggregation table	173	1,402	0.87	0.0	0.1	3.3	202	1,016	0.79	0.0	0.0	1.0
<i>Panel B: participation</i>												
All estimates	939	362	0.86	12.0	13.5	28.8	1,124	311	0.83	13.8	15.9	29.6
Passing the elasticity and s.e. screens	362	694	0.86	0.7	1.4	13.9	422	706	0.79	0.3	1.8	10.9
Displayed: income and event-time profiles	356	697	0.86	0.3	1.0	13.6	419	706	0.79	0.3	1.7	10.9
of which in the maps	307	728	0.86	0.1	0.5	11.7	342	759	0.78	0.2	1.3	8.1
of which in the aggregation table	353	697	0.86	0.3	1.0	13.6	419	706	0.79	0.3	1.7	10.9

Notes: First-stage strength of the cell-dose-period two-stage least squares estimates behind the exhibits of Section 4, $r = 0$ to 3 (in the €0–10,000 bin, the cells of the earner and entrant sub-samples). Each regression has one instrument and one endogenous regressor; F is the heteroskedasticity-robust first-stage F , i.e. the squared robust first-stage t -ratio, which with one instrument equals the effective F of Montiel Olea and Pflueger (2013) (the classical F for the one displayed cell screened on the classical t -ratio). Median F and Ratio, the median of the robust over the classical F , weight each cell by its regression sample; the last three columns give the share of that weight, in percent, in cells below the Stock and Yogo (2005) critical value for a maximal two-stage least squares size of 10 percent (16.38), below the Montiel Olea and Pflueger (2013) critical value for a worst-case bias of 10 percent (23.1), and below 104.7, the value above which the two-stage least squares t -ratio has correct 5 percent size (Lee et al., 2022). Displayed cells pass every cell screen of Appendix F.2, including the robust first-stage screen $F \geq 9$ ($|t| \geq 3$); the map and aggregation-table rows are the subsets that enter a displayed tile and the aggregation table.

Table F.10: Income profiles under a stricter first-stage requirement

Own income	Displayed ($F \geq 9$)		$F \geq 104.7$			Difference
	Estimate (s.e.)	Cells	Estimate	Cells	Weight kept	
<i>Panel A: earnings</i>						
Women						
0–10k	0.45 (0.11)	18	0.45	18	100.0	0.00
10–20k	0.18 (0.07)	58	0.18	58	100.0	0.00
20–30k	0.30 (0.07)	57	0.26	48	92.4	-0.04
30–40k	0.02 (0.05)	28	0.02	20	94.8	0.00
40–50k	0.19 (0.09)	9	0.19	8	99.1	0.01
50–60k	0.31 (0.21)	3	0.31	3	100.0	0.00
60–70k	-0.28 (0.32)	1	-0.28	1	100.0	0.00
Men						
0–10k	0.45 (0.48)	1	0.45	1	100.0	0.00
10–20k	0.32 (0.08)	41	0.32	41	100.0	0.00
20–30k	0.17 (0.05)	74	0.18	68	99.3	0.00
30–40k	0.13 (0.06)	44	0.13	36	97.2	0.01
40–50k	0.02 (0.07)	22	0.02	20	99.6	0.00
50–60k	0.05 (0.09)	13	0.05	13	100.0	0.00
60–70k	0.19 (0.13)	8	0.20	7	98.1	0.01
<i>Panel B: participation</i>						
Women						
0–10k	0.33 (0.06)	67	0.33	64	98.3	0.00
10–20k	0.14 (0.04)	102	0.05	72	86.7	-0.09
20–30k	0.20 (0.04)	102	0.13	53	77.1	-0.08
30–40k	0.06 (0.03)	60	0.05	28	85.6	0.00
40–50k	0.07 (0.04)	17	0.08	9	83.7	0.01
50–60k	0.19 (0.10)	5	0.20	4	96.5	0.01
60–70k	-0.11 (0.23)	3	-0.02	1	33.8	0.09
Men						
0–10k	0.29 (0.11)	27	0.29	27	100.0	0.00
10–20k	0.17 (0.04)	87	0.20	59	82.3	0.03
20–30k	0.15 (0.03)	122	0.17	83	91.2	0.01
30–40k	0.15 (0.03)	91	0.17	55	88.6	0.02
40–50k	0.08 (0.03)	49	0.08	26	90.0	0.00
50–60k	0.06 (0.04)	25	0.06	16	89.9	0.00
60–70k	0.14 (0.07)	18	0.08	9	77.6	-0.05

Notes: Own-income bins of the income profiles (Figure 6 and Appendix Figure F.4), frequency weights, $r = 0$ to 3: the displayed estimate with its person-clustered standard error, and the same construction with the first-stage requirement raised from $F \geq 9$ to $F \geq 104.7$ (Lee et al., 2022), all other screens kept. Weight kept is the share of the bin's frequency weight retained, in percent; Difference is the change in the bin estimate. The stricter construction is reported as a point estimate only. Largest move: 0.041 (earnings) and 0.090 (participation) over bins with at least two displayed cells. The two bins that move by more than one displayed standard error, women's participation at €10,000–20,000 and €20,000–30,000 (by -0.090 and -0.078), move because of cells of households at the 14 percent rate predicted to cross into the 30 percent bracket whose F lies below 104.7: dropping only those cells moves the bins by -0.079 and -0.097 . These cells hold 7.0 and 10.5 percent of the bins' frequency weight but 0.2 percent of their predicted treated; with predicted-treated weights the stricter screen moves the bins by -0.013 and 0.002 .

Table F.11: tF -adjusted standard errors of the income-profile bins

Own income	Estimate	s.e.	s.e. tF	Ratio	Max factor	Weight adjusted	$ t $	$ t tF$
<i>Panel A: earnings</i>								
Women								
0–10k	0.45	0.108	0.108	1.000	1.00	0.0	4.2	4.2
10–20k	0.18	0.066	0.066	1.000	1.00	0.0	2.8	2.8
20–30k	0.30	0.065	0.066	1.007	1.11	7.6	4.7	4.6
30–40k	0.02	0.051	0.051	1.004	1.36	5.2	0.4	0.4
40–50k	0.19	0.092	0.092	1.001	1.39	0.9	2.0	2.0
50–60k	0.31	0.215	0.215	1.000	1.00	0.0	1.4	1.4
60–70k	-0.28	0.319	0.319	1.000	1.00	0.0	0.9	0.9
Men								
0–10k	0.45	0.483	0.483	1.000	1.00	0.0	0.9	0.9
10–20k	0.32	0.084	0.084	1.000	1.00	0.0	3.8	3.8
20–30k	0.17	0.049	0.049	1.000	1.20	0.7	3.5	3.5
30–40k	0.13	0.057	0.057	1.001	1.40	2.8	2.2	2.2
40–50k	0.02	0.068	0.068	1.000	1.20	0.4	0.3	0.3
50–60k	0.05	0.091	0.091	1.000	1.00	0.0	0.5	0.5
60–70k	0.19	0.134	0.134	1.000	1.01	1.9	1.4	1.4
<i>Panel B: participation</i>								
Women								
0–10k	0.33	0.061	0.061	1.000	1.04	1.7	5.4	5.4
10–20k	0.14	0.039	0.040	1.032	1.44	13.3	3.7	3.5
20–30k	0.20	0.042	0.043	1.037	1.79	22.9	4.9	4.7
30–40k	0.06	0.027	0.028	1.062	1.52	14.4	2.2	2.1
40–50k	0.07	0.045	0.046	1.031	1.40	16.3	1.6	1.5
50–60k	0.19	0.096	0.098	1.013	1.39	3.5	2.0	2.0
60–70k	-0.11	0.226	0.262	1.158	1.20	66.2	0.5	0.4
Men								
0–10k	0.29	0.110	0.110	1.000	1.00	0.0	2.6	2.6
10–20k	0.17	0.043	0.044	1.015	1.23	17.7	3.9	3.8
20–30k	0.15	0.028	0.033	1.176	1.58	8.8	5.4	4.6
30–40k	0.15	0.029	0.031	1.041	1.76	11.4	5.2	4.9
40–50k	0.08	0.034	0.034	1.018	1.66	10.0	2.4	2.3
50–60k	0.06	0.042	0.043	1.023	1.65	10.1	1.4	1.4
60–70k	0.14	0.071	0.074	1.043	1.25	22.4	1.9	1.8

Notes: The displayed person-clustered bin standard error (s.e.) against a tF -adjusted version (s.e. tF). Each cell's robust standard error is multiplied by the 5 percent adjustment factor $c_{0.05}(F)/1.96$ of Lee et al. (2022, Table 3, Panel A) at its robust first-stage F ($c_{0.05}(F)$ is the tF critical value for $|t|$), linearly interpolated between the published grid points and equal to 1 for $F \geq 104.7$; Ratio is the bin standard error aggregated from the adjusted cell standard errors over the one aggregated from the unadjusted ones (cells treated as independent), and s.e. tF is the displayed standard error times Ratio. Max factor is the largest cell factor in the bin; Weight adjusted the share of the bin's frequency weight in cells with $F < 104.7$, in percent; $|t|$ the bin estimate over the respective standard error. The adjustment is a heuristic bound: the tF size guarantee is for a single just-identified t -ratio, not for a weighted mean of cells. Largest ratio over bins with at least two cells: 1.007 (earnings), 1.176 (participation).

Table F.12: Anderson–Rubin sets and Wald intervals, displayed cells

Own income	Cells	Bounded	Unbounded	Width ratio, median	Width ratio, p90	Both reject	Wald only	AR only
<i>Panel A: displayed cells, earnings</i>								
Women								
0–10k	18	18	0	1.003	1.006	4	0	0
10–20k	58	58	0	1.003	1.013	8	0	0
20–30k	57	57	0	1.004	1.025	8	0	0
30–40k	28	28	0	1.007	1.065	5	0	0
40–50k	9	9	0	1.003	1.134	1	0	0
50–60k	3	3	0	1.005	1.008	0	0	0
60–70k	1	1	0	1.013	1.013	0	0	0
All bins	174	174	0	1.004	1.023	26	0	0
Men								
0–10k	1	1	0	1.003	1.003	0	0	0
10–20k	41	41	0	1.005	1.010	1	0	0
20–30k	74	74	0	1.004	1.018	11	0	1
30–40k	44	44	0	1.004	1.041	4	0	0
40–50k	22	22	0	1.003	1.018	2	0	0
50–60k	13	13	0	1.005	1.010	2	0	0
60–70k	8	8	0	1.009	1.023	0	0	0
All bins	203	203	0	1.004	1.018	20	0	1
<i>Panel B: displayed cells, participation</i>								
Women								
0–10k	67	67	0	1.005	1.015	22	0	0
10–20k	102	102	0	1.008	1.056	29	0	1
20–30k	102	102	0	1.016	1.091	18	0	0
30–40k	60	60	0	1.024	1.148	8	0	2
40–50k	17	17	0	1.017	1.135	1	0	0
50–60k	5	5	0	1.008	1.140	1	0	0
60–70k	3	3	0	1.031	1.075	0	0	0
All bins	356	356	0	1.010	1.080	79	0	3
Men								
0–10k	27	27	0	1.003	1.014	5	0	0
10–20k	87	87	0	1.010	1.035	18	0	0
20–30k	121	121	0	1.010	1.062	31	0	1
30–40k	91	91	0	1.011	1.127	19	0	0
40–50k	49	49	0	1.019	1.134	6	0	0
50–60k	25	25	0	1.011	1.052	1	0	1
60–70k	18	18	0	1.021	1.081	1	0	0
All bins	418	418	0	1.011	1.079	81	0	2

Notes: Heteroskedasticity-robust Anderson–Rubin 95 percent sets (Anderson & Rubin, 1949) per cell, computed from the robust reduced-form and first-stage estimates of each regression on its instrumental-variable sample: with $\hat{\gamma} = \hat{\beta}\hat{\pi}$ the reduced form, V_γ and V_π the robust reduced-form and first-stage variances and C their robust covariance, the set is $\{\beta : (\hat{\gamma} - \beta\hat{\pi})^2 \leq 3.84(V_\gamma - 2\beta C + \beta^2 V_\pi)\}$, bounded if and only if $F > 3.84$. The robust two-stage least squares standard error satisfies $\text{s.e.}(\hat{\beta})^2 \hat{\pi}^2 = V_\gamma - 2\hat{\beta}C + \hat{\beta}^2 V_\pi$ on every displayed cell (largest relative deviation 2.2e-12). Cells counts the displayed cells with a robust first stage (all but the one screened on the classical t -ratio). Width ratio is the length of the Anderson–Rubin set over the length of the robust Wald interval $\hat{\beta} \pm 1.96 \text{s.e.}(\hat{\beta})$; Both reject, Wald only and AR only count cells where both, only the Wald interval, or only the Anderson–Rubin set exclude zero (the Anderson–Rubin test of a zero effect is the robust reduced-form t -test). Over all estimated cells with a robust first stage, 73.4 percent of the sets are bounded for earnings and 77.1 percent for participation; the others are unbounded, which happens exactly when $F \leq 3.84$.

F.12 Age and retirement restrictions

Figures F.8 and F.9 compare the earnings and participation profiles of individuals younger than 59 at the event with the headline profiles. Table F.13 reports the bins from €10,000 up and their aggregates in three restricted samples (Section 4.3).




Figure F.8: Earnings elasticities: all ages versus below 59

Notes: The panels overlay the headline profile (all ages, dark navy, solid) with the profile estimated on individuals below age 59 at the event (light blue, dashed). In the €0–10,000 bin the all-ages points (women 0.64, men 0.72) are not Figure 6's (0.45 for each); from €10,000 up the all-ages series is Figure 6. Below 59, the women's profile shifts down with its shape intact (lowest bin 0.64 to 0.48, s.e. 0.11 in both), and the men's lies below 0.10 in every bin from €10,000 to €60,000 (0.16, s.e. 0.11, above), with no bottom-bin estimate passing the screens. The age cut removes 4 percent of the predicted-treated women and 9 percent of the predicted-treated men. Construction, screens and pooling window as in Figure 6. Shaded areas are 95% confidence intervals from person-clustered standard errors, in the colour of the series; hollow markers flag a failed pre-event placebo of the series; grey markers are bins resting on a single cell-dose-period estimate, with the interval cut at the axis: the men's all-ages €0–10,000 bin, 0.72 with interval –0.20 to 1.64, and the women's €60,000–70,000 bin, –0.28 with interval –0.90 to 0.35 (all ages) and –0.01 with interval –0.69 to 0.66 (below 59).




Figure F.9: Participation elasticities: all ages versus below 59

Notes: As Figure F.8, for participation; no bin rests on a single cell-dose-period estimate. In the €0–10,000 bin the lowest-earning men below 59 keep a large participation elasticity (0.84, s.e. 0.15, on 10 estimates, against 0.72, s.e. 0.08, at all ages, a point that is not Appendix Figure F.4's 0.29), although the cut removes 57 percent of the predicted-treated men in that bin. Above it, men's elasticities below 59 range from –0.01 to 0.14 (all ages 0.06 to 0.17). Women's are within 0.03 of the all-ages ones from €10,000 to €50,000 and within one standard error in the two top bins; their lowest bin falls from 0.20 to 0.03. The age cut removes 7 percent of the predicted-treated women and 15 percent of the men.

Table F.13: Responses from €10,000 up in three restricted samples

	Women				Men			
	Headline	Below 59	No pension	Floor 400	Headline	Below 59	No pension	Floor 400
Own income								
<i>A. Earnings elasticity</i>								
10–20k	0.18 [†] (0.07)	0.08 [†] (0.05)	0.03 (0.05)	0.16 [†] (0.06)	0.32 (0.08)	0.09 (0.07)	−0.01 (0.08)	0.09 (0.06)
20–30k	0.30 (0.07)	0.18 (0.05)	0.08 (0.04)	0.18 [†] (0.06)	0.17 (0.05)	0.09 (0.04)	0.09 (0.04)	0.14 (0.04)
30–40k	0.02 (0.05)	−0.07 (0.05)	−0.04 (0.05)	−0.04 [†] (0.04)	0.13 (0.06)	0.05 (0.05)	0.05 (0.04)	0.02 (0.04)
40–50k	0.19 (0.09)	0.13 (0.08)	0.15 [†] (0.09)	0.14 [†] (0.07)	0.02 (0.07)	−0.11 (0.06)	−0.10 (0.06)	0.03 (0.05)
50–60k	0.31 (0.21)	0.20 (0.22)	0.32 (0.21)	0.20 (0.18)	0.05 (0.09)	0.04 (0.08)	0.00 (0.08)	0.05 (0.07)
60–70k	−0.28 ^a (0.32)	−0.01 ^a (0.34)	0.05 ^a (0.34)	−0.33 ^a (0.26)	0.19 (0.13)	0.16 (0.11)	0.18 (0.13)	0.21 (0.14)
Aggregate ^b	0.23 (0.05)	0.10 (0.03)	0.06 (0.03)	0.15 (0.04)	0.17 (0.03)	0.06 (0.03)	0.01 (0.03)	0.08 (0.03)
Removed (%) ^c		3.5	1.5	0.2		7.7	2.6	0.3
<i>B. Participation elasticity</i>								
10–20k	0.14 (0.04)	0.12 (0.04)	0.08 (0.05)	0.14 (0.04)	0.17 (0.04)	0.06 (0.04)	−0.01 (0.04)	0.17 (0.04)
20–30k	0.20 (0.04)	0.18 (0.04)	0.11 (0.04)	0.20 (0.04)	0.15 (0.03)	0.12 (0.02)	0.14 (0.04)	0.15 (0.03)
30–40k	0.06 [†] (0.03)	0.03 (0.03)	0.03 [†] (0.03)	0.06 [†] (0.03)	0.15 (0.03)	0.11 [†] (0.03)	0.11 [†] (0.03)	0.15 (0.03)
40–50k	0.07 (0.04)	0.06 (0.04)	0.06 (0.04)	0.07 (0.04)	0.08 (0.03)	0.02 (0.03)	−0.03 (0.03)	0.08 (0.03)
50–60k	0.19 (0.10)	0.13 (0.09)	0.16 (0.09)	0.19 (0.10)	0.06 (0.04)	−0.01 (0.04)	0.01 (0.04)	0.06 (0.04)
60–70k	−0.11 (0.23)	−0.02 (0.24)	−0.01 (0.21)	−0.11 (0.23)	0.14 [†] (0.07)	0.14 (0.07)	0.10 (0.07)	0.14 (0.07)
Aggregate ^b	0.18 (0.02)	0.15 (0.02)	0.11 (0.03)	0.18 (0.02)	0.15 (0.02)	0.11 (0.02)	0.07 (0.02)	0.15 (0.02)
Removed (%) ^c		3.7	1.9	0.0		8.3	3.1	0.0

Notes: Each column re-estimates the profiles of Figure 6 (Panel A) and Appendix Figure F.4 (Panel B) on a restricted sample, with the same cells, screens, pooling window ($r = 0$ to 3) and weights; *Headline* repeats those figures. *Below 59*: individuals younger than 59 at the event. *No pension*: individuals with no pension income of their own, and no missing pension record, in the year before the event ($r = -1$). *Floor 400*: the earnings outcome is kept only for individuals with more than €400 of work income at $r = -1$ and in the outcome year; participation is not restricted, so Panel B repeats the headline estimates to two decimals (men's €60,000–70,000 placebo: $t = 1.96$ there, 1.94 under the floor). The €0–10,000 bin is omitted (under *Floor 400*, the men's earnings bin is 0.34, s.e. 0.21, on six estimates, conditional on outcome-year work income). Standard errors, clustered by person, in parentheses. [†] The bin's pooled pre-event placebo fails ($|t| \geq 1.96$). ^a A single cell-dose-period estimate. ^b Bins entering Table 4 from €10,000 up, with its weights. ^c Share of the panel's estimation sample from €10,000 up that the restriction removes, counted on the dose groups both runs estimate.

F.13 Heterogeneity

Figure F.10 gives the gender breakdown of the subgroup estimates of Section 4.4.




Figure F.10: Elasticities by household subgroup and gender

Notes: As Figure 8, separately by gender: each gender's rows pool that gender's cells only. A hollow marker means that the gender's own pre-event placebo fails at the 5 percent level, tested with the estimate's weights and covariance.

F.14 Welfare effects under near-Rawlsian weights

Figure F.11 repeats the exercise of Figure 10 with near-Rawlsian welfare weights.




Figure F.11: Welfare effects of local reforms under near-Rawlsian weights

Notes: As Figure 10, using near-Rawlsian welfare weights concentrated on the lowest-income households. Panels (a) and (b) show the constant- and estimated-response calibrations; Panel (c) reports estimated minus constant. Colour scales are in model welfare units under the €500-by-€500 local-cell and one-percentage-point reform normalization, not euro equivalents. Blue denotes gains and red denotes losses in Panels (a) and (b); in Panel (c), blue denotes a larger and red a smaller welfare effect under the estimated-response calibration.

G The marriage design

The marriage design uses first marriages of spouses who were childless and separately taxed in the year before the wedding and file jointly the year after, restricted to spouses for whom the merger predicts a rate increase and compared within cells with the zero-dose spouses. Below €70,000 of own work income the estimation sample holds 1,249,585 women and 1,126,070 men, of whom 308,483 (25 percent) and 121,154 (11 percent) are predicted a rate increase. The predicted-increase share falls with own income, from 58 (women) and 36 (men) percent in the lowest bin to under 13 and 6 percent above €30,000, which is why the design identifies little above that level; 93 percent of both genders have positive work income at $r = -1$, and 42 and 40 percent of the weddings precede the 2011 reform. The design draws on the 2005 to 2017 cohorts, of 83,000 to 113,000 women each; later cohorts are much smaller.

The design follows the child's fiscal exit design of Sections 3.2 and 3.5, with the following differences. The instrument applies the pre-event formula to the two spouses' $r = -1$ incomes pooled into one unit with quotient 2, on the statutory schedule; computed on the effective schedule of Section 2.1, as at the child's exit, it removes the women's response at €10,000–20,000 and fails its placebo on the strict sample (both spouses also never married or in a civil union before), and is not used. Only spouses whom the merger moves to a higher marginal

rate (“marrying up”) are kept, and they are compared within cells with the zero-dose spouses; after the residualisation described below, the predicted-decrease side is indistinguishable from zero for earnings and participation. Cells fix the own-income bin, the partner-income tercile within the bin and the pre-event bracket, as elsewhere; a cell enters if, in addition to the floors of Appendix F.1, it holds at least 30 zero-dose spouses. The controls replace civil status and the number of children, which the sample’s screens fix, by the number of fiscal parts of each spouse before the wedding. The plain estimates carry a post-event contrast of the same sign for predicted increases and predicted decreases, so the outcomes are residualised before estimation: within each cell and period the outcome is regressed, in the zero-dose group, on the partner’s income at $r = -1$ (capped at its within-cell 95th percentile), both spouses’ ages and own work income at $r = -2$ and $r = -3$, with the fixed effects of the estimation, and the fitted values are subtracted from every observation. For weddings before 2011, when the wedding year was declared on separate returns before the wedding and a joint one after, the $r = 0$ incomes are imputed from the split-year returns and are not used, so $r = 0$ rests on the 2011 to 2017 cohorts alone; on the strict sample the women’s participation reduced form at $r = 0$ has the opposite sign to the later years, which it does not on the sample used here. The exhibits pool $r = 0$ to 3 as elsewhere, and we report no event-time profile for this design; a participation build-up after the event would be read over $r = 1$ to 3.

The estimates are pooled with the frequency weights of Section 3.4. In this design the weights matter more than at the child’s exit: the marrying-up spouses are a small share of a cell (a median of 5.4 percent for women and 3.6 percent for men), the zero-dose spouses are shared by the doses of a cell, and a dose group with few marrying-up spouses carries the weight of the whole comparison pool. Where the two weightings differ by more than 0.02, both are given, and the treated-weighted value is the one to quote.


Within an own-income bin the marrying-up spouses face rate jumps of different sizes, and the mix changes from bin to bin, so the marriage response is one number per bracket jump rather than a profile. For the jump from the 14 to the 30 percent bracket below €30,000 (9,331 and 46,556 marrying-up women and 2,820 and 18,529 men in the €10,000–20,000 and €20,000–30,000 bins), the earnings elasticity is 0.14 and 0.18 for women and 0.32 and 0.16 for men (standard errors 0.02 to 0.04), within 0.02 of these values under treated weights; the participation elasticity for the same jump is 0.11 and 0.08 for women and 0.09 and 0.09 for men

under treated weights (under frequency weights the women’s €20,000–30,000 value is 0.75, carried by cells with few marrying-up spouses and large comparison pools). No jump-level estimate of the 14 to 30 percent transition below €30,000 fails its pre-event placebo under frequency weights; under treated weights, the women’s €20,000–30,000 estimates fail it for earnings ($t = 4.61$) and for participation ($t = 2.29$). Above €30,000 no jump gives a number that is stable across weightings: the 14 to 30 percent jump is indistinguishable from zero for earnings under both weightings and changes sign or significance with them for participation, and the 30 to 41 percent jump reaches 500 marrying-up spouses only for women (502 to 576), where its estimates move by 0.1 to 0.3 with the weights.


Figures G.1 and G.2 show the design in the paper’s two formats, maps by own and partner income and profiles along own income, under the display conventions of Appendix F.2 with three differences, stated in the notes: own income from €10,000 on every exhibit, because below it fewer than half of the spouses are in work two years before the wedding, against nearly all above, so that a participation response there would measure entry into work rather than the exit from work that the design measures elsewhere; a floor of 500 marrying-up spouses behind the passing cells on maps and profiles alike, in place of the floor of 1,000 individuals, which never binds here because every marrying-up cell carries a large shared comparison pool; and a map grid from €10,000 to €50,000, the bins with at least 500 marrying-up spouses behind displayed cells for at least one gender. The profiles are the paper’s format, not a gradient to read: their bins mix bracket jumps and, above €30,000, rest on 534 to 4,961 marrying-up spouses. Below €30,000 they show the weighting at work: the women’s earnings bins are 0.39 and 0.33 under frequency weights and 0.18 and 0.19 under treated weights, the men’s 0.30 and 0.04 against 0.31 and 0.17. Two of the 16 shown tiles and two of the 14 shown bins fail their pre-event placebo, pooled with the weights of the estimate: the women’s earnings tiles at €20,000–30,000 in the high partner tercile (0.31) and at €30,000–40,000 in the middle tercile (−0.35), the women’s earnings €20,000–30,000 bin and the men’s participation €10,000–20,000 bin.

Inference in the marriage design is classical. Cells are screened on the first-stage rule of Appendix F.2 in its classical form, which no cell fails, and standard errors and placebo tests treat the cells as independent. Because a cell pools four post-event years of the same people, these standard errors are too small: in the child’s fiscal exit design, from €10,000 up, person-


clustered standard errors exceed the classical ones by a median factor of 1.45 over the income-profile bins and by between 1.18 and 1.91 for four in five map tiles.




(a) Men, earnings elasticity



(b) Women, earnings elasticity



(c) Men, participation elasticity



(d) Women, participation elasticity

Figure G.1: Marriage design: earnings and participation elasticities by own and partner income, predicted rate increases

Notes: Sample: the marriage sample of Table F.1, restricted to spouses whom the merger moves to a higher marginal rate (marrying up) and compared within cells with the zero-dose spouses. Outcome: work-income growth relative to $r = -1$ in Panels (a) and (b), an indicator for positive work income in Panels (c) and (d), residualised on pre-event covariates fitted in the zero-dose group as described in the text, pooled over $r = 0$ to 3 with frequency weights; the participation estimates are divided by pre-event participation. Display conventions as in Figure 7 and Appendix F.2, with three differences: own income from €10,000; a floor of 500 marrying-up spouses in place of 1,000 individuals (the counts above dashed cells are marrying-up spouses); and a grid to €50,000. Every cell of an n.a. rectangle here fails the outcome-scale standard-error screen; a blank rectangle has no marrying-up cell. Two tiles fail their pre-event placebo (corner ticks and a grey value): the women's earnings tiles at €20,000–30,000 in the high tercile and at €30,000–40,000 in the middle tercile. Cells are screened on the classical first-stage t -ratio, which no cell fails, and standard errors and placebo tests are classical, with the cells treated as independent. These standard errors are too small: in the child's fiscal exit design, person-clustered standard errors exceed the classical ones by a median factor of 1.45 from €10,000 up.




Figure G.2: Marriage design: elasticities along the own-income distribution, predicted rate increases

Notes: Sample, outcomes and pooling as in Figure G.1. Each marker pools the cell-dose elasticities of an own-income bin with frequency weights; the band is the 95 percent confidence interval. A navy marker passes the cell screens of Appendix F.2, the standard-error cap of 0.5 and the floor of 500 marrying-up spouses; a hollow marker fails its pre-event placebo: the women's earnings bin at €20,000–30,000, the men's participation bin at €10,000–20,000 and the screened women's earnings bin at €60,000–70,000. Grey diamonds with dashed intervals are bins below the floor (114 to 392 marrying-up spouses), drawn so that no estimate is dropped silently; none enters a number quoted in the text. Their intervals are cut at the axis, and a diamond on the edge with a bracketed value lies beyond it. Every displayed bin pools at least eight cell-dose-period estimates. Bins above €30,000 mix bracket jumps; the text reads the design by jump. First-stage screen, standard errors and placebo tests as in Figure G.1: classical, which understates the uncertainty.

H First-stage monotonicity

This appendix (i) states the monotonicity condition in the canonical reform-based simulated-IV setting (Gruber & Saez, 2002; Jakobsen & Søgaaard, 2022), (ii) characterises when the condition fails and illustrates a violation in a calibrated example, and (iii) closes by arguing that violations are far less likely under schedule *shifts*, as in our design, than under reforms that change the schedule's gradient.

H.1 The condition in the reform-based setting

Consider a reform that raises statutory rates on part of the schedule. Let $g \in \{0, 1\}$ indicate whether the reform is in force, and let $\text{MTR}_i(g)$ denote the marginal tax rate individual i faces under regime g , inclusive of behavioural responses. The simulated instrument is the mechanical prediction of $\text{MTR}_i(1) - \text{MTR}_i(0)$, holding income fixed at its pre-reform value. The IV estimand averages individual responses with weights proportional to the realised first-stage responses $\pi_i \equiv \text{MTR}_i(1) - \text{MTR}_i(0)$ (Angrist & Imbens, 1995; Borusyak et al., 2022). For the weights to be non-negative—and the estimand a convex average that cannot reverse the sign of a uniformly signed effect (Blandhol et al., 2026)—the π_i must share a common sign: first-stage monotonicity requires

$$(H.1) \quad \text{MTR}_i(1) \geq \text{MTR}_i(0) \quad \forall i,$$

with the reverse inequality for a reform that cuts rates. A violation—individuals whose realised marginal tax rate moves *against* the direction of the reform—puts negative weight on their responses, and the estimand loses its interpretation as a convex average.

H.2 When the condition fails

Under the strong assumptions standard in the literature of (i) a quasi-linear utility model, implying no income effects, (ii) continuous labour supply, and (iii) no heterogeneity in elasticities, condition (H.1) is not violated. The intuition is that, for any reform—whether it increases or decreases the marginal tax rate in a given bracket—the utility ranking between the kink point and interior solutions in brackets unaffected by the reform remains unchanged. As a result, any new interior optimum in a bracket unaffected by the reform is dominated by the kink point at the reform-affected bracket.

When either of the first two assumptions is relaxed—allowing income effects or labour-supply rigidities—and elasticities or preferences are heterogeneous, first-stage monotonicity can be violated.

We illustrate a violation of first-stage monotonicity using a model with income effects and heterogeneity in the curvature of the utility function. Instead of the standard quasi-linear

utility function, we consider a more general specification that allows for income effects and non-linear labour-supply responses.⁶⁶ Individuals derive utility from consumption c and hours worked h according to

$$u_i(c, h) = \frac{c^{1-\sigma_i}}{1-\sigma_i} - \chi_i \frac{h^{1+\varphi_i}}{1+\varphi_i},$$

where $\sigma_i > 0$ governs the curvature of utility from consumption, and hence the strength of income effects, $\varphi_i > 0$ governs the curvature of the disutility from labour, and $\chi_i > 0$ scales the marginal disutility of hours worked. Labour income is given by $y = wh$. We assume that hours rigidities restrict individuals to four discrete choices: 0, 20, 30, or 40 hours per week.

For the tax schedule, we consider the French income tax system as of 2015. In the part of the schedule on which we focus, labour income up to €26,791 per year is taxed at a marginal rate of 14%, and income above this threshold at 30%.⁶⁷ Because utility is defined over hours worked, we can equivalently express the tax schedule as a function of hours, with a kink at $h = k$ corresponding to the income threshold. In the example, we set the hourly wage to €20.6, so that the threshold at €26,791 corresponds to approximately 25 hours of work per week.

We study a hypothetical reform ($g = 1$) that eliminates the 30% bracket by raising the marginal tax rate above the threshold from 30% to 41%, equal to the next statutory rate in the French schedule.

We consider two preference types, Type A and Type B, and calibrate the model such that both share the same discrete optimum of 30 hours per week under no reform. The key parameter governing the relative strength of income and substitution effects is the consumption-curvature parameter σ_i . Type A is calibrated with $\sigma_A < 1$, so that the substitution effect dominates, whereas Type B is calibrated with $\sigma_B \gg 1$, so that the income effect dominates. For each type, the scale χ_i is set to deliver this common optimum. The resulting preference parameters are

$$(\sigma_A, \varphi_A, \chi_A) = (0.696, 0.300, 0.272), \quad (\sigma_B, \varphi_B, \chi_B) = (7.000, 0.300, 9.17 \times 10^{-30}).$$

⁶⁶Hours rigidities would be sufficient to illustrate a violation. To make the argument applicable more broadly, however, we consider a more general model with income effects.

⁶⁷To simplify the exposition, the example uses the statutory schedule, without the tax reduction for small liabilities of Section 2.1; the reduction does not affect the example, and results including it are available on request.

Under no reform, both types choose 30 hours per week, placing them in the 30% bracket:

$$\text{MTR}_A(0) = \text{MTR}_B(0) = 0.30.$$

The simulated instrument, computed by holding both types' incomes fixed at their 30-hour values, accordingly predicts the same treatment change for both: a rise in the marginal tax rate from 30% to 41%, that is, a predicted first-stage response of +0.11.

The optimal hours choices under the reform and under no reform are shown in Figure H.1.⁶⁸ Under the reform, the two types' optimal choices diverge:

$$h_A^* = 20 \text{ hours/week}, \quad h_B^* = 40 \text{ hours/week}.$$

The realised treatments therefore respond to the same shock in opposite directions:

$$\pi_A = \text{MTR}_A(1) - \text{MTR}_A(0) = 0.14 - 0.30 = -0.16 < 0,$$

$$\pi_B = \text{MTR}_B(1) - \text{MTR}_B(0) = 0.41 - 0.30 = +0.11 > 0.$$

Type A's marginal tax rate moves against the direction of the reform: despite being exposed to a shock that raises statutory rates—and despite the instrument predicting a tax increase for it—its realised marginal tax rate *falls*, violating condition (H.1). Type B's realised treatment moves with the reform.

First-stage monotonicity is therefore violated: the estimand places negative weight on Type A's response and, with heterogeneous elasticities, can even reverse the sign of a uniformly signed response—it is not weakly causal in the sense of Blandhol et al. (2026).

A strong income effect for Type B is not necessary for the violation: if Type B simply stays at 30 hours under the reform ($\pi_B = +0.11$) while Type A moves to 20 ($\pi_A = -0.16$), the same violation arises.

Models with participation responses can also generate violations of first-stage monotonicity. Leaving work in response to, for example, a reform that increases tax rates moves an individual

⁶⁸The optimal hours choice under each regime is determined by maximising $u_i(c, h)$ subject to the piecewise-linear budget constraint. There is no closed-form solution in general, so the optimum is obtained numerically by comparing utility across all candidate solutions: interior first-order-condition roots on each linear segment, the kink, and the corner at zero hours.




Figure H.1: Violation of first-stage monotonicity under heterogeneous labour-supply responses

Notes: Top panel. Budget constraints under the reform (red) and under no reform (blue). The reform raises the marginal tax rate above the second bracket threshold from 30% to 41%. Both Type A and Type B optimally choose 30 h/week under no reform. Under the reform, Type A chooses 20 h/week and is below the kink; Type B ($\sigma_B = 7$) chooses 40 h/week and is above the kink.

Bottom panel. Marginal tax rates under the reform (red) and under no reform (blue), as a function of hours worked, with the same optimal choices indicated.

out of the tax schedule altogether, that is, in France, to a marginal tax rate of zero, so that $\pi_i < 0$. If, at the same time, other individuals j remain employed and face the higher rate under the reform, $\pi_j > 0$, the realised treatment responds to the shock in opposite directions

across individuals, violating condition (H.1). Because participation responses to tax increases are empirically common, as we also document in this paper, this is a realistic and quantitatively important source of monotonicity violations in reform-based IV strategies.

We stress that the reform-based setting of this example is a foil: it is the canonical design from which our identification strategy departs, not our identifying variation. We now turn to our design.

H.3 The condition in the event-based setting

In the main design the schedule change is the fiscal exit of the first-born, which lowers the quotient—from 2.5 to 2 with one dependent child, from 3 to 2.5 with two (Section 3.2). Because household income is divided by the quotient, the marginal rate at every level of household income y satisfies $T'_t(y/Q') \geq T'_t(y/Q)$ for $Q' < Q$: the exit weakly raises the marginal rate at every point of the budget set, so only the rate-increase direction of condition (H.1) is at stake, and a violation cannot arise from the shift itself. It requires a behavioural relocation: a household reducing its income by enough that its realised marginal rate falls strictly below the rate it would have faced under the old quotient. The required relocation is large. An individual predicted into the next bracket has per-part income under the new quotient at or above the upper threshold of the pre-event bracket, and a violation requires it to fall below the lower threshold: household income must fall to less than the ratio of the two thresholds times its pre-event level. Under a gradient reform a small response suffices. Partner income makes even this channel unlikely: where partner income alone, divided by the post-event quotient, places the household above the tax-free threshold, a complete exit of the individual leaves the household in a positive bracket—a condition that can be checked cell by cell, since the cells fix the partner-income range.⁶⁹ The same floor arises at marriage, where the partner's income is added to the tax base, as the figure below illustrates.

Figure H.2 illustrates why such a monotonicity violation is highly unlikely with our instrument. The figure shows the marginal tax rate an individual faces as a function of own income, both when single (solid) and when married to a partner earning €40,000 (dashed). Crucially, the

⁶⁹Two residual channels remain: a simultaneous collapse of both spouses' incomes below the household tax-free threshold, and the partner's own exit from work, which lowers household income through the same mechanism and is not rare at the ages at which the event occurs.

marginal tax rate for a married individual remains strictly positive—at 14%—even when own income is zero, because the partner’s income already places the household above the tax-free threshold: even the most extreme reduction in labour supply, a complete exit, leaves the individual facing a 14% rate rather than a 0% rate. For an individual predicted to move from the 14% to the 30% bracket, a violation would require the realised marginal rate to fall strictly below 14%, which the partner’s income prevents unless it collapses as well. The logic extends to all individuals facing predicted tax increases; for predicted decreases, violations are likewise unlikely, since they would require a dramatic increase in labour supply.




Figure H.2: Marginal tax rate under marriage, partner earns €40,000

Notes: The figure shows the marginal tax rate as a function of own gross income under individual taxation (blue solid) and under joint taxation when married to a partner earning €40,000 (red dashed), based on the statutory 2019 French income tax schedule (Section 2.1). Income is not pooled, so own consumption corresponds to own after-tax income only. As a result, the marginal tax rate for a married individual is 14% even at zero own income, because the partner’s income of €40,000 already places the household above the tax-free threshold. Consequently, even a complete exit from the labour market cannot lower the individual’s marginal tax rate below 14%, which makes a monotonicity violation unlikely for individuals predicted to move from the 14% to the 30% bracket.

I Published estimates of work-income responses

Appendix Table I.1 records the estimates and distributional locations underlying Figure 1, together with each estimate’s design, estimand, horizon and population; ranking universes and source versions are documented in the notes. It includes only studies represented in the figure, while retaining alternative estimates from those studies for context.

Table I.1: Evidence base for the literature map of work income responses

Study	Outcome	Design	Estimand	Horizon	Population	Estimate (SE)	Identifying variation
Saez (2010)	Work income	Bunching	Compensated	1995-2004 pooled	EITC filers with one child, wage earners only	0.025 (0.036)^a	P13.4–P15.3 (C)
					EITC filers with two or more children, wage earners only	0.003 (0.022)^a	P19.1–P21.1 (C)
	Earnings				All EITC filers with one child	0.213 (0.033)	P13.4–P15.3 (C)
					Self-employed EITC filers with one child	1.101 (0.092)	
					All EITC filers with two or more children	0.152 (0.02)	P19.1–P21.1 (C)
Chetty, Friedman and Saez (2013)	Work income	Reform DiD (EITC around first-child birth)	Frictionless-model elasticity matching observed EITC refund changes	Event study around first-child birth, 2000-2005	EITC-eligible parents around the birth of the first child, 2000-2005	0.755 (0.054)	
						0.31 (0.018)^b	P0–P19.4 (C)
						0.14 (0.015)^b	P24.8–P54.5 (C)
						0.19 (0.019)	P0–P54.5 (C)
Chetty et al. (2011)	Work income	Reform IV (small reforms)	Observed elasticity; the paper argues frictions attenuate it	1994-2001 pooled	All wage earners	0.21 (0.012)	
	Work income (taxable base)	Bunching	Compensated (observed, frictions not netted out)			-0.005 (0.003)	not locatable
	Taxable income	Structural calibration	Structural lower bound (calibrated, frictions ≤ 5% of consumption)		Wage earners	~0.01 ^j	P75 (A)
	Work income (taxable base)	Bunching	Compensated (observed)		Married women, wage earners	≥0.34	not locatable
Bastani and Selin (2014)	Work income (taxable base)	Bunching	Compensated	1999-2005 pooled	Wage earners aged 15-70; anyone with active business income or closely-held-corporation ties removed	~0.02	P75 (A)
			Structural upper bound (tolerated loss = 1% of disposable income)		Wage earners aged 15-70	~0.00 ^k	P80 (A)
			Compensated			≤0.70	
Lehmann et al. (2013)	Taxable income	Reform IV	Marginal-rate elasticity, average net-of-tax rates controlled	2000-2008 pooled	All self-employed	0.024	
	Work income (taxable base)				Wage earners aged 15-70	~0.00	
Saez (2003)	Work income incl. contributions	Reform IV (small reforms)	Compensated	1-year, 2003-2006 pooled	Private-sector employees earning less than twice the minimum wage; public sector and self-employed excluded	0.225 (0.085)^g	P10–P60 (C)
						0.016 (0.096)	
						0.084 (0.173)^f	P44–P94 (C)
Saez (2003)	Work income	Reform IV (bracket creep, 1979-1981 inflation)	Compensated	1-year, 1979-1981	Tax returns in the NBER panel, 1979-1981; singles with AGI \$3,000-40,000 and married \$5,000-70,000 (1979 USD)	0.120 (0.196)	P55–P94 (C)

Continued on next page

Table I.1 continued

Study	Outcome	Design	Estimand	Horizon	Population	Estimate (SE)	Identifying variation	
Kleven and Schultz (2014)	Earnings	Reform IV	No income effects	3-year, 1984-2005	All workers, top 10 percent earners	0.085 (0.012)^c	P90–P100 (A)	
	Work income				All workers, top 20 percent earners	0.076 (0.008)^c	P80–P100 (A)	
					Wage earners, top 10 percent earners	0.081 (0.012)	P90–P100 (A)	
					Wage earners, top 20 percent earners	0.073 (0.009)	P80–P100 (A)	
Jakobsen and Sogaard (2022)	Work income incl. contributions plus transfers	Local reform IV (2SLS local linear)	No income effects	4-year, 2008-2012 (pre-period 2004-2008)	Aged 15+, pre-reform personal income above DKK 250k; self-employed excluded	~0.038	P74.9–P76.9 (C)	
						~0.051	P82.5–P83.8 (C)	
						~0.047	P87.7–P88.5 (C)	
						~0.031	P91.1–P91.5 (C)	
						~–0.004	P93.3–P93.6 (C)	
	Taxable income	Reform IV			4-year, 2008-2012		~0	P94.9–P95 (C)
							~0.028	P95.6–P95.7 (C)
							~–0.001	P96.2–P96.3 (C)
							0.016 (0.007)	P74.9–P100 (C)
							0.202 (0.007)	
Miao, Selin and Soderstrom (2025)	Earnings	Reform difference-in-differences (grouped)	No income effects (grouped Wald / DiD)	4-year, 2015 vs 2019	Percentile groups 96-100 (treatment) vs 88-95 (control), aged 20-65	0.156 (0.023)	P96–P100 (A)	
					Percentile groups 96-100 vs 88-95, groups 94-97 excluded	0.171 (0.019)		
Kleven et al. (2025)	Work income incl. contributions	Reform IV, job switchers	Long-run population elasticity under orthogonality of switching timing	Four-year	Treated job switchers above the threshold	0.487 (0.026)^d	P70–P98 (A)	
				Two-year	Treated job switchers in the top decile	0.642 (0.029)^d	P90–P98 (A)	
					All treated job switchers	0.486 (0.022)	P70–P98 (A)	
Gelber (2014)	Work income	Reform IV	Uncompensated	1-year, 1989-90 and 1990-91 pooled	Treated job switchers in the top decile	0.725 (0.029)	P90–P98 (A)	
					Married men aged 18-65 with positive earnings in 1989 or 1990	0.210 (0.02)^e	not locatable	
			Compensated		Married women aged 18-65 with positive earnings in 1989 or 1990	0.230 (0.03)^e		
					Married men aged 18-65 with positive earnings in 1989 or 1990	0.410 (0.02)		
					Married women aged 18-65 with positive earnings in 1989 or 1990	0.470 (0.02)		
Sigaard (2023)	Work income	Reform IV	No income effects	3-year	Danish wage earners with reported contract and performed hours, 2002-2016	0.058 (0.0037)	not locatable	
	Hours worked Taxable income					0.082 (0.0091) 0.095 (0.0042)		
Blomquist and Selin (2010)	Work income (taxable base)	Reform IV	Uncompensated	10-year, 1981 vs 1991	Married or cohabiting men; self-employed excluded	0.210 (0.074)¹	not locatable	
					Married or cohabiting women; self-employed excluded	1.442 (0.832)¹		
He et al. (2021)	Work income (taxable base)	Reform IV	Compensated (income effects small and insignificant)	One-year	Wage earners aged 18-60 in one (unnamed) Chinese city	2.423 (0.188)^m	not locatable	
		Bunching		Pooled months		0.1 (0.02) 0.41 (0.07) 0.09 (0.03)		
					Reform IV (revised, long-run)	Long-run	~4	

Continued on next page

Table I.1 continued

Study	Outcome	Design	Estimand	Horizon	Population	Estimate (SE)	Identifying variation
Kleven–Waseem (2013)	Work income (taxable base)	Bunching (notch)	Observed (friction-attenuated) elasticity	Annual, 2006-2007	Pakistani wage-earner filers 2006-07 (under 2% of the working-age population, over 99% male)	~ 0 (0.000-0.004	not locatable across notches) ⁿ
			Structural (frictionless) elasticity, bunching-hole lower bound			0.024 (0.009)	
						0.035 (0.010)	
						0.034 (0.019)	
						0.001 (0.013)	

Notes: One row per coded estimate; **bold** estimates are plotted in Figure 1, with superscripts matching the figure’s label letters. Within a study block, repeated entries are shown once. (–) indicates that no standard error is reported.

Placement. Ranks in the identifying-variation column: (A) reported by the study; (B) computed by us from a closely matching official distribution; (C) computed from an official distribution with a related income concept, drawn dashed in the figure; (D) weak reconstruction, (Q) qualitative statement, (U) unrecoverable — not placed, shown in the figure’s right panel. The column gives the income range of the treated group only, never the estimation sample: control groups located elsewhere in the income distribution (Miao et al.: P88–95; Kleven et al.: below the reform threshold; Jakobsen–Søgaard: the validation region) are not part of the stated range, while for Kleven–Schultz the sample and the treated band coincide because identification uses within-sample variation in reform exposure.

Estimands. Not harmonised: bunching estimates are compensated elasticities; Gelber’s plotted pair is uncompensated (the study’s compensated estimates are 0.41 and 0.47); most register-based reform estimates are specified without income effects; and Chetty–Friedman–Saez report the elasticity a frictionless model requires to match observed EITC refund changes. Reported income effects are small throughout, so these differences are second order relative to design and horizon. Statutory-taxable-income estimates on wage-earner samples enter as work income (taxable base); population-wide taxable-income studies remain in the project archive.

Sources. Values are taken from the published articles, with exceptions: Chetty et al. (2011) from the revised NBER working paper; Sigurdsson (2025) from the November 2024 manuscript; Blomquist–Selin verified from CESifo Working Paper 2644. Jakobsen–Søgaard points and confidence bands are digitized from their Figure 3, Panel D; their local estimates use triangular weights within DKK 50,000 of each centre. Kleven–Waseem’s observed elasticity is derived as the paper instructs (structural lower bound times the responding share); no standard error applies.